

# The Part-Time Parliament

Leslie Lamport, 1998

Carlos Eduardo B. Bezerra

Faculty of Informatics  
Università della Svizzera italiana

May 25th, 2011

# Outline

- 1 Introduction
  - A long way to Paxos
  - The consensus problem
- 2 About correctness
  - Notations used
  - Consistency conditions
  - Proof sketch
- 3 The Synod Protocol
  - Preliminary protocol
  - Basic protocol
  - Complete protocol
- 4 The Paxos Protocol
  - Running multiple consensus instances
  - Relevance to computer science

# Introduction

Interesting information about ‘The Part-Time Parliament’:

- the algorithm was first described in 1989 as a technical report [1]
- the paper was submitted in 1990...and accepted in 1998
- several other publications, such as [2, 3, 4], had explaining it as their main purpose...

Some people do use the described **Paxos** protocol:

- Google uses Paxos in their Chubby distributed lock service
- IBM supposedly uses Paxos in their SAN Volume Controller product
- Microsoft uses Paxos in the Autopilot cluster management service
- WANDisco uses Paxos in their replication technology
- Keyspace uses Paxos as its basic replication primitive

The paper mostly uses an analogy with the fictional ancient island of Paxos, where:

- The parliament's (group's) legislators (processes) and their messengers (communication channels) may be unpredictably absent;
- There may be several ballots for each decree to be passed (rounds in each consensus instance);
- Legislators – and messengers – may be lazy or absent, but not dishonest (non-Byzantine failure model\*)

\* The algorithm was extended later, however, to tolerate Byzantine failures [5]

# Definition

**Consensus** is the process of agreeing on one single result among a group of participants. Assuming a collection of processes that can propose values, a consensus algorithm ensures that:

- only a single value among the proposed ones may be chosen;
- if no value has been proposed, then no value is chosen;
- a process never learns that a value has been chosen unless it really has been.

## Some notations used

A value is chosen after a series of numbered **ballots**. For each ballot  $B$ , we have that:

- $B_{bal}$  is its ballot number;
- $B_{qrm}$  is a nonempty finite set of processes (quorum);
- $B_{vot}$  is the set of processes who voted in  $B$  – each process may decide to participate (vote) or not in  $B$ ;
- $B_{val}$  is the value being voted on in that ballot.

A ballot is said to be *successful* when  $B_{qrm} \subseteq B_{vot}$  – which means that every quorum member voted.

## Some more notations

In each ballot, each process can cast a vote (accept or not a given proposed value). A vote  $v$  has the components:

- $v_p$ , which is the process  $p$  who cast it;
- $v_{bal}$ , the number of the ballot for which it was cast;
- $v_{val}$ , a value voted on.

For the set  $\beta$  of ballots, the set  $Votes(\beta)$  is defined as  $\{v : \exists B \in \beta \mid v_p \in B_{vot} \wedge v_{bal} = B_{bal} \wedge v_{val} = B_{val} \}$

## Still more notations. . .

We then define  $MaxVote(b, p, \beta)$  as the vote  $v$  with highest  $v_{bal}$  in  $\{v \in Votes(\beta) : (v_p = p) \wedge (v_{bal} < b)\} \cup \{null\}$ ,  
 where  $null$  means that  $p$  didn't vote in any  $B \in \beta$  ( $null_{bal} = -\infty$ )

We also define  $MaxVote(b, Q, \beta)$  as the vote  $v$ , such that  $v_{bal}$  is the highest among all  $MaxVote(b, p, \beta)$  with  $p \in Q$

With such notations, we can specify the conditions which ensure consistency and allow for progress.



# Consistency conditions

Consistency is guaranteed if the following conditions are ensured:

- $B1(\beta)$   
 $\forall B, B' \in \beta : (B \neq B') \Rightarrow (B_{bal} \neq B'_{bal})$
- $B2(\beta)$   
 $\forall B, B' \in \beta : B_{qrm} \cap B'_{qrm} \neq \emptyset$
- $B3(\beta)$   
 $\forall B, B' \in \beta :$   
 $(MaxVote(B_{bal}, B_{qrm}, \beta)_{bal} \neq -\infty) \Rightarrow$   
 $\Rightarrow (B_{val} = MaxVote(B_{bal}, B_{qrm}, \beta)_{val})$

# Proof sketch

**Lemma.** If  $B1(\beta)$ ,  $B2(\beta)$  and  $B3(\beta)$  hold, then

$$((B_{qrm} \subseteq B_{vot}) \wedge (B'_{bal} > B_{bal})) \Rightarrow (B'_{val} = B_{val})$$

for any  $B, B'$  in  $\beta$ .

**Proof of Lemma.** Let  $\Psi(B, \beta)$  be defined as:

$$\Psi(B, \beta) = \{B' \in \beta : (B'_{bal} > B_{bal}) \wedge (B'_{val} \neq B_{val})\}$$

Assume that  $B_{qrm} \subseteq B_{vot}$ . By contradiction, assume that  $\Psi(B, \beta) \neq \emptyset$ .

(1) As  $\Psi(B, \beta) \neq \emptyset$ , we pick  $C \in \Psi(B, \beta)$ , such that

$$C_{bal} = \min\{B'_{bal} : B' \in \Psi(B, \beta)\}$$

(2) From (1) and from the definition of  $\Psi(B, \beta)$ , we have that  $C_{bal} > B_{bal}$

# Proof sketch

continuing...

(3) As  $B_{qrm} \subseteq B_{vot}$  and  $\forall B, B' \in \beta : B_{qrm} \cap B'_{qrm} \neq \emptyset$ , we have that  $B_{vot} \cap C_{qrm} \neq \emptyset$

(4) As  $C_{bal} > B_{bal}$  and  $B_{vot} \cap C_{qrm} \neq \emptyset$ , we have that  $MaxVote(C_{bal}, C_{qrm}, \beta)_{bal} \geq B_{bal}$

(5) From (4) and the definition of  $MaxVote(C_{bal}, C_{qrm}, \beta)$ , we know that such vote exists in  $Votes(\beta)$  (i.e. it is not *null*).

(6) From (5) and  $B3(\beta)$ , we have that  $MaxVote(C_{bal}, C_{qrm}, \beta)_{val} = C_{val}$

(7) From (6), and as  $C_{val} \neq B_{val}$  (from  $\Psi(B, \beta)$ ), we have that  $MaxVote(C_{bal}, C_{qrm}, \beta)_{val} \neq B_{val}$

# Proof sketch

continuing...

(8) As  $MaxVote(C_{bal}, C_{qrm}, \beta)_{bal} \geq B_{bal}$  (from (4)) and  $MaxVote(C_{bal}, C_{qrm}, \beta)_{val} \neq B_{val}$  (from (7)), we have that  $MaxVote(C_{bal}, C_{qrm}, \beta)_{bal} > B_{bal}$

(9) From (8), and since  $MaxVote(C_{bal}, C_{qrm}, \beta)_{val} \neq B_{val}$ , we have that  $MaxVote(C_{bal}, C_{qrm}, \beta) \in Votes(\Psi(B, \beta))$

(10) By the definition of  $MaxVote(C_{bal}, C_{qrm}, \beta)$ , we have that  $MaxVote(C_{bal}, C_{qrm}, \beta)_{bal} < C_{bal}$

(11) From (10), we have that  $\exists B' : B' \in \Psi(B, \beta) \wedge B'_{bal} < C_{bal}$ , which means that

$$C_{bal} \neq \min\{B'_{bal} : B' \in \Psi(B, \beta)\}$$

which is a contradiction with (1).  $\square$

# Proof sketch

continuing...

With the lemma, we show that, for a given set  $\beta$  of ballots which hold  $B1$ ,  $B2$  and  $B3$ , then any two successful ballots decided the same value.

**Theorem 1.** If  $B1(\beta)$ ,  $B2(\beta)$  and  $B3(\beta)$  hold, then

$$((B_{qrm} \subseteq B_{vot}) \wedge (B'_{qrm} \subseteq B'_{vot})) \Rightarrow (B'_{val} = B_{val})$$

for any  $B, B'$  in  $\beta$ .

**Proof of Theorem.** If  $B'_{bal} = B_{bal}$  then  $B1(\beta)$  implies  $B' = B$ . If  $B'_{bal} \neq B_{bal}$ , the theorem follows immediately from the lemma.  $\square$

# Proof sketch

continuing...

If there are enough correct (non-faulty) processes, then a new ballot may be conducted while  $B1$ ,  $B2$  and  $B3$  are preserved.

**Theorem 2.** Let  $b$  be a ballot number, and let  $Q$  be a set of processes such that  $b > B_{bal}$  and  $Q \cap B_{qrm} \neq \emptyset$ , for all  $B \in \beta$ . If  $B1(\beta)$ ,  $B2(\beta)$  and  $B3(\beta)$  hold, then there is a ballot  $B'$  with  $B'_{bal} = b$  and  $B'_{qrm} = B'_{vot} = Q$  such that  $B1(\beta \cup \{B'\})$ ,  $B2(\beta \cup \{B'\})$  and  $B3(\beta \cup \{B'\})$  hold.

## Proof of Theorem.

- $B1(\beta \cup \{B'\})$  follows from  $B1(\beta)$  and the fact that  $B'_{bal} = b$ , and the assumption about  $b$ .
- $B2(\beta \cup \{B'\})$  follows from  $B2(\beta)$  and the fact that  $B'_{qrm} = Q$ , and the assumption about  $Q$ .
- For  $B3(\beta \cup \{B'\})$ , if  $MaxVote(b, Q, B) = -\infty$  then let  $B'_{val}$  be any value; otherwise, let  $B'_{val} = MaxVote(b, Q, B)_{val}$ .  $B3(\beta \cup \{B'\})$  then follows from  $B3(\beta)$ .  $\square$

## Preliminary protocol

A protocol can be derived directly from  $B1(\beta)$ ,  $B2(\beta)$  and  $B3(\beta)$ , so that all these properties are maintained throughout its execution and a value may\* be agreed upon.

- To maintain  $B1(\beta)$ , whenever  $p$  starts a new ballot  $B$ , it assigns a sequence number  $seq$ , which was never used by  $p$ . To avoid collisions,  $p$ 's unique id is appended. Then  $B_{bal} = (seq, p_{id})$ .
- To maintain  $B2(\beta)$ , each ballot  $B$  has  $B_{qrm}$  as the majority set of the processes. As it is impossible to choose two disjoint majority sets from the same group\*,  $B2(\beta)$  is satisfied.
- To maintain  $B3(\beta)$ , before initiating a ballot  $B$ ,  $p$  has to find out what is  $MaxVote(B, B_{qrm}, \beta)$ , which is done by finding the  $MaxVote(B, q, \beta)$  for each  $q$  in  $B_{qrm}$ .
  - If  $MaxVote(B, B_{qrm}, \beta) = null$ , then  $B_{val}$  can be anything; otherwise  $B_{val} = MaxVote(B, B_{qrm}, \beta)_{val}$ .

\* If half of the processes is faulty, consistency is still respected and there is no deadlock, but the ballot  $B$  does not succeed

## Preliminary protocol

Phase 1) Find  $MaxVote(B, B_{qrm}, \beta)_{val}$

**Phase 1a)**  $p$  picks a ballot number  $b$  and sends  $NextBallot(b)$  to other processes

**Phase 1b)** each process  $q$  responds with  $LastVote(b, v)$  to  $p$ , where  $v$  is the vote cast by  $q$  with the highest ballot number less than  $b$  (or *null*)

Phase 2) Execute the ballot  $B$

**Phase 2a)** after receiving  $LastVote(b, v)$  from a majority  $Q$  of the processes,  $p$  initiates the ballot with number  $b$  and value  $val$  such that  $B3(\beta)$  is maintained, by sending  $BeginBallot(b, val)$  to every process in  $Q$

**Phase 2b)** upon receiving  $BeginBallot(b, val)$ ,  $q$  will not vote on it if it has already sent  $LastVote(b', v')$ , such that  $b' > b$ ; otherwise, it sends  $Voted(b, q)$  to  $p$  and records this vote in its memory



## Preliminary protocol

Phase 3) Announce the result of consensus

**Phase 3a)** if  $p$  receives  $Voted(b, q)$  from a majority  $Q$  of processes, then it sends  $Success(val)$  to all processes

**Phase 3b)** upon receiving  $Success(val)$ , the process learns that the group achieved consensus on  $val$ .

Every phase of this protocol is optional. Consistency is still maintained, although progress is not guaranteed.

Each process  $q$  must keep  $MaxVote(b, q, \beta)$  for each ballot numbered  $b$  and remember all  $LastVote(b, v)$  messages it has sent.

# Basic protocol

In the preliminary protocol, much information must be kept. A simpler, yet correct, protocol can be derived. Each process  $p$  keeps then only:

***lastTried*** $[p]$  – the number of the last ballot  $p$  tried to initiate ( $-\infty$  if none)

***prevVote*** $[p]$  – the highest ballot for which  $p$  ever voted ( $-\infty$  if none)

***nextBal*** $[p]$  – the highest ballot number  $b$  for which  $p$  sent  $LastVote(b, v)$  ( $-\infty$  if none)

In the preliminary protocol,  $LastVote(b, v)$  represented a “promise” from  $q$  not to vote in any ballot numbered  $b' : LastVote(b, v)_{bal} < b' < b$

Here,  $LastVote(b, v)$  represents a stronger “promise” from  $q$  not to vote in any ballot numbered  $b' < b$

# Basic protocol

## Phase 1)

**Phase 1a)**  $p$  picks a ballot number  $b > lastTried[p]$  and sends  $NextBallot(b)$  to other processes

**Phase 1b)** upon receipt of  $NextBallot(b)$  from  $p$ , if  $b > nextBal[q]$ ,  $q$  sets  $nextBal[q]$  to  $b$  and sends  $LastVote(b, v)$  to  $p$ , where  $v = prevVote[q]$

## Phase 2)

**Phase 2a)** after receiving  $LastVote(b, v)$  from a majority  $Q$ , where  $b = lastTried[p]$ ,  $p$  initiates ballot number  $b$  with value  $val$  such that  $B3(\beta)$  is maintained, by sending  $BeginBallot(b, val)$  to every process in  $Q$

**Phase 2b)** upon receiving  $BeginBallot(b, val)$  with  $b = nextBal[q]$ ,  $q$  votes on ballot number  $b$ , set  $prevVote[q]$  to this vote and sends  $Voted(b, q)$  to  $p$

## Basic protocol

Phase 3)

**Phase 3a)** if  $p$  receives  $Voted(b, q)$  from a majority  $Q$  of processes, then it sends  $Success(val)$  to all processes

**Phase 3b)** upon receiving  $Success(val)$ , the process learns that the group achieved consensus on  $val$ .

Consistency is still maintained, and progress is still not guaranteed.

Anyway, it is impossible to guarantee termination of consensus in an asynchronous system [6]. . .

# Complete protocol

- Same steps of the basic protocol
- Each process performs phases 1b – 3b as soon as they can
- Some process must execute phase 1a
  - Never initiating a ballot prevents termination
  - However, too frequent initiation also does it
  - A single process (leader) should be the only one initiating ballots
- Knowledge of time, and of a maximum message delay  $\Delta$ , is needed
  - Not exactly an asynchronous system anymore

# Complete protocol

- A minimal execution of the basic protocol would take at most  $5\Delta$
- A leader would ensure progress, but that requires leader election
  - when a leader seems to have failed, a new one should be elected (the election takes time  $T_{el}$ )
  - if not enough responses to (1a) or (2a) arrive after  $2\Delta$ , a ballot  $B$  had started before  $T_{el}$ 
    - or too many failures or message losses occurred...
  - in this case, a ballot with number  $b' > B_{bal}$  may be started by the leader – *NACKs* to (1a) or (2a) may have carried  $B_{bal}$
  - in the worst case, the leader receives a *NACK* for (2a),  $4\Delta$  after  $T_{el}$
  - as no other process initiated any ballot after  $T_{el}$ , the ballot numbered  $b'$  would succeed  $5\Delta$  after the *NACK*
- The complete protocol, with leader election, terminates at most after  $T_{el} + 9\Delta$ 
  - if no message was lost, every correct process learns the decided value

# The Paxos Protocol

## Multiple consensus instances

- Each instance has an id  $I_{id}$  – and there should be a leader
- Phases (1a) and (1b) can be executed for many of them – e.g.  $p$  sends  $NextBallot(b, n)$ , for all instances  $I$  with  $I_{id} > n$
- The response from  $q$  (1b) would contain a  $LastVote(I_{id}, b, v)$  message for each instance  $I : I_{id} > n$  that  $q$  voted in
- Another process can send  $NextBallot(b', n')$ , with  $b' > b$ , and  $p$ 's (2a) messages would be ignored from then on
- This brings the protocol latency, on most cases, down to  $3\Delta$ 
  - no leader changes, enough correct processes and enough messages
- If each process sent messages (2b) to every other one, we'd have  $2\Delta$
- If instance  $I^A$  was passed (phase (3b) concluded) before  $I^B$  had a value proposed (phase (2a)), then  $I_{id}^A < I_{id}^B$

# The Paxos Protocol

## Relevance to computer science

- It may be used to implement state-machine replication
  - A system (e.g. a database) may be implemented as a set of states and deterministic state changes
  - The group of processes would be the set of replicas of the system
  - Each state changed would be agreed upon by means of consensus (an ordered instance of Paxos)
  - One of the replicas might be elected as leader, and client requests (state changes) would be sent to it
  - As all replicas would have a consistent set of state changes, so would be their state
- Many derivations from Paxos have been published since it was first described
  - Cheap Paxos, Fast Paxos, Generalized Paxos, Multi-Paxos, Byzantine Paxos, Fast Byzantine Paxos, Fast Byzantine Multi-Paxos, Ring Paxos



## References

- [1] Lamport, L. **The part-time parliament**, Technical Report 49, Systems Research Center, Digital Equipment Corp., 1989
- [2] Lamport, L. **Paxos made simple**, ACM SIGACT News, 2001
- [3] De Prisco, R., Lampon, B., Lynch, N. **Revisiting the Paxos algorithm**, Distributed Algorithms, 1997
- [4] Lampon, B. **How to build a highly available system using consensus**, Distributed Algorithms, 1996
- [5] Castro, M., Liskov, B., **Practical Byzantine fault tolerance**, Operating Systems Design and Implementation, 1999
- [6] Fischer, M., Lynch, N., Paterson, M., **Impossibility of distributed consensus with one faulty process**, J. ACM 32 1 (Jan.), 374-382, 1985

# Thank you!

Questions?