



virtutech

Simics User Guide for Unix

Simics Version 3.0

Revision 1376
Date 2007-01-24

© 1998–2006 Virtutech AB
Norrtullsgatan 15, SE-113 27 STOCKHOLM, Sweden

Trademarks

Virtutech, the Virtutech logo, Simics, and Hindsight are trademarks or registered trademarks of Virtutech AB or Virtutech, Inc. in the United States and/or other countries.

The contents herein are Documentation which are a subset of Licensed Software pursuant to the terms of the Virtutech Simics Software License Agreement (the “Agreement”), and are being distributed under the Agreement, and use of this Documentation is subject to the terms the Agreement.

This Publication is provided “as is” without warranty of any kind, either express or implied, including, but not limited to, the implied warranties of merchantability, fitness for a particular purpose, or non-infringement.

This Publication could include technical inaccuracies or typographical errors. Changes are periodically added to the information herein; these changes will be incorporated in new editions of the Publication. Virtutech may make improvements and/or changes in the product(s) and/or the program(s) described in this Publication at any time.

Contents

I	Simics Documentation	11
1	About Simics Documentation	13
1.1	Conventions	13
1.2	Simics Guides and Manuals	13
	Simics Installation Guide for Unix and for Windows	13
	Simics User Guide for Unix and for Windows	14
	Simics Eclipse User Guide	14
	Simics Target Guides	14
	Simics Programming Guide	14
	DML Tutorial	14
	DML Reference Manual	14
	Simics Reference Manual	14
	Simics Micro-Architectural Interface	14
	RELEASENOTES and LIMITATIONS files	15
	Simics Technical FAQ	15
	Simics Support Forum	15
	Other Interesting Documents	15
2	Glossary	17
II	Simulating with Simics	21
3	Introduction	23
3.1	Hosts and Targets	24
3.2	Host Recommendations	24
3.3	Simics Targets	25
3.3.1	AlphaPC 164LX	25
3.3.2	ARM SA1110	25
3.3.3	Ebony	25
3.3.4	Fiesta	25
3.3.5	IA-64 460GX	25
3.3.6	Malta/MIPS4kc	25
3.3.7	PM/PPC	26
3.3.8	Simple PPC64	26

3.3.9	Serengeti	26
3.3.10	SunFire	26
3.3.11	x86 440BX	26
3.4	Simics Version Number	26
3.5	Simics Compatibility	27
4	First Steps	29
4.1	Launch Simulation	29
4.2	Running the Simulation	30
4.3	Checkpointing	31
4.4	Hindsight	31
4.5	Getting Files into a Simulated System	33
4.6	Debugging	34
4.7	Tracing	39
4.8	Scripting	42
4.9	Simple Virtual Network	43
4.10	Connect to a Real Network	45
5	Command-line Interface: Basics	49
6	Configuration and Checkpointing	53
6.1	Basics	53
6.2	Checkpointing	54
6.2.1	Attributes	55
6.2.2	Images	56
	Image Search Path	57
6.2.3	Saving and Restoring Persistent Data	58
6.2.4	Modifying Checkpoints	58
6.2.5	Merging Checkpoints	59
6.3	Inspecting the Configuration	59
6.4	Components	60
6.4.1	Component Definitions	60
6.4.2	Importing Component Commands	60
6.4.3	Creating Components	60
6.4.4	Connectors	61
6.4.5	Instantiation	62
6.4.6	Inspecting Component Configurations	63
6.4.7	Accessing Objects from Components	63
6.4.8	Available Components	64
6.5	Ready-to-run Configurations	64
6.5.1	Customizing the Configurations	65
6.5.2	Adding Devices to Existing Configurations	66

CONTENTS

7	Managing Disks, Floppies, and CD-ROMs	69
7.1	Working with Images	70
7.1.1	Saving Changes to an Image	70
7.1.2	Reducing Memory Usage Due to Images	71
7.1.3	Using Read/Write Images	72
7.1.4	Editing Images Using Mtools	73
7.1.5	Editing Images Using Loopback Mounting	73
7.1.6	Constructing a Disk from Multiple Files	74
7.1.7	The Craff Utility	75
7.2	CD-ROMs and Floppies	76
7.2.1	Accessing a Host CD-ROM Drive	76
7.2.2	Accessing a CD-ROM Image File	76
7.2.3	Accessing a Host Floppy Drive	77
7.2.4	Accessing a Floppy Image File	77
7.3	Using SimicsFS	78
7.3.1	Installing SimicsFS on a Simulated Linux System	78
7.3.2	Installing SimicsFS on a Simulated Solaris System	79
7.3.3	Using SimicsFS	80
7.4	Importing a Real Disk into Simics	80
8	Simics Scripting Environment	83
8.1	Script Support in CLI	83
8.1.1	Variables	83
8.1.2	Command Return Values	84
8.1.3	Control Flow Commands	84
8.1.4	Integer Conversion	86
8.1.5	Accessing Configuration Attributes	86
8.1.6	Script Branches	86
	Introduction to Script Branches	86
	Waiting for Haps in Script Branches	87
	How Script Branches Work	87
	Script Branch Commands	87
	Variables in Script Branches	88
	Canceling Script Branches	89
	Script Branch Limitations	89
8.2	Scripting Using Python	89
8.2.1	Python in Simics	89
8.2.2	Accessing CLI Variables from Python	90
8.2.3	Accessing the Configuration from Python	90
	Configuration Objects	90
	Creating Configurations in Python	91
8.2.4	Accessing Command-Line Commands from Python	92
8.2.5	The Simics API	92
8.2.6	Haps	93
	Example of Python Callback on a Hap	93

III Simics Networking	95
9 Network Simulation	97
9.1 Ethernet Links	97
9.2 Link Object Timing	98
9.3 IP Services	99
9.3.1 IP Based Routing	100
9.3.2 DHCP and BOOTP	101
9.3.3 DNS	102
9.3.4 TFTP	102
9.4 Distributed Network Simulation	103
9.5 Serial Links	103
10 Connecting to a Real Network	105
10.1 Accessing Host Ethernet Interfaces	105
10.1.1 Raw Access	106
10.1.2 TAP Access	106
10.2 Selecting Host Ethernet Interface	107
10.3 Preparing for the Examples	108
10.4 Connection Types	111
10.4.1 Port Forwarding	113
The connect-real-network Command	113
Example	114
Incoming Port Forwarding	114
Example	115
Outgoing Port Forwarding	116
Example	116
NAPT	117
Example	118
DNS Forwarding	119
Example	119
10.4.2 Ethernet Bridging	120
Example	122
10.4.3 IP Routing	123
Example	124
10.4.4 Host Connection	126
Example	127
10.5 Performance	128
10.6 Troubleshooting	128
11 Distributed Simulation	131
11.1 Synchronization	131
11.2 Architecture	132
11.3 Running distributed	132
11.4 Example of Distributed Simulation and Network	133

CONTENTS

IV	Developing with Simics	137
12	Debugging Tools	139
12.1	Breakpoints	139
12.1.1	Memory Breakpoints	139
12.1.2	Temporal Breakpoints	141
12.1.3	Control Register Breakpoints	141
12.1.4	I/O Breakpoints	141
12.1.5	Graphics Breakpoints	142
12.1.6	Text Output Breakpoints	142
12.1.7	Magic Instructions and Magic Breakpoints	143
12.2	Using GDB with Simics	146
12.2.1	Remote GDB and Shared Libraries	147
12.2.2	Using GDB with Hindsight	148
12.2.3	Compiling GDB	150
12.3	Symbolic Debugging Using Symtable	151
12.3.1	Symtables and Contexts	151
12.3.2	Sample Session	151
12.3.3	Source Code Stepping	154
12.3.4	Symbolic Breakpoints	155
12.3.5	Reading Debug Information from Binaries	156
12.3.6	Loading Symbols from Alternate Sources	156
12.3.7	Multiple Debugging Contexts	156
12.3.8	Scripted Debugging	159
13	Profiling Tools	161
13.1	Instruction Profiling	161
13.1.1	Virtual Instruction Profiling	162
13.2	Data Profiling	163
13.3	Examining the Profile	164
V	Advanced Simics Usage	167
14	Startup Options	169
14.1	Simulation Modes	169
14.1.1	Normal mode	170
14.1.2	Memory Timing with -stall	170
14.1.3	Micro Architectural Simulation with -ma	170
14.2	Common Options	170
15	The Command Line Interface	173
15.1	Invoking Commands	173
15.1.1	How are Arguments Resolved?	174
15.1.2	Namespace Commands	175
15.1.3	Expressions	175

15.1.4	Interrupting Commands	176
15.2	Tab Completion	176
15.3	Help System	176
15.4	Simics's Search Path	179
15.5	Using the Pipe Command	181
16	Memory Transactions	183
16.1	Observing Memory Transactions	184
16.2	Stalling Memory Transactions	185
16.3	Observing Instruction Fetches	185
16.4	Simulator Translation Cache (STC)	186
16.5	Summary of Simics Memory System	186
17	Understanding Simics Timing	189
17.1	Events	189
17.2	Instruction Execution Timing	190
	Simics in-order	190
	Stalling	190
	Simics MAI	190
	Choosing an Execution Mode	190
	Changing the Step Rate	191
	Suspending Time or Execution	192
17.3	Multiprocessor Simulation	193
18	Cache Simulation	197
18.1	Introduction to Cache Simulation with Simics	197
18.2	Simulating a Simple Cache	198
18.3	Example Machines	199
18.4	A More Complex Cache System	199
18.5	Workload Positioning and Cache Models	203
18.6	Using g-cache	203
18.7	Understanding g-cache Statistics	204
18.8	Speeding up g-cache simulation	205
18.9	Cache Miss Profiling	206
18.10	Using g-cache with Several Processors	208
18.11	g-cache Limitations	208
19	Memory Spaces	211
19.1	Memory Space Basics	211
19.2	Memory Space Commands	213
19.3	Memory Mapping Types	213
19.4	Avoiding Circular Mappings	214

20	PCI Support in Simics	215
20.1	Introduction	215
20.2	Configuration Space	217
20.3	Memory and I/O Spaces	218
20.4	PCI Interrupts	219
20.5	Expansion ROM	219
20.6	PCI Express	219
20.7	Other PCI Features	219
20.7.1	Master Abort	219
20.7.2	Target Abort	220
20.7.3	Message Signaling Interrupt	220
20.7.4	Special Cycles	220
20.7.5	System Error (SERR#)	220
20.7.6	Parity Error (PERR#)	220
20.7.7	Interrupt Acknowledge	220
20.7.8	VGA Palette Snooping	220
21	Driving Context Changes	221
21.1	Switching Contexts Manually	221
21.2	Process Trackers	222
21.3	Switching Contexts Automatically	225
22	Understanding Hindsight	227
22.1	Command Overview	227
22.2	Performance	228
	Index	229

Part I

Simics Documentation

Chapter 1

About Simics Documentation

1.1 Conventions

Let us take a quick look at the conventions used throughout the Simics documentation. Scripts, screen dumps and code fragments are presented in a `monospace` font. In screen dumps, user input is always presented in bold font, as in:

```
Welcome to the Simics prompt
simics> this is something that you should type
```

Sometimes, artificial line breaks may be introduced to prevent the text from being too wide. When such a break occurs, it is indicated by a small arrow pointing down, showing that the interrupted text continues on the next line:

```
This is an artificial ?
line break that shouldn't be there.
```

The directory where Simics is installed is referred to as `[simics]`, for example when mentioning the `[simics]/README` file. In the same way, the shortcut `[workspace]` is used to point at the user's workspace directory.

1.2 Simics Guides and Manuals

Simics comes with several guides and manuals, which will be briefly described here. All documentation can be found in `[simics]/doc` as Windows Help files (on Windows), HTML files (on Unix) and PDF files (on both platforms). The new Eclipse-based interface also includes Simics documentation in its own help system.

Simics Installation Guide for Unix and for Windows

These guides describe how to install Simics and provide a short description of an installed Simics package. They also cover the additional steps needed for certain features of Simics to work (connection to real network, building new Simics modules, ...).

Simics User Guide for Unix and for Windows

These guides focus on getting a new user up to speed with Simics, providing information on Simics features such as debugging, profiling, networks, machine configuration and scripting.

Simics Eclipse User Guide

This is an alternative User Guide describing Simics and its new Eclipse-based graphical user interface.

Simics Target Guides

These guides provide more specific information on the different architectures simulated by Simics and the example machines that are provided. They explain how the machine configurations are built and how they can be changed, as well as how to install new operating systems. They also list potential limitations of the models.

Simics Programming Guide

This guide explains how to extend Simics by creating new devices and new commands. It gives a broad overview of how to work with modules and how to develop new classes and objects that fit in the Simics environment. It is only available when the DML add-on package has been installed.

DML Tutorial

This tutorial will give you a gentle and practical introduction to the Device Modeling Language (DML), guiding you through the creation of a simple device. It is only available when the DML add-on package has been installed.

DML Reference Manual

This manual provides a complete reference of DML used for developing new devices with Simics. It is only available when the DML add-on package has been installed.

Simics Reference Manual

This manual provides complete information on all commands, modules, classes and haps implemented by Simics as well as the functions and data types defined in the Simics API.

Simics Micro-Architectural Interface

This guide describes the cycle-accurate extensions of Simics (Micro-Architecture Interface or MAI) and provides information on how to write your own processor timing models. It is only available when the DML add-on package has been installed.

RELEASENOTES and LIMITATIONS files

These files are located in Simics's main directory (i.e., `[simics]`). They list limitations, changes and improvements on a per-version basis. They are the best source of information on new functionalities and specific bug fixes.

Simics Technical FAQ

This document is available on the Virtutech website at <http://www.simics.net/support>. It answers many questions that come up regularly on the support forums.

Simics Support Forum

The Simics Support Forum is the main support tool for Simics. You can access it at <http://www.simics.net>.

Other Interesting Documents

Simics uses Python as its main script language. A Python tutorial is available at <http://www.python.org/doc/2.4/tut/tut.html>. The complete Python documentation is located at <http://www.python.org/doc/2.4/>.

Chapter 2

Glossary

This is a list of terms that have a special meaning in Simics documentation.

- **callback** — A user-defined function installed so that it will be called from Simics, for example when a *hap* occurs.
- **checkpoint** — The state of simulation, saved as a number of files, that can be loaded to continue simulation at the point the checkpoint was saved.
- **CLI** — See Command Line Interface.
- **Command Line Interface** — The default Simics command-line (user interface). It uses a simple language implemented in Python, and is variously called the Simics “front end” or the “CLI”.
- **component** — A component is typically the smallest hardware unit that can be used when configuring a real machine, and examples include motherboards, PCI cards, hard disks, and backplanes. Components are usually implemented in Simics using several *configuration objects*.
- **configuration** — A configuration is a description of a target architecture, and is loaded into Simics with the *read-configuration* command. Note that a configuration can also include the state of the target, and saved from within Simics using the *write-configuration* command, in which case it provides a portable *checkpoint* facility.
- **configuration object** — Simics’s configuration system is object-oriented: a simulated machine is represented as a set of *components*, which are implemented by *configuration objects*.
- **context** — Each processor has a *current context*, which represents the virtual address space currently visible to code running on the processor.
- **craft** — **C**ompressed **R**andom **A**ccess **F**ile **F**ormat, used to save raw data for objects, such as disk dumps and target memory contents. A separate utility allows you to compress input data to Simics, such as target disk dumps, in a format that Simics can use directly.

- **cycle** — The smallest unit of time in Simics. When using Simics in its default mode, the cycle count is usually the same as the *step* count, but this can be changed in various ways to improve the fidelity of the simulation.
- **device** — A module modeling a hardware device.
- **event** — A Simics event occurs at some predefined point of simulated time. Time can be specified either as a number of steps on a simulated processor, or a number of simulated clock cycles.
- **extension** — A module which is not a device, but adds features to Simics; e.g., a statistics collection module.
- **hap** — Defined simulation or simulator state changes that may trigger callback functions.
- **Hindsight™** — The technology utilized by Simics to achieve reverse execution.
- **host addresses** — Logical memory addresses on the host machine; i.e., addresses in the virtual memory space that the simulator itself is running in.
- **host** — The machine the simulator (Simics) is running on.
- **logical addresses** — Memory addresses corresponding to virtual or logical addresses on the target machine. These are typically translated by a memory management unit (MMU) to physical addresses.
- **memory hierarchy** — A user defined module that simulates caches and timing of memory operations. Interfaces to Simics using a memory transaction data structure.
- **module** — A dynamically linked library or a script that interfaces to Simics and extends the functionality of the simulator. A module is either an *extension* or a *device*.
- **object** — See Configuration object.
- **physical addresses** — Memory addresses corresponding to physical/real addresses on the target machine; i.e., the actual address bits put out on the memory bus.
- **Python** — An object oriented script language. See <http://www.python.org> for more info.
- **reverse execution** — Execution of a simulation backwards in simulated time.
- **Simics console** — The Simics console is a text console where you can issue commands to Simics, and where Simics will display status information, log messages, and print-out from issued commands. The Simics console is available in all Simics user interfaces but in slightly different versions.
- **SimicsFS** — A “magic” feature allowing simple accesses from a simulated system to the host’s real file system. Pre-configured disk dumps use the `/host` mount point for the magic device. This is presently only supported when running Solaris or Linux on the target machine.

- **Simics object** — See Configuration object.
- **STC** — Simulator Translation Cache. A mechanism in the simulator to improve simulator performance. Among other things, it filters uninteresting memory accesses from the memory hierarchy.
- **step** — An issued instruction that completes or causes an exception, or an external interrupt.
- **system level instruction set simulation** — The effect of every single instruction is simulated both on user and supervisor level. At any instruction, the simulation can be stopped and state can be inspected (and changed).
- **target** — The simulated machine.
- **virtual machine** — The simulated machine.
- **workspace** — The primary area where the user's files (scripts, checkpoints, source code, etc.) are placed.

Part II

Simulating with Simics

Chapter 3

Introduction

Simics is an *efficient, instrumented, system level instruction set simulator*.

- Whereas an *emulator* is focused on executing a program as quickly and accurately as possible, a **simulator**, in addition, is designed from the ground up to gather information on the execution and in general be a flexible tool. Simics is not a monolithic program, it is a complete platform on which simulation-based tools can be built.
- **efficient** means that Simics is designed to run simulations very fast. Often, a Simics simulation will run as fast as the real hardware, or even faster. In fact, in its class of tools, Simics is the fastest simulator ever implemented.
- **instrumented** means that Simics was designed not to run just the target system programs, but to gather a great deal of information during runtime. Many simulators are simply fast emulators with instrumentation added as an afterthought, making the tool slow down considerably when running with high levels of information-gathering enabled.

Simics, by contrast, is specifically designed to achieve good performance even with a high level of instrumentation and very large workloads. Simics provides a variety of statistics in its default configuration and allows various add-ons to be developed by power users and plugged into the simulator.

- Simics is **system level**, meaning that it models a target computer at the level that an operating system acts. Thus, Simics models the binary interfaces to buses, interrupt controllers, disks, video memory, etc. This means that Simics can run anything that the target system can, i.e., arbitrary workloads. Simics can boot unmodified operating system kernels from raw disk dumps.
- **instruction set simulator** means that Simics models the target system at the level of individual instructions, executing them one at a time. This is the lowest level of the hardware that software has ready access to. Simulating at this level allows Simics to be *system level*, yet still permits an *efficient* design.

Simics fully virtualizes the target computer, allowing simulation of multiprocessor systems as well as a cluster of independent systems, and even networks, regardless of the

simulator host type. The virtualization also allows Simics to be *cross platform*. For instance, Simics/SunFire can run on a Linux/x86 system, thus simulating a 64-bit big-endian system on a 32-bit little endian host.

The end uses for Simics include program analysis, computer architecture research, and kernel debugging. The analysis support includes code profiling and memory hierarchy simulation (i.e., cache hierarchies). Debugging support includes a wide variety of breakpoint types. The support for system-level simulation allows operating system code to be developed and analyzed.

3.1 Hosts and Targets

Two fundamental terms used throughout the Simics manuals are *host* and *target*:

- *host* defines the computer on which you are running Simics. It can also refer to the architecture (x86) or the model (Sun SunFire) of computer that runs the simulator, or to the host *platform* (combination of an architecture and an operating system, like x86-linux).
- *target* refers to the computer simulated by Simics. It can also refer to the architecture or the model of computer simulated.

You can find a full list of terms used in the glossary (Chapter 2).

The standard host platforms for Simics are:

Linux/x86

Built for Red Hat Linux 7.3. Simics also runs on many other Linux distributions.

Linux/AMD64

Built for SuSE Linux 9.0. Simics also runs on many other AMD64 Linux distributions.

Solaris/UltraSPARC 64-bit

Built for Solaris 8. Simics also runs on Solaris 9 and 10.

Windows/x86

Built for Windows 2000. Simics also runs on newer versions of Windows.

3.2 Host Recommendations

Requirements on memory and disk sizes depends on the workload, but at least 512MB RAM and several GB of free disk is recommended. In general, it helps to have at least as much memory as what is being used on the simulated machine to avoid unnecessary swapping. A fast disk system speeds up loading and saving of disk images and checkpoints.

If you are running Simics on a laptop and want to reduce the power consumption, you can try to use the **enable-real-time-mode** command to prevent Simics from simulating faster than real time (see the *Simics Reference Manual* for more information on this command). To further reduce the CPU usage of Simics, you can pass an argument to **enable-real-time-mode** that is less than 100 (in percent of real time).

3.3 Simics Targets

The following section lists all Simics targets that are publicly available, i.e., available for evaluation and academic use. Not all processor models are available in the public distribution; contact Virtutech for information on other models. Refer to the Simics target guide corresponding to a specific target to get a full list of all the processor models available.

3.3.1 AlphaPC 164LX

This target models the Alpha 21164 (also known as EV5) implementation of the Alpha architecture. It includes device models and configurations for systems based on the 21174 (Pyxis) chipset, similar to the AlphaPC 164LX.

This target is capable of booting Red Hat Linux 6.0 and 6.2. It has been tested with Linux Miniloader (MILO).

3.3.2 ARM SA1110

This target models a generic ARMv5 processor with minimal implementations of the devices from the Intel StrongARM processor that are needed to boot a minimal Linux configuration.

3.3.3 Ebony

This target models a PPC-based Ebony card with a PPC440GP 32-bits processor. Other processor models are available, like the 405GP and 440GX. It boots Linux 2.4 and VxWorks. This target supports Hindsight.

3.3.4 Fiesta

This target simulates the Sun Blade 1500 workstation from Sun Microsystems, running Solaris. The processor modeled is UltraSPARC IIIi. A variety of PCI based devices are supported, such as SCSI and Fibre Channel. This target supports Hindsight.

The *Micro Architectural Interface* (MAI) is available for the Fiesta target (see [chapter 17](#) and the *Simics Micro-Architectural Interface* for more information).

3.3.5 IA-64 460GX

This target models the Intel Itanium and Itanium2 processors. It includes device models to simulate a complete system based on the 460GX chipset, and is capable of booting Linux 2.4.

3.3.6 Malta/MIPS4kc

This target models the 32 bit MIPS-4Kc processor with a limited set of devices from the MIPS Malta reference board. Only the devices needed in order to boot Linux have been implemented. This targets boots Linux 2.4.

3.3.7 PM/PPC

This target models a 32 bit PowerPC 750 processor on an Artesyn's PM/PPC card with PCI support. It boots Linux 2.4. Other processors available include the 74xx family. This target supports Hindsight.

3.3.8 Simple PPC64

This target models a 64-bit PowerPC 970FX processor along with a few standard devices necessary to boot Linux. It boots Linux 2.4. This target supports Hindsight.

3.3.9 Serengeti

This target simulates the Sun Fire 3800–6800 server series from Sun Microsystems, running Solaris. The processors modeled are UltraSPARC III, UltraSPARC III Cu and UltraSPARC IV. A variety of PCI based devices are supported, such as SCSI and Fibre-Channel. This target supports Hindsight.

The *Micro Architectural Interface* (MAI) is available for the Serengeti target (see chapter 17 and the *Simics Micro-Architectural Interface* for more information).

3.3.10 SunFire

This target simulates the Sun Enterprise 3000–6500 server series from Sun Microsystems, running Solaris or Linux. The processor modeled is UltraSPARC II. A variety of SBus and PCI based devices are supported, such as SCSI, Fibre Channel, and graphics cards. This target supports Hindsight.

The *Micro Architectural Interface* (MAI) is available for the SunFire target (see chapter 17 and the *Simics Micro-Architectural Interface* for more information).

3.3.11 x86 440BX

This target simulates various x86 compatible processors, ranging from 486sx to Pentium 4 and AMD64 processors, and it is capable of booting several Linux versions, Windows NT 4.0, 2000 and XP in both single-processor and multi-processors (SMP) configurations. It includes standard PC devices, such as graphic devices, north and south bridges, floppy and hard disks.

The *Micro Architectural Interface* (MAI) is available for this target (see chapter 17 and the *Simics Micro-Architectural Interface* for more information).

3.4 Simics Version Number

Each release of Simics has a unique three digit version number, e.g. 2.2.7. A change in the first or second number is defined as a major release. Changing the third number is a minor release. Major releases are generally distributed to provide significant product improvements. Minor releases are generally distributed when adding new simulation models or for error corrections.

3.5. *Simics Compatibility*

A release with an odd second number or with a letter in the version e.g. 2.3.2 or 3.1b.0 is a development release and is not subject to any support commitment. Development releases are generally distributed to provide early access to new functionality.

3.5 Simics Compatibility

In general, Simics strives for compatibility between major releases. However, different parts of Simics will provide different level of compatibility.

The following interfaces are backward compatible between major releases. Any exceptions is listed in the release notes:

- Configurations and checkpoints
- Simics API
- CLI commands

The following parts may change in major release and warnings may not be listed in the release notes:

- Simics file structure
- Simics library source code

The following parts are likely to change in major release without any warning:

- Configuration scripts
- Module source code
- Example code
- Simics ABI

Platform support may change at a major release.

For minor releases, the Simics ABI is backward compatible with rare exceptions due to error corrections. Error corrections causing incompatibility are noted in the release notes while backward compatible ABI extensions are not guaranteed to be mentioned. Thus, device models will usually only have to be re-compiled for a new major release and it should be possible to load them directly into a newer minor release. Note that there may be direct dependencies between modules unrelated to the Simics ABI that force two or more modules to be updated simultaneously.

Processor models are not guaranteed to be be movable even between minor releases. Instead a new version of the processor module is supplied with each new minor release.

Communication between Simics simulation processes is only guaranteed to work when running the same version of Simics.

DML code has a version number which is recognized or rejected by the DML compiler. DMLC generates code for the Simics version it is distributed with. However, the user can

write DML code that runs with an older minor version of Simics by avoiding using interfaces introduced in later minor versions.

Target model functionality may have been extended and corrected between minor releases. Target timing may have been changed between minor releases. Important updates and changes will be noted in the release notes.

Chapter 4

First Steps

First Steps is a step-by-step guide describing how to perform some common tasks with Simics. The guide is based on a target computer known as *Ebony*. *Ebony* is a simulated PowerPC 440GP machine running a very small Linux Montavista installation.

4.1 Launch Simulation

The installation of Simics is made to be shared among several users, and therefore you first need to create your own *workspace*. A workspace is a place where you have write permission and can save your own Simics files.

Type the following commands in a terminal, to create a workspace called `simics-workspace` in your home directory. Replace `[simics]` with your Simics installation directory.

```
joe@computer: ~$ [simics]/bin/workspace-setup ~/simics-workspace
Setting up Simics workspace directory: /home/joe/simics-workspace
[..]
joe@computer: ~$ cd simics-workspace
joe@computer: simics-workspace$
```

The preconfigured *Ebony* machines reside under the directory `targets/ebony/`. Launch the `firststeps` configuration in the command-line environment of Simics.

```
joe@computer: simics-workspace$ ./simics targets/ebony/ebony-linux-firststeps.simics
```

Simulation starts in suspended mode, and the Simics prompt will appear, waiting for you to give further commands.

In this guide, interaction with the Simics console will be presented in monospace font. User input will be presented in **bold font**.

```
simics> this is something that you should type
This is output from Simics
```

Ebony's screen will be shown in an additional window. It is initially empty before the machine is started.

4.2 Running the Simulation

You start the simulation with the command **continue**.

```
[...]
simics> continue
```

As the simulation is running, you will see the boot process in the terminal connected to *Ebony*.

At any time you can pause the simulation by pressing control-C. Simics will stop between two instructions and bring you back to the prompt.

```
[press control-C]
[cpu0] v:0xc0003d0c p:0x000003d0c beq- cr4,0xc0003d1c
simics> continue
```

After some time Linux will be up and running and you arrive at the command line prompt. See figure 4.1.

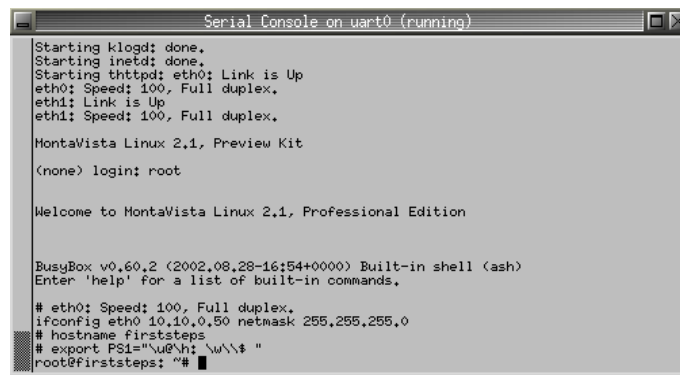


Figure 4.1: A booted Montavista Linux

You can now try typing a few commands on the prompt:

```
root@firststeps: ~# pwd
/root
root@firststeps: ~# ls /
LICENSE      dev          host         lost+found  proc        tmp
bin          etc          lib          mnt         root        usr
boot        home        linuxrc      opt         sbin        var
```

4.3. Checkpointing

```
root@firststeps: ~#
```

Note: If the terminal console does not respond, make sure that the simulation is actually running.

4.3 Checkpointing

In order to avoid booting First Steps every time, we use the facility known as *configuration checkpointing* or simply *checkpointing*. This enables Simics to save the entire state of the simulation to disk. The files are saved in a portable format and can be loaded at a later time.

To write a checkpoint, use the **write-configuration** command. The command expects a filename as argument. Name the checkpoint `after_boot.conf`.

```
[control-C]
simics> write-configuration after_boot.conf
simics>
```

The state is now saved, and you can safely terminate the simulation.

```
simics> quit
joe@computer: simics-workspace$
```

All that remains of the simulation is the checkpoint we just created. To load the checkpoint, use the `-c` flag when starting Simics.

```
joe@computer: ebony$ ./simics -c after_boot.conf
[...]
simics>
```

The terminal contents is also restored and you should end up in the same state as before. Remember to resume the simulation (by entering **continue**) before typing more commands at the console.

You can read more about configuration and checkpointing in [chapter 6](#).

4.4 Hindsight

Note: Hindsight requires a special license. This section can be skipped if Hindsight is not available.

Hindsight enables the simulation to run backwards in time. To enable Hindsight, we have to set an initial *time bookmark*. Time bookmarks are set with the **set-bookmark** command.

```
simics> set-bookmark booted
simics> c
```

Note: Typing **c** is a shorthand for **continue**.

To demonstrate the possibilities Hindsight gives, we will accidentally remove an important file on the simulated system. Enter the following commands in the simulated system's terminal:

```
root@firststeps: ~# rm /bin/ls
root@firststeps: ~# ls /
ls: No such file or directory
```

The program **ls** has been removed. You can no longer list the contents of a directory. Let us use Hindsight to recover **ls**.

ptime is a useful command that shows the number of executed instructions and the current simulated time. We will use it to show that the time have advanced backwards.

```
[press control-C]
simics> ptime
processor          steps          cycles    time [s]
cpu0              16667617210      16667617210    166.676
simics>
```

The **skip-to** command can be used to quickly jump to a previous point in time. We will use the bookmark we created before as our time-travel destination. Note that it is not possible to reverse past the first bookmark.

```
simics> skip-to bookmark = booted
simics> ptime
processor          steps          cycles    time [s]
cpu0              13586370338      13586370338    135.864
simics>
```

The system is now in the state it was before the file was erased. Now, we run forward again.

```
simics> c
```

When you type something in the terminal, you will notice that it does not respond any longer! Instead the same commands as before will be replayed.

This behavior is intentional, and keeps the deterministic property of the simulation, which is invaluable when debugging. Keystrokes, network traffic and any other input is replayed until the last known time is reached.

4.5. Getting Files into a Simulated System

In our example, this is not what we want. To erase all knowledge about the future, run the **clear-recorder** command.

```
[press control-C]
simics> skip-to bookmark = booted
simics> clear-recorder
Replay of recorded input finished; Simics is now running normally
simics> c
```

Resume the simulation and enter the following command:

```
root@firststeps: ~# ls /
LICENSE      dev          host          lost+found   proc         tmp
bin          etc          lib           mnt          root         usr
boot         home         linuxrc       opt          sbin         var
root@firststeps: ~#
```

ls works again!

You can read more about Hindsight in [chapter 22](#).

4.5 Getting Files into a Simulated System

A full systems simulation can be run completely isolated, but Simics also provides a way for the simulated system to access files that are located on the host, i.e, your real computer.

SimicsFS is a Linux kernel filesystem module that talks to a simulated device, namely a corresponding Simics module. Our *firststeps* machine is prepared with **simicsfs** support; just mount `/host` to access it.

```
root@firststeps: ~# mount /host
[simicsfs] mounted
root@firststeps: ~# ls /host
bin          etc          lost+found   net          root         sys
boot         home         media        opt          sbin         tmp
[...]
root@firststeps: ~#
```

Note: Support for SimicsFS have to be installed in the target system, which means it have to be modified. Modules exists for the Linux 2.4 and 2.6 series. On other systems a module has to be written in order to get SimicsFS to work.

As write support is experimental, SimicsFS is mounted read-only by default. To be able to transfer files from the simulation to the host, mount it read-write.

```
root@firststeps: ~# mount /host -o remount,rw
```

```

root@firststeps: ~# cp /proc/cpuinfo /host/tmp/cpuinfo
root@firststeps: ~# umount /host
root@firststeps: ~#

```

Now, we can read the file on our host machine.

```

[press control-C]
simics> !cat /tmp/cpuinfo
processor      : 0
cpu           : 440GP Rev. C
revision      : 4.129 (pvr 4012 0481)
bogomips      : 799.53
vendor        : IBM
machine       : Ebony
simics>

```

Note: The `!` command will interpret the rest of the command line as if it was given in a shell.

Sometimes it is desirable to limit the accessible files to a certain directory. This is done with the `<hostfs>.root` command. Set the root to the Simics installation directory.

```

[press control-C]
simics> hfs0.root sim->simics_base
simics> c

```

Now transfer the file to be used in next section into the simulated system.

```

root@firststeps: ~# mount /host
[simicsfs] mounted
root@firststeps: ~# cp /host/targets/ebony/images/debug_example .
root@firststeps: ~#

```

You can read more about SimicsFS/`hostfs` and other ways to transfer files in [chapter 7](#).

4.6 Debugging

Note: Some steps in this section require Hindsight. However, you will grasp the basic debugging commands even if Hindsight is not available.

This section demonstrates some source-level debugging facilities that Simics provides. Your Simics distribution contains an example code snippet called `debug_example.c`.

4.6. Debugging

Copy the source file and corresponding compiled executable into your workspace using the following commands.

```
simics> !cp [simics]/targets/ebony/debug_example.c .
simics> !cp [simics]/targets/ebony/images/debug_example .
simics> c
```

Replace `[simics]` with the Simics installation directory.

We recommend opening the file `debug_example.c` in an editor of your choice to easier follow the debugging example. This file contains the code that we are going to debug. The program is supposed to print some information about the users on a system.

In the previous section we copied the executable into the simulated system by using *SimicsFS*. Now, run that program by starting it in the simulated terminal.

```
root@firststeps: ~# ./debug_example
[...]
Got segmentation fault!
root@firststeps: ~#
```

This output indicates that our program crashed. Let us use Simics features to debug it.

Simics needs to know the mapping between addresses and line numbers, and this information is stored in the executable. The **symtable** module in Simics contains commands related to symbolic debugging.

```
[press control-C]
simics> new-symtable file = debug_example
Created symbol table 'debug_example'
[symtable] Symbols loaded at 0x10000000
ABI for debug_example is ppc-elf-32
debug_example set for context primary_context
simics>
```

Note: Remember to write **file =**, otherwise you will create an empty **symtable** named `debug_example`. You can add files to an existing **symtable** with the `<symtable>.add-symbol-file` command.

To help debugging your programs, we have introduced *magic instructions*. These are instructions that have no side-effects on a real machine, but can be programmed to do things when run inside Simics, for example stop the simulation.

The debug example code contains such an instruction in the beginning of the `main` function. Enable break on magic instruction by the command **magic-break-enable**:

```
simics> magic-break-enable
simics> c
```

Now rerun `debug_example`:

```
root@firststeps: ~# ./debug_example
```

Simics will stop at the magic instruction and show the corresponding source code line and assembly opcode.

```
[cpu0] v:0x10000634 p:0x00737b634 magic instruction (or r0,r0,r0)
main (argc=1, argv=0x7ffffe34) at /tmp/debug_example.c:73
73             MAGIC_BREAKPOINT;
simics>
```

As you can see, this magic instruction is effectively a no-operation, which means that the simulation will run as usual on real hardware, or when magic breakpoints are disabled in Simics.

Now let us find the cause of the segmentation fault. Place a breakpoint in the *`sigsegv_handler()`* function. The *`sigsegv_handler()`* function is called when the program receives the segmentation fault and will allow the program to gracefully exit.

```
simics> break (sym sigsegv_handler)
Breakpoint 1 set on address 0x10000520 with access mode 'x'
1
simics>
```

Resume the simulation. It will stop at the signal handler, and by giving the **stack-trace** command, you can also see the chain of function calls leading up to this point. This list gives you useful hints about where the crash occurred.

```
simics> c
Code breakpoint 1 reached.
[cpu0] v:0x10000520 p:0x00737b520 stwu r1,-32(r1)
sigsegv_handler (sig=0) at /tmp/debug_example.c:35
35     {
simics> stack-trace
#0 0x10000520 in sigsegv_handler (sig=0) at /tmp/debug_example.c:35
#1 0x7ffff3d8 in ?? ()
#2 0xff06a44 in ?? ()
#3 0xff0ff60 in ?? ()
#4 0x100006ac in main (argc=1, argv=0x7ffffe34) at /tmp/debug_example.c:82
#5 0xfed5fdc in ?? ()
#6 0x0 in ?? ()
simics>
```

4.6. Debugging

Simics prints question marks when no symbol could be found at the address. This can either be a bogus address or a function inside the standard library, to which no symbols have been loaded.

A few frames down you have the *main()* function, which caused the crash. Now we run the simulation backward into that function. **reverse-step-line** will run backwards until the previous known source line is reached.

```
simics> reverse-step-line

[cpu0] v:0x100006a8 p:0x00737b6a8 bl 0x10010c54
main (argc=1, argv=0x7ffffe34) at /tmp/debug_example.c:82
82                                printf("Type: %s\n", user.type);
simics>
```

This line cause the crash. Let us examine what `user.type` contains:

```
simics> psym user.type
(char *) 0xa94 (unreadable)
simics> psym user
{name = 0x7ffffdc0 "shutdown", type = (char *) 0xa94 (unreadable)}
simics>
```

As you can see, the `type` member points to an unreadable address, which caused the crash. Where does this pointer come from? What we want to do is to find where the last write to this pointer occurred.

Using Hindsight, we can first set a write-access breakpoint on the memory of interest, and run backward (using **reverse**) until the breakpoint is reached. After some time will find the place where the write takes place.

```
simics> break -w (sym "&user.type") (sym "sizeof user.type")
Breakpoint 2 set on address 0x7ffffdc8, length 4 with access mode 'w'
2
simics> reverse
Breakpoint on write to address 0x7ffffdc8 in primary_context.
Completing instruction @ 0xff365e0 on cpu0.
[cpu0] v:0x0ff365e4 p:0x0075275e4 beqlr
simics>
```

Now, examine the stack trace:

```
simics> stack-trace
#0 0xff365e4 in ?? ()
#1 0x100005d8 in read_developer (p=0x7ffffdc0, f=0x10010ca0)
    at /tmp/debug_example.c:60
#2 0x10000674 in main (argc=1, argv=0x7ffffe34) at /tmp/debug_example.c:80
```

```
#3 0xfed5fdc in ?? ()
#4 0x0 in ?? ()
simics>
```

In the stack trace, you will see that a call from `read_developer()` have caused the crash. Switch to that frame and display the code being run.

```
simics> frame 1
#1 0x100005d8 in read_developer (p=0x7ffffdc0, f=0x10010ca0)
    at /tmp/debug_example.c:60
simics> list read_developer 15
48  {
49      char line[100], *colon;
50
51      if (fgets(line, 100, f) == NULL)
52          return 0;          /* end of file */
53
54      /* Type is always developer */
55      p->type = "developer";
56
57      /* Everything until the first colon is the name */
58      colon = strchr(line, ':');
59      *colon = '\0';
60      strcpy(p->name, line);
61      return 1;
62  }
simics>
```

On line 60 you can see that while the *name* field was filled in using `strcpy`, our failing pointer was accidentally overwritten (remember that the breakpoint was placed on the *type* member). If you write **psym line** you will see that the string copied is "shutdown". A look into the declaration of `struct person` shows that the *name* field is only 8 bytes big, and hence has no space for the trailing null byte.

Check the contents of *p* after and before the actual write to verify it is overwritten.

```
simics> psym "*p"
{name = 0x7ffffdc0 "shutdown", type = (char *) 0xa94 (unreadable)}
simics> reverse-step-instruction
Completing instruction @ 0xff365e0 on cpu0.
Breakpoint on write to address 0x7ffffdc8 in primary_context.
[cpu0] v:0x0ff365e0 p:0x0075275e0 stb r0,4(r5)
simics> frame 1; psym "*p"
#1 0x100005d8 in read_developer (p=0x7ffffdc0, f=0x10010ca0)
    at /tmp/debug_example.c:60
{name = 0x7ffffdc0 "shutdown\020", type = (char *) 0x10000a94 "developer"}
```

4.7. Tracing

```
simics>
```

To clean up after our debug session, we must remove the breakpoints that we have set. They are identified with a number and can be shown using **list-breakpoints**.

```
simics> delete 1
simics> delete 2
simics> magic-break-disable
simics>
```

You can read more about debugging in chapter 12. Magic instructions are described in section 12.1.7.

4.7 Tracing

Tracing is a way to observe what is going on during the simulation. This section describes how to trace memory accesses, I/O accesses, control register writes, and exceptions in Simics.

The tracing facility provided by the **trace** module will display all memory accesses, both instruction fetches and data accesses.

First, launch the `ebony-linux-firststeps.simics` configuration, but not boot it. Second, create a tracer:

```
simics> new-tracer
Trace object 'trace0' created. Enable tracing with 'trace0.start'.
simics>
```

Now we are going to trace a few of instructions executed when booting *Ebony*. We execute 300 instructions without tracing first to reach a sequence of instructions that includes memory accesses:

```
simics> continue 300
[cpu0] v:0xffffffff160 p:0xffffffff160 tlbwe r1,r4,1
simics> trace0.start
Tracing enabled. Writing text output to standard output.
simics> continue 6
inst: [ 1] CPU 0 <v:0xffffffff160> [...] 7c240fa4 tlbwe r1,r4,1
inst: [ 2] CPU 0 <v:0xffffffff164> [...] 7c4417a4 tlbwe r2,r4,2
inst: [ 3] CPU 0 <v:0xffffffff168> [...] 38840001 addi r4,r4,1
inst: [ 4] CPU 0 <v:0xffffffff16c> [...] 4200ffdc bdnz+ 0xffffffff148
inst: [ 5] CPU 0 <v:0xffffffff148> [...] 84050004 lwzu r0,4(r5)
data: [ 1] CPU 0 <v:0xffffffff1b8> [...] Read 1 bytes 0xc0
data: [ 2] CPU 0 <v:0xffffffff1b9> [...] Read 1 bytes 0x0
data: [ 3] CPU 0 <v:0xffffffff1ba> [...] Read 1 bytes 0x12
data: [ 4] CPU 0 <v:0xffffffff1bb> [...] Read 1 bytes 0x10
```

```
inst: [    6] CPU  0 <v:0xfffff14c> [...] 2c000000 cmpwi r0,0
[cpu0] v:0xfffff150 p:0x1fffff150 beq- 0xfffff170
simics>
```

Lines beginning with `inst:` are executed instructions. Each line contains the address (both virtual and physical) and the instruction itself, in both hexadecimal form and the mnemonic.

Lines beginning with `data:` indicate that some instructions are performing memory operations. Each line contains the operation address (again, both virtual and physical), the type of operation (read or write), the size and the value.

Note: In the trace you can see one four-byte memory read that is split into four single-byte reads. This reflects how Simics models the accesses to the flash memory it reads from.

It is also possible to only trace accesses to a certain device. This is done with the **trace-io** command. In this example we are looking at in the interaction with the UART device.

```
simics> trace0.stop
Tracing disabled
simics> trace-io uart0
simics> continue 30_000
[cpu0 -> uart0] Write: 0x140000203 1 0x80
[cpu0 -> uart0] Write: 0x140000200 1 0x48
[cpu0 -> uart0] Write: 0x140000201 1 0
[cpu0 -> uart0] Write: 0x140000203 1 0x3
[cpu0 -> uart0] Write: 0x140000202 1 0
[cpu0 -> uart0] Write: 0x140000204 1 0
[cpu0 -> uart0] Read: 0x140000205 1 0x60
[cpu0 -> uart0] Read: 0x140000200 1 0
[cpu0 -> uart0] Write: 0x140000207 1 0
[cpu0 -> uart0] Write: 0x140000201 1 0
[cpu0 -> uart0] Read: 0x140000205 1 0x60
[cpu0 -> uart0] Write: 0x140000200 1 0xd
[cpu0 -> uart0] Read: 0x140000205 1 0
[cpu0] v:0xffff8a634 p:0x1fff8a634 bl 0xffff8a608
simics>
```

Note: You can use underscores anywhere in numbers to make them more readable. The underscores have no other meaning and are ignored when the number is read.

trace-cr turns on tracing of changes in the processor's control registers.

```
simics> untrace-io uart0
simics> trace-cr -all
```


4.7. Tracing

```
simics> continue 6_500_000
[cpu0] tcr <- 0
[cpu0] dec <- 0
[cpu0] decar <- 0
[cpu0] tsr <- 0x8000000
[cpu0] decar <- 0x51615
[cpu0] dec <- 0x51615
[cpu0] tcr <- 0x4400000
[cpu0] ivpr <- 0
[cpu0] msr <- 0x29000
[cpu0] msr <- 0
[...]
[cpu0] [cpu0] v:0x07fd8634 p:0x007fd8634 bl 0x7fd8608
simics>
```

We can single-step with the *-r* flag, to see what registers each instruction changes.

```
simics> untrace-cr -all
simics> step-instruction -r 10
[cpu0] v:0x07fd8608 p:0x007fd8608 mftbu r3
[cpu0] v:0x07fd860c p:0x007fd860c mftbl r4
      r4 <- 6529678
[cpu0] v:0x07fd8610 p:0x007fd8610 mftbu r5
[cpu0] v:0x07fd8614 p:0x007fd8614 cmpw r3,r5
[cpu0] v:0x07fd8618 p:0x007fd8618 bne+ 0x7fd8608
[cpu0] v:0x07fd861c p:0x007fd861c blr
[cpu0] v:0x07fd8638 p:0x007fd8638 subfc r4,r4,r7
      r4 <- 1384
[cpu0] v:0x07fd863c p:0x007fd863c subfe. r3,r3,r6
[cpu0] v:0x07fd8640 p:0x007fd8640 bge+ 0x7fd8634
[cpu0] v:0x07fd8634 p:0x007fd8634 bl 0x7fd8608
simics>
```

Simics can also monitor exceptions. Here we will trace all system calls.

```
simics> trace-exception System_call
simics> c
[cpu0] (@ cycle 203963506) Exception 8: System_call
[cpu0] (@ cycle 205608905) Exception 8: System_call
[cpu0] (@ cycle 205617423) Exception 8: System_call
[...]
[control-C]
[cpu0] v:0xc01bc12c p:0x0001bc12c addi r10,r10,1
simics> untrace-exception -all
```

```
simics>
```

Note:

There are variants of **trace-io**, **trace-cr** and **trace-exception** that will stop the simulation when respective event occur. These commands begin with **break-**.

Why do we need to create an object to trace instructions? This is because the tracing should be easy to customize to your needs. By changing some attributes, it is possible to choose what to trace on, and control the output format. You can also build your own tracer based on the source of this tracer. All attributes are described in the Reference Manual.

4.8 Scripting

The Simics *command line interface* (CLI) has some built-in scripting capabilities. When that is not enough, Python can be used instead.

We will use a **trace** object as our example object:

```
simics> new-tracer
Trace object 'trace0' created. Enable tracing with 'trace0.start'.
simics>
```

There are two *commands* available for this object: **<trace>.start** and **<trace>.stop**.

For example, to start tracing:

```
simics> trace0.start
Tracing enabled. Writing text output to standard output.
simics>
```

It is also possible to access an object's *attributes* using CLI. The state of an object is contained in its attributes.

```
simics> trace0->classname
base-trace-mem-hier
simics> trace0->enabled
1
simics>
```

Variables in CLI are prefixed with \$, and can hold a string, a number, or an object reference. In the following example the variable *my_tracer* references our **trace0** object (i.e. it is *not* a copy).

```
simics> $my_tracer = trace0
simics> $my_tracer->enabled
1
```

4.9. Simple Virtual Network

```
simics>
```

It is also possible to access the tracer from Python. All lines begin with a @ are evaluated as a Python statement.

```
simics> @trace_obj = SIM_get_object("trace0")
simics> @trace_obj
<the base-trace-mem-hier 'trace0'>
simics> @trace_obj.enabled
1
simics>
```

The Simics API is directly accessible from Python. The script below counts the number of instructions that are executed until the register `msr` is modified. It imitates the functionality of `break-cr msr`.

```
simics> @start_cycle = SIM_cycle_count(conf.cpu0)
simics> @msr = conf.cpu0.msr
simics> @while conf.cpu0.msr == msr: SIM_continue(1)
[...]
simics> @end_cycle = SIM_cycle_count(conf.cpu0)
simics> @print "Executed", end_cycle - start_cycle, "instructions"
Executed 56 instructions
simics>
```

After you enter `@while conf.cpu0.msr [...] command`, the simulation starts, and continues until the `msr` register is modified. When that happens (which should not take any noticeable time), the simulation stops and the rest of the commands can be entered.

You can read more about scripting in chapter 8. The full description of the Simics API is available in the Reference Manual.

4.9 Simple Virtual Network

Simics can simulate several machines simultaneously. These machines can be connected together using a simulated Ethernet link.

This chapter uses a different launch configuration than the rest of the First Steps chapters. Before continuing, launch the `ebony-linux-firststeps-multi.simics` configuration.

```
joe@computer: simics-workspace$ ./simics targets/ebony/ebony-linux-firststeps-multi.simics
```

This configuration contains two similar machines, with their own CPU. Simics will run the machines synchronized, by executing some instructions on one CPU and then switch to the other.

Boot the two simulated machines by starting the simulation.

```
simics> c
```

One of them will be assigned the Internet Protocol (IP) address 10.10.0.50 and the other one 10.10.0.51. If you try to ping from one of the machines to the other, it will fail. This is because the two machines have not yet been connected to the simulated network.

An Ethernet link is simulated using **ethernet-link** module. You can connect any number of simulated network cards to the link.

[control-C]

```
simics> new-ethernet-link
[ethlink0 info] Adjusting latency to 1e-05 s [...]
Created ethernet-link ethlink0
simics>
```

Next, we must connect the simulated network cards to this link:

```
simics> ebony0_emac0.connect ethlink0
simics> ebony1_emac0.connect ethlink0
simics> ethlink0.info
Information about ethlink0 [class ethernet-link]
=====

                Latency : 10 us
            Distribution : local
                Filtering : enabled

Devices:
    Local devices : <0:0> ebony0_emac0, <1:1> ebony1_emac0
    Remote devices : none

Real network connection:
                Connected : No

simics> c
```

Now it should be possible to ping between the two simulated machines. Enter the following commands in the first machine.

```
root@firststeps: ~# ping -c 2 10.10.0.51
PING 10.10.0.51 (10.10.0.51): 56 data bytes
64 bytes from 10.10.0.51: icmp_seq=0 ttl=64 time=0.0 ms
64 bytes from 10.10.0.51: icmp_seq=1 ttl=64 time=0.0 ms

--- 10.10.0.51 ping statistics ---
2 packets transmitted, 2 packets received, 0% packet loss
```

4.10. Connect to a Real Network

```
round-trip min/avg/max = 0.0/0.0/0.0 ms
```

```
root@firststeps: ~#
```

You can read more about network simulation in chapter 9.

4.10 Connect to a Real Network

A simulation can be connected to a real network. By doing this, simulated computers and real computers are able to communicate with each other.

Note: If you experience timing problems, for example TCP timeouts, when using real network, the simulation is running too fast. In these cases, slow it down by using the **enable-real-time-mode** command.

Before following the steps in this example, launch a new `ebony-linux-firststeps.simics`. Don't boot it yet.

Connecting a simulated machine to a real network is done with one command:

```
simics> connect-real-network 10.10.0.50
No ethernet-link found, creating 'ethlink0'.
No service-node found, creating 'ethlink0_sn0' with IP '10.10.0.1'.
Connecting device 'emac1' to 'ethlink0'
Connecting device 'emac0' to 'ethlink0'
NAPT enabled with gateway 10.10.0.1 on link ethlink0.
Host TCP port 4021 -> 10.10.0.50:21 on link ethlink0
Host TCP port 4023 -> 10.10.0.50:23 on link ethlink0
Host TCP port 4080 -> 10.10.0.50:80 on link ethlink0
Real DNS enabled at 10.10.0.1 on link ethlink0.
simics> continue
```

Note: Note that the IP address specified in the **connect-real-network** command is the IP address of the simulated machine

This command will create a new **ethernet-link**, connect it to both the simulated network cards, and to the real network. It will also enable NAPT, *network address port translation*. Finally, it will forward ports 4021, 4023 and 4080 to the simulated machine's telnet, FTP and HTTP ports.

To actually make the real network accessible from the simulated machine, the simulated system must be configured. These commands set up the service node as gateway and domain name server:

```
root@firststeps: ~# route add default gw 10.10.0.1
root@firststeps: ~# echo nameserver 10.10.0.1 > /etc/resolv.conf
```

```
root@firststeps: ~#
```

This tells the simulated system to direct all accesses outside the local network to the gateway at 10.10.0.1 (which is the service node **connect-real-network** created for us).

Now, if your computer is connected to Internet, we can try to telnet to a real computer on Internet. In this example we use `gnu.org`.

```
root@firststeps: ~# telnet gnu.org 80
GET /
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<HTML><HEAD>
<TITLE>301 Moved Permanently</TITLE>
</HEAD><BODY>
<H1>Moved Permanently</H1>
The document has moved <A HREF="http://www.gnu.org/">here</A>.<P>
<HR>
<ADDRESS>Apache/1.3.31 Server at gnu.org Port 80</ADDRESS>
</BODY></HTML>
Connection closed by foreign host.
root@firststeps: ~#
```

Since **connect-real-network** forwards ports to the telnet, FTP and HTTP ports of the simulated machine, it is possible to telnet into the simulated machine, or access its web server from a web browser. To access the web server, enter the address `http://localhost:4080` in a web browser on your real machine.

In an new shell on your computer, you can also try to telnet into the simulated machine:

```
joe@computer: ~$ telnet -l root localhost 4023
Trying 127.0.0.1...
Connected to localhost.localdomain.
Escape character is '^]'.

Welcome to MontaVista Linux 2.1, Professional Edition

BusyBox v0.60.2 (2002.08.28-16:54+0000) Built-in shell (ash)
Enter 'help' for a list of built-in commands.

# exit
Connection closed by foreign host.
joe@computer: ~$
```

You can read more about real network in chapter 10.

4.10. Connect to a Real Network



Figure 4.2: Montavista Linux HTTP server front page

4.10. Connect to a Real Network

Chapter 5

Command-line Interface: Basics

The Simics command-line uses GNU Readline to provide command-line editing and a command history buffer.

Let us go through a list of handy commands for Simics common usage:

Simulation

At the Simics prompt, you can run the simulation with the **continue** (**c**) command, followed by the number of steps you want to run. If no argument is provided, the simulation will run until `control-C` is pressed in the Simics console.

You can also use **stepi** (**si**) to make the simulation progress step by step. You can use the **step-cycle** (**sc**) command to make the simulation progress cycle by cycle. Note that as long as you do not use timing models, **si** and **sc** are equivalent. For more information, see chapter 17.

Simulation State

All processors support the **pregs** command to show the state of the main registers. The argument *-all* will allow you to inspect all the registers instead of the usual working set.

If the processor has floating-point registers, **pfregs** can be used to print them out. By default, the **pregs** command (and other commands dependent on the CPU) uses the current CPU scheduled by the simulation. You can change the current selected CPU by using the **pselect** command. This won't affect the simulation scheduling.

You can inspect and change the contents of a register by using the **read/write-reg** commands. Using **%** is equivalent to the **read-reg** command (i.e., **%pc**, **%eax**, ...). You can inspect and change the contents of the memory by using the commands **get**, **x** and **set**.

To get statistics concerning the current CPU, you can use the **ptime** and **pstats** commands.

Scripts

Simics can execute Python commands directly at the prompt; they just need to be prefixed by the symbol **@**. You can also evaluate Python expressions in the middle of a command by placing them between backquotes:

```

    --- this is Python's print()
simics> @print conf.cpu0.name
cpu0
    --- this is CLI's echo
simics> echo `conf.cpu0.name`
cpu0
simics>

```

Note: '@' will only be recognized as starting a python command if it is the *first* character on the command-line. It will be treated normally otherwise.

Simics can load both Simics scripts (a file containing a list of acceptable commands for the Simics frontend) or Python script with the commands **run-command-file** and **run-python-file**. You can read more about Simics scripting in chapter 8.

Modules

Simics usually handles modules automatically when loading a configuration. You may however want to load a specific module by yourself with the **load-module** command.

The modules currently loaded are available through the **list-modules** command. If a module fails to be loaded, it will be listed by the **list-failed-modules** command, along with an explanation of the problem. You can use the *-v* argument to get the exact error message preventing the module from loading.

Command-Line Interface

You can ask Simics to run command every time it returns at the prompt, using the **display** command. For example:

```

simics> display ptime
display 1: ptime
simics> si
[cpu0] v:0xffffffff00000024 p:0x000007fff0000024  nop
processor                steps                cycles    time [s]
cpu0                      1                      1          0.0
simics> c 100
[cpu0] v:0x000007fff0102190 p:0x000007fff0102190  sub %o2, 8, %o2
processor                steps                cycles    time [s]
cpu0                      101                    101        0.0
simics> undisplay 1
simics>

```

To stop the command from being run every time, use **undisplay *number*** where *number* is the number that was returned by **display** for this specific command (in the example, 1).

The **print** command allows you to print out a value in different ways (binary, hexadecimal, etc.). You can control the way the command-line handles numbers with the following commands: **output-radix** allows you to choose a default base for printing numbers; **digit-grouping** can be used to group digits in a number in a more readable way.

The CLI also supports the following Unix shell-like commands: **cd**, **date**, **dirs**, **echo**, **ls**, **popd**, **pushd**, **pwd**. All of them are performed on the **host** machine, not on the simulated machine.

Chapter 6

Configuration and Checkpointing

Simics includes a *configuration system* used to describe the state of the simulated machines.

Simics's configuration system is object-oriented: a simulated machine is represented as a set of *objects* that interact when the simulated time advances. Typically, processors, memories and devices are modeled as objects. Each object is defined by a number of properties that are called *attributes*. A processor object, for example, will have an attribute called *freq_mhz* to define its clock frequency (in MHz).

Simics's configuration system is flexible: it is possible to create and delete objects dynamically, as well as access the attributes of all objects for reading and writing at any time.

Simics's configuration system allows its elements to be saved so that a complete simulated machine state can be written to a set of files. This set of files is called a *checkpoint*.

This chapter describes the Simics configuration system as well as the different layers built on top of it to handle more specific tasks:

- Checkpoints: how a simulated state is saved and restored via checkpointing;
- Inspection: how the simulated state can be examined and changed during simulation;
- Start Scripts: the components and scripts used to define the initial state of a machine in the examples provided with Simics.

6.1 Basics

As mentioned above, Simics's configuration system is object-oriented. A Simics object is instantiated from a Simics *class*. The core of Simics defines some useful classes, but most of the classes (processors, device models, statistic gathering extensions) are provided by *modules* that are loaded by the simulator when requested.

For example, the **x86-p4** module defines, surprisingly, the **x86-p4** class. Note that a module may define several classes. Since modules advertise the classes they define, Simics can load modules transparently as objects are instantiated.

A class defines *attributes* that represent both the static and dynamic state of the instantiated objects. The static state includes information that doesn't change during the simulation (like a version number in a register) while the dynamic state covers the part of the device that are affected by the simulation (registers, internal state, buffers, etc.).

Let us take the example of an x86-p4 processor and have a closer look:

- We can create an object instantiated from the class **x86-p4**. Let us call it **cpu0**
- The attribute *freq_mhz* can be set to 1500. It defines the processor clock frequency (in MHz)
- The attribute *physical_memory* can be set to a memory space object, such as **phys_mem0**. This attribute points to the object that will answer to the memory accesses coming from the processor.
- etc.

As you noticed, attributes may be of various types. A complete description is available in the next section.

6.2 Checkpointing

Simics's configuration system can save the complete state of a simulation in a portable way. This functionality is known as *checkpointing*, and the set of files that represent the elements of the systems are called a *checkpoint*.

Saving and restoring a checkpoint can be done on the command-line with the **write-configuration** and **read-configuration** commands.

A checkpoint consists of the following files:

- A main file called *configuration file* (usually terminated by *.conf*, but this is only a convention). This file is a text representation of the objects present in the system.
- An optional raw data file (described below), having the same name as the configuration file, with the suffix *.raw* appended.
- Optional image files (described in section 6.2.2), having the same name as the configuration file, with the suffix *.object_image* appended.

Below is a portion of a checkpoint file showing an object. Saved objects are always represented as `OBJECT object-name TYPE class-name { attributes }`. In this case we have an instance of the **AM79C960** class (a 10Mbits ISA Ethernet card with on-board DMA) named **lance0**. The *irq_dev* and the *irq_level* attributes connect the device to a controller that will handle the interrupts it generates. Since this device has on-board DMA, it is also connected to memory with the *memory* attribute.

```
...
}
OBJECT lance0 TYPE AM79C960 {
    mac_address: "10:10:10:10:10:30"
    irq_dev: isa0
    irq_level: 7
```

6.2. Checkpointing

```
        memory: pci_mem0
        ...
    }
OBJECT ... TYPE ... {
    ...
```

Objects are saved in the main checkpoint file in no specific order.

6.2.1 Attributes

The short example of the **lance0** description only uses three types of attribute values: strings, objects, and (signed 64-bit) integers. The possible attribute types are:

string

Strings are enclosed in double quotes, with C-style control characters: "a string\n"

integer

Integers can be in hexadecimal (0xfce2) or signed decimal (-17) notation.

boolean

One of TRUE or FALSE.

floating-point

Specified in decimal (1.0e-2) or hexadecimal (0x5.a21p-32) style, just like in C.

object

The name of a configuration object: cpu0.

raw data

Arbitrary data; typically used to save large dumps of binary information. The data itself is stored in the raw data file. The syntax is [*R length-in-bytes filename file-offset*].

list

Comma-separated list of any attribute values, enclosed in parentheses. Example: ("a string", 4711, (1, 2, 3), cpu0)

dictionary

The format is a comma-separated list of key/value pairs, like in: { "master-cpu" : cpu0, "slave-cpu" : cpu1 }. The key should be a string, integer or object, while the value can be of any attribute type. Dictionaries are typically used to save Python dictionaries in a checkpoint. Keys must be unique, although Simics does not enforce this.

Each attribute belongs to one of the following categories. Note that only attributes of the first two categories are saved in checkpoints.

Required

Required attributes must be set when creating an object. They are saved in checkpoints. If you edit a checkpoint, you should never remove a required attribute—Simics will complain and refuse to load the checkpoint if you do.

Optional

If no other value is provided, optional attributes take their default value when the object is created. They are saved in checkpoints, but if you edit them out they will revert to their default value when the checkpoint is loaded.

Session

Session attributes are only valid during a Simics session. They are not saved in checkpoints. They are usually used for statistics gathering or values that can be computed from the rest of the object state.

Pseudo

Pseudo attributes are not saved in checkpoints and usually contain read-only information that does not change, or that is calculated when the attribute is accessed. Pseudo attributes are in some cases used to trigger state changes in the object when written.

There are two special cases in the attribute checkpointing process. The first one, described above, concerns raw data. Raw data is saved in a separate file that belongs to the checkpoint file set. The main configuration file only contains a pointer to the corresponding data in the raw data file.

The second one concerns *images*.

6.2.2 Images

Simics implements a special class called **image** for objects that potentially need to save a huge amount of state, like memories and disks. An image represents a big amount of raw data using pages and compression to minimize disk usage.

To save space and time, images do not save their entire state every time a checkpoint is written. They can work in two ways:

- Images can save their state *incrementally*. At each checkpoint, an image saves the difference between its current state and the previously saved state (either the previous checkpoint or the initial state). This is the default behavior implemented by Simics. This allows several checkpoints to be saved and restored using the same base image and a series of difference files.
- Images can be used as *read-write* media. In that case the file representing the data is always up to date to the current state. However, this prevents the image from being used in a previously saved checkpoint or initial state, since its contents are modified as the simulation advances.

It is important to understand that when used in incremental mode, images create *dependencies* between checkpoints. A checkpoint can only be loaded if all previous checkpoints are intact.

To re-use the example above, let us have a look at the disk image of the same x86 computer:

...

6.2. Checkpointing

```
}  
OBJECT disk0_image TYPE image {  
    ...  
    files: (("enterprise3-rh73.craff", "ro", 0, 0x4c5abc000, 0),  
           ("checkpoint-1.disk0_image", "ro", 0, 0x4c5abc000, 0))  
    size: 0x4c5abc000  
    ...  
}  
...
```

The checkpointed image is based on the file `enterprise3-rh73.craff`, on top of which is added the file `checkpoint-1.disk0_image` that contains the difference between the checkpoint *checkpoint-1* and the initial state.

Files like `checkpoint-1.disk0_image` are often called *diff files* because they contain the difference between the new state and the previous state.

Image Search Path

This section contains more in-depth explanations about image handling that you may skip when reading this guide for the first time.

When successive checkpoints are saved, an image object may become dependent on several diff files present in different directories. To keep track of all files, Simics stores in the checkpoint a *checkpoint path* list that contains the absolute directory paths where image files may be found. Images filenames are then saved as `%n%/filename` where `%n%` represents the number of the entry in the checkpoint path, counting from zero.

Note: Simics's checkpoint path is different from Simics's search path (see section [15.4](#)), although both will be used when looking for image files, as show below.

To summarize, when loading a checkpoint or a new configuration, Simics looks for images in the following way:

- If the filename doesn't contain any path information (like `image.craff`) or contains a relative path (like `test/image.craff`), the file is looked up *first* from the checkpoint directory, *then* from all the path entries in Simics's search path, *in order* (see also section [15.4](#) for more information).

For example, if Simics's search path contains `[workspace]/targets/sunfire/` and the checkpoint is located in `/home/joe/checkpoints/`, Simics will look for the file `test/image.craff` in the following places:

1. `/home/joe/checkpoints/test/image.craff`
2. `[workspace]/targets/sunfire/test/image.craff`

- If the filename contains a checkpoint path marker (`%n%`), the marker is translated using Simics's checkpoint path and the file is looked up in the corresponding path.

For example, if Simics's checkpoint path contains `/home/joe/c1:/home/joe/c2`, the file `%1%/image.craff` will be translated into `/home/joe/c2/image.craff`.

- If the filename contains an absolute path (like `/home/joe/image.craff`), the file path is used as is.

Note: The reason why Simics's search path is involved in the process is that it makes writing new configurations easier. Adding a path to the place where all initial images are located allows you to just specify the image names.

6.2.3 Saving and Restoring Persistent Data

As an alternative to checkpointing, Simics allows you to only save the *persistent* state of a machine, i.e., data that survive when the machine is powered-down. This typically consists of disk images and flash memory or NVRAM contents. A persistent data checkpoint is handled exactly like any other checkpoint and contains the same file set, but only objects containing persistent data are saved. This persistent data checkpoint can be loaded on top of a corresponding configuration later on.

The commands **save-persistent-state** and **load-persistent-state** respectively save and load the persistent data in a configuration.

Note: These commands are often used to save the state and reboot a machine after the disk contents have been modified. Remember that the target OS **might have cached disk contents in memory**. In order to have a clean disk that can be used at boot, you should synchronize the disk, for example by running `init 0` on a Unix target system, or shutting down the operating system, before you issue the **save-persistent-state** command.

6.2.4 Modifying Checkpoints

Checkpoints are usually created by saving a configuration inside Simics, but it is possible to edit or even create checkpoints yourself.

Because a minimal checkpoint only has to include required attributes, *creating a checkpoint from scratch* works relatively well for small configurations. We suggest you use an existing checkpoint as a starting point if you wish to do that. Note that more advanced layers have been built on top of the configuration system to make the creation of a new machine much easier. Refer to section 6.5 for more information.

Modifying checkpoints require some extra care. Adding or removing devices may confuse the operating system, which does not expect devices to appear or disappear while the system is running, and cause it to crash.

Changing the processor frequency may be enough to confuse the operating system. Many operating systems check the CPU frequency at boot time, and base their waiting loops and timing on the value they got. Saving a checkpoint and changing the frequency after boot may affect the simulation and confuse the system. Devices that use processor frequency to

6.3. Inspecting the Configuration

trigger events at specific times may also behave strangely when the frequency suddenly changes.

6.2.5 Merging Checkpoints

If you want to make a checkpoint independent from all previous checkpoints, for example to distribute it, you can use the small **checkpoint-merge** program in `[simics]/bin` from your system command line. It merges the checkpoint with all its ancestors to create a checkpoint that has no dependencies. Specify the checkpoint you want to distribute as the first parameter and the name of the new stand-alone checkpoint as the second. This tool can be used in both Unix and Windows environments, but it can not handle Cygwin paths on Windows.

6.3 Inspecting the Configuration

Object attributes that are of type `integer`, `string` or `object` are directly accessible at the *command-line* with the notation `object->attribute`:

```
# reading the EAX register in an x86 processor
simics> cpu0->eax
0
# writing a new value to EAX
simics> cpu0->eax = 10
simics> cpu0->eax
10
simics>
```

More information about the command-line and scripting is available in [chapter 8](#).

The Eclipse-based User Interface includes a *Configuration Browser* that allows you to visually browse the complete state of the simulation.

Finally, objects and attributes (of all types) are also available when *scripting Simics directly in Python*. Configuration objects are available under the `conf` namespace:

```
# reading the EAX register in an x86 processor
simics> @conf.cpu0.eax
0
# writing a new value to EAX
simics> @conf.cpu0.eax = 10
simics> @conf.cpu0.eax
10
simics>
```

More information about scripting Simics in Python is available in [chapter 8](#).

6.4 Components

All machines in `[simics]/targets/architecture` use components to create configurations. A component is typically the smallest hardware unit that can be used when configuring a real machine, and examples include motherboards, PCI cards, hard disks, and backplanes. Components are usually implemented in Simics using several configuration objects.

Components are intended to reduce the large configuration space provided by Simics's objects and attributes, by only allowing combinations that match real hardware. This greatly simplifies the creation of different systems by catching many misconfigurations.

Components themselves are also configuration objects in Simics. But to avoid confusion, they will always be referred to as components and the objects implementing the actual functionality will be called objects.

6.4.1 Component Definitions

All machines are based on a *top-level* component. The top-level component is the root of the component hierarchy and is often a motherboard, backplane, or system chassis.

When a component is created, it is in a *non-instantiated* state. At this stage only the component itself exists, not the configuration objects that will implement the actual functionality. Once a complete configuration has been created, all included components can be *instantiated*. When this happens, all objects are created and their attributes are set.

A *standalone* component is a component that can be instantiated without being connected to a component hierarchy. A typical example is a hotplug device, such as a PC Card (PCMCIA) or an Ethernet link.

6.4.2 Importing Component Commands

Components in Simics are grouped by machine architecture, or by type, into several collections. Before a component can be used in Simics, the corresponding component collection has to be imported. What the import actually does, is to add CLI commands for creating components in the collection. The most common collections, that are not architecture specific, are `memory`, `pci`, `std` and `timing`. To import all collections that are used by the *Ebony* machine for example, issue the following commands:

```
simics> import-pci-components
simics> import-std-components
simics> import-memory-components
simics> import-ppc440gp-components
```

6.4.3 Creating Components

The `create-<component>` command is used to create non-instantiated components. There is one create command for each component class. The arguments to the create command

6.4. Components

represent attributes in the component. Standalone components can be created both non-instantiated and instantiated. To create instantiated components, there are **new-** commands, similar to the **create-** commands.

The following code creates a non-instantiated 'ebony-board' component representing an *Ebony Reference Board*, called 'board'

```
simics> (create-ebony-board board cpu_frequency = 100
.....          mac_address0 = "00:04:ac:00:50:00"
.....          mac_address1 = "00:04:ac:00:50:01"
.....          rtc_time = "2003-09-03 11:17:00 UTC")
.....
```

The parentheses are needed to allow multi-line input. The command arguments set the processor frequency, the MAC addresses of the on-board network adapters, and the time and date of the real-time clock.

In multi machine configurations it is often useful to separate objects from the different machines by name. The command **set-component-prefix** *str* causes all following **create-** commands to prefix all created object names—including the names of the components themselves—with the string *str*.

6.4.4 Connectors

A connector provides a means for a component to connect to other components. Connectors have a defined direction: *up*, *down*, or *any*. The direction is *up* if it needs an existing hierarchy to connect to; for example, the PCI-bus connector in a PCI device must connect to a PCI slot. A connector has a *down* direction if it extends the hierarchy downwards; for example, a PCI slot is a connection downward from a board to a PCI device. There are also non-directed connectors, with direction *any*. You can only connect an *up* to a *down* connector or to an *any* connector, and similar for *down* connectors. Connectors with the *any* direction can not be connected together.

Many connectors have to be connected before the component is instantiated, while others can be empty. A standalone component, as described above, may have all connectors empty.

A *hotplug* connector supports connect and disconnect after instantiation. Other connectors can only be connected, or left unconnected, when the configuration is created and may not be modified after that point. A *multi* connector supports connections to several other connectors. A typical example of a *hotplug multi* connector is the Ethernet link.

It is not possible to connect instantiated components with non-instantiated ones. The reason is that the instantiated component expects the other to have all objects already created, and need to access some of them to finish the connection.

The **info** command of a component lists all connectors and some information about them:

```
simics> board.info
Information about board [class ebony-board]
=====
```

Implementing objects:

Connectors:

```

ddr-slot0 : mem-bus          down
ddr-slot1 : mem-bus          down
    emac0 : ethernet-link    down  hotplug
    emac1 : ethernet-link    down  hotplug
pci-slot0 : pci-bus          down
pci-slot1 : pci-bus          down
pci-slot2 : pci-bus          down
pci-slot3 : pci-bus          down
    uart0 : serial           down  hotplug
    uart1 : serial           down  hotplug

```

Since the component in the example isn't instantiated yet, the list of implementing objects is empty. The *ebony-board* has two slots for DDR SDRAM modules, four PCI slots, two serial ports and two Ethernet ports. The memory slot is not listed as *hotplug* since DDR modules have to be inserted when the machine is configured initially, while serial and Ethernet ports support connect and disconnect in run-time. As the *ebony-board* is a top-level component, there are no *up* connectors.

To enable input and output for the simulated machine, the following commands create a serial text console and connects it to the `uart0` connector of the *Ebony* board. Since the text console only has a single connector, it does not have to be specified in the connect command.

```

simics> board.connect uart0 (create-std-text-console)
Created non-instantiated 'std-text-console' component 'text_console_cmp0'.

```

Since no name is given the console component, a unique name will be assigned to it by the configuration system. The **create-** command returns the name of the newly created component, allowing it to be used as in the example above.

If the board only had a single serial connector, the `uart0` connector name could have been left out as well.

Since the machine needs some memory to run, we also add a DDR memory module to our example. A CLI variable is used to hold the name of the memory component.

```

simics> $ddr = (create-ddr-memory-module rank_density = 128 module_data_width = 64)
Created non-instantiated 'ddr-memory-module' component 'ddr_memory_module_cmp0'.
simics> board.connect ddr-slot0 $ddr

```

6.4.5 Instantiation

When a component hierarchy has been created, it can be instantiated using the **instantiate-components** command. This command will look for all non-instantiated top-level components and instantiate all components below them. The **instantiate-components** command

6.4. Components

can also be given a specific component as argument. Then only that component will be instantiated, including its hierarchy if it is a top-level component.

```
simics> instantiate-components
```

If there are unconnected connectors left that may not be empty, the command will return with an error.

When the instantiation is ready, all object and attributes have been created and initialized. In our example, a text console window should have opened. The hardware of the simulated Ebony board is now properly configured, but since no software is loaded, it will not show any output on the console if the machine is started.

6.4.6 Inspecting Component Configurations

The **list-components** command prints a list of all components in the system. All connectors are included, and information about existing connections between them.

The **info** name-space command provides static information about a component, such as the implementing objects and a list of connectors.

The **status** name-space command provides dynamic information about a component, such as attribute values and a list of all current connections. The output from status in the Ebony example:

```
simics> board.status
Status of board [class ebony-board]
=====

Attributes:
            rtc_time : 2003-09-03 11:17:00 UTC
cpu_frequency : 100
mac_address0 : 00:04:ac:00:50:00
mac_address1 : 00:04:ac:00:50:01

Connections:
            uart0 : text_console_cmp0
            ddr-slot0 : ddr_memory_module_cmp0
```

The Eclipse-based User Interface includes a *Configuration Browser* where the configuration can be browsed visually.

6.4.7 Accessing Objects from Components

When doing more advanced configuration of a machine, it may be necessary to access configuration objects and their attributes directly. Since the objects are created by the component system automatically during instantiation, the names of the implementing objects are not known in advance. For this reason it is possible to query a component about what configuration objects are used to implement it, using **get-component-object**. The argument is

a name for the object that is the same for all components of the same class. A list of such names, and their mapping to actual configuration object names, is available in the output from the **info** command. The next example prints the *pc* attribute from the *cpu* object.

```
simics> p -x (board.get-component-object cpu)->pc
```

The **get-component-object** is mainly useful in scripts, when running interactively it is easy to find object names using the **info** command, or **list-objects**.

The **get-component-object** command does not work for non-instantiated components since they do not have any associated configuration objects. But it is possible to access the `pre_conf_objects` of the non-instantiated component from Python using the `get_component_object()` function. This function works for both instantiated and non-instantiated components, and Python code as in the following example works in both cases:

```
@print "0x%x" % get_component_object(conf.board, 'cpu').pc
```

6.4.8 Available Components

The *Simics Target Guide* for each architecture lists and describes all components that are available.

6.5 Ready-to-run Configurations

Simics includes many customizable ready-to-run configurations. Because checkpoint files are by definition very static, these example configurations are not checkpoint-based, but rather build on *components* and *scripts* to generate a working simulated machine.

The example configurations are located in separate directories for each system *architecture*: `[simics]/targets/architecture`. Each of these directories contains a number of Simics scripts (i.e., files containing Simics commands):

<machine>-common.simics

Script that defines a complete simulated machine, i.e., both hardware and software, that can be run by Simics directly. The `common` script use the `-system.include` script to define the hardware, and the `-setup.include` script for software configuration. The `-common.simics` scripts may add additional hardware in some cases.

These are the files you want to use to start the standard example machines in this directory.

<machine> in the script name is either a Unix machine name, or a some other name that defines the hardware/software combination.

<machine>-<feature>-common.simics

A script that extends the `-common.simics` script with a new feature. Many minor features, such as the processor frequency, can be controlled using parameters to the `common` script, but features that are mutually exclusive are added as separate scripts. Typical examples are scripts that add different diff-files to the same disk image in the system setup.

6.5. Ready-to-run Configurations

<architecture-variant>-system.include

Script that defines the hardware of a machine. This script can be shared by several simulated machines that are based on the same hardware. The hardware setup is typically configurable using some standard parameters.

<machine>-setup.include

Script that defines the software and possibly configures the machine to run the selected software, for example setting boot path, and scripting automatic login.

The example configurations are designed to work with the disk images distributed by Virtutech. The machines are described in the *Target Guides* corresponding to each architecture.

Several machines may be defined for a given architecture, and thus the corresponding architecture directory will contain several *machine-common.simics* scripts.

Note: The `[simics]/home/machine` directory contains deprecated example machines provided for backward compatibility reasons. These machine configurations are based on a combination of Simics and Python scripts. It is highly recommended to use the component-based scripts defined in `[simics]/targets/architecture` rather than the scripts in `[simics]/home/machine`.

6.5.1 Customizing the Configurations

There are several ways to customize the examples provided with Simics. The table below shows them sorted by how simple they are to use.

Parameters

The machine scripts distributed with Simics can be modified by setting parameters (CLI variables) before the script is actually run. This is the easiest way to change the default configuration. Parameters can typically be used to change properties such as the amount of memory, the number of processors and the primary MAC address. The available parameters are listed in each *Target Guide*.

Scripts

A simulated machine is defined by several scripts, as described above. By replacing the `-common.simics` file with a user defined one, the system can be configured in more details while keeping the machine definition provided by the `-system.include` file. Similarly the `-setup.include` can be replaced to configure different software on the same machine.

Components

Components represents real hardware items such as PCI cards, motherboards, and disks. Using components to configure a machine provides freedom to set up the simulated machine in any way that is supported by the architecture. The `-system.include` files use components to create their machine definitions. A complete description of components is provided earlier in this chapter.

Objects and Attributes

A component is implemented by one or more configuration objects in Simics, and each object has several attributes describing it. Configuring machines on the object and attribute level is *not* supported in Simics, and such configurations may not work in future versions.

Below is an example of a simple configuration based on *ebony*, that uses parameters to configure two machines slightly differently that both run in the same Simics session.

```
$freq_mhz = 100
$memory_megs = 128
$host_name = "ebony0"
set-component-prefix "ebony0_"
run-command-file "ebony-linux-common.simics"

$freq_mhz = 200
$memory_megs = 256
$host_name = "ebony1"
set-component-prefix "ebony1_"
run-command-file "ebony-linux-common.simics"
```

6.5.2 Adding Devices to Existing Configurations

The parameters available for each predefined machine allows the user to do minor modifications. It is also possible to extend the ready-to-run configurations with additional components without creating new machine setups from scratch. This is done by adding components and connecting them to the machine before the **instantiate-components** command is run.

Since the machine setup scripts are located in the read-only master installation of Simics, they should not be modified. User files that add new components should instead be placed in the corresponding `[workspace]/targets/architecture` directory.

The following commands adds a SCSI controller, and one SCSI disk, to one of the PCI slots in the predefined `ebony-linux-common.simics` file. The commands should be placed in a file in the target directory for *ebony* in the workspace, for example in `[workspace]/targets/ebony/ebony-linux-scsi.simics`

```
script-branch {
  wait-for-variable machine_defined
  $sym = (create-pci-sym53c810)
  $scsi_bus = (create-std-scsi-bus)
  $system.connect pci-slot2 $sym
  $scsi_bus.connect $sym
  $scsi_bus.connect (create-std-scsi-disk scsi_id = 0 size = 3221225472)
}
```

6.5. Ready-to-run Configurations

```
run-command-file "%script%/ebony-linux-common.simics"
```

The script works by starting a *script branch* that waits for the CLI variable `$machine_` defined to be written. All machine defining scripts distributed with Simics sets this variable, by convention, when the complete machine has been created but before the components are instantiated. Since many script branches can wait for the same variable, it is possible to extend the same configuration from different scripts in this way.

When the machine has been created, and `machine_defined` variable is written, the script branch continues to execute. It creates a SCSI controller, a SCSI bus and a 3GB disk, and then connects them. The `$system` variable is defined in the `ebony-system.include` where the machine is created. When the script branch has finished executing, the main script will continue and run `instantiate-components`.

The example above extends the `ebony-linux-common.simics` file in the workspace target directory. To access files in the target directory of the main Simics installation, the `%simics%` path shortcut can be used:

```
run-command-file "%simics%/targets/ebony/ebony-system.include"
```

Here's another example that adds a network card in PCI slot 3 and connects it to the same network as the primary ethernet cards:

```
$create_network = yes
```

```
script-branch {  
    # wait for machine to be setup and connect the new eth before instantiation  
    wait-for-variable machine_defined  
    $add_eth = (create-pci-dec21143 mac_address = "10:10:10:10:10:28")  
    $system.connect pci-slot3 $add_eth  
}
```

```
run-command-file "%script%/ebony-linux-common.simics"  
ethernet_link_cmp0.connect $add_eth
```

6.5. *Ready-to-run Configurations*

Chapter 7

Managing Disks, Floppies, and CD-ROMs

In order to use Simics, you must have *images* (also called *disk dumps*) with the operating system and the applications you plan to run. Depending on the type of machine you are using, these images will correspond to the contents of a disk, a flash memory, a CD-ROM, etc. Virtutech provides images that work with the example machines located in the `targets` directory.

Simics images are usually stored in a special format called `craff` (for Compressed Random Access File Format) to save disk space. Simics also accepts a raw binary dump as an image. This is sometimes more practical if you are manipulating images outside Simics. Simics comes with the `craff` utility to manipulate and convert images in `craff` format (see section 7.1.7).

This chapter will explain the following:

- How to work with images in general
- How to use CD-ROMs and floppies with Simics
- How to use the SimicsFS filesystem
- How to import the contents of a real disk inside Simics

To provide you with a more practical overview, here are the ways you can install and modify the operating system and the applications you wish to run:

Using Simics:

- You can install a completely new OS or simply copy files using a simulated CD-ROM drive, by linking it to a real CD-ROM drive on your host machine or by using a CD image file (refer to sections 7.2.1 and 7.2.2).
- You can copy files from the simulated floppy drive by linking it to the real host floppy device or by using a floppy image file (see sections 7.2.3 and 7.2.4).
- You can use SimicsFS to directly access your real file systems from the simulated machine (see section 7.3).

- You can download files over the simulated network (see chapter 10).

Don't forget to read more about *images* in section 7.1.1 to learn how to save or re-use your changes.

Using External Programs

- You can install a new OS along with new programs on a real machine and create an image from the real machine storage (disk, flash memory, etc.). Section 7.4 shows how to perform this with a disk.
- You can modify an image directly with Mtools (see section 7.1.4).
- You can modify an image directly via a loopback device (see section 7.1.5).

7.1 Working with Images

7.1.1 Saving Changes to an Image

If you modify or create new files on a storage device within Simics, you should remember that by default images are *read-only*. This means that the alterations made when running Simics are *not* written to the image, and will last only for the current Simics session. As described in the 6.2.2 section, this behavior has many advantages. You may however want to save your changes to the image to re-use them in future simulations.

The first thing you should do is to make sure that all the changes are actually written to the media you are using. In many cases, the simulated operating system caches operations to disks or floppies. A safe way to ensure that all changes are written back is to shutdown the simulated machine.

When all changes have been written to the media *in* the simulation, you can save the new contents of the image in different ways:

- Using the **save-persistent-state** command, all image changes for persistent storage media are saved to disk as a persistent state. **This is the recommended way of saving your image changes.**
- Using the `<image>.save-diff-file` command, you can manually save a diff file for the images you are interested in.
- Using the `<image>.save` command, you can create a raw binary dump of the image that is completely independent of all previous images and diff files. Note that this is **not** the best way to get a new and shiny image, since the image is saved uncompressed, which can take a lot of time and disk space.

Note: The `<image>.save` actually allows you to save a *partial* dump of an image, which may be useful to dump a specific part of a disk or a floppy.

Once you have saved the images, you can do the following:

7.1. Working with Images

- If you used **save-persistent-state**, you can issue the **load-persistent-state** command just after starting the original configuration. This will add the new changes to the persistent storage media images and the machine will boot with the changes included. **This is the recommended way of using a saved persistent state.**

For example, let us suppose that you saved some new files on the disk of the *enterprise* machine (started with the `enterprise-common.simics` script). You saved the persistent state of the machine after stopping it to the persistent state file `new-files-added`. You can easily create a small script to start *enterprise* with the new files:

```
# enterprise-new-files.simics
run-command-file enterprise-common.simics
load-persistent-state new-files-added
```

- You can also load the original configuration and add the diff files manually to the images, using the `<image>.add-diff-file` command.
- If you are building your own configurations (either as scripts or as checkpoints), you can add the diff files to the *files* attribute of the corresponding **image** object. This corresponds to what the `<image>.add-diff-file` command does.

If you save several persistent states or image diff files that are dependent on each other, it may become cumbersome to take care of all these dependencies and to remember which files are important or not. You can *merge* the states or image diff files to create a new independent state:

- If you are working with persistent states, you can use the `checkpoint-merge` utility to create a persistent state that is independent of all previous files, including the original images provided by Virtutech. **This is the recommended way of creating a new independent image.** You can load it as usual with the **load-persistent-state** command.
- If you saved some image diff files manually, you can use the `craff` utility described below to merge the diff files yourself.

7.1.2 Reducing Memory Usage Due to Images

Although images are divided into pages that are only loaded on request, Simics can run out of host memory if very big images are used, or if the footprint of the software running on the simulated system is bigger than the host memory. To prevent these kind of problems, Simics implements a global image memory limitation controlled by the **set-memory-limit** command.

When the memory limit is reached, Simics will start swapping out pages to disk very much like an operating system would do. The **set-memory-limit** command let you specify the maximum amount of memory that can be used, and where swapping should occur.

Note: This memory limitation only applies to *images*. Although this is unlikely, Simics can run out of memory due to other data structures becoming too large (for example memory profiling information) even though a memory limit has been set.

7.1.3 Using Read/Write Images

As mentioned in section 6.2.2, images can also work as read-write media, although this is **not** recommended. It can be useful sometimes when planning to make large changes to an image (like installing an operating system on a disk).

To make an image read-write in your own configurations, simply set the second parameter (the “read-only” flag) of the *files* attribute in the image object to “rw”.

In a ready-to-run example like *enterprise*, you can change this attribute after the configuration is completed:

```
# read the 'files' attribute
simics> @files = conf.disk0_image.files
simics> @files
[['enterprise3-rh73.craff', 'ro', 0, 20496236544L, 0]]
# provide the complete path to the file
simics> @files[-1][0] = "[workspace]/targets/enterprise/images/enterprise3-rh73.craff"
# change the second element to make the file read-write
simics> @files[-1][1] = "rw"
simics> @files
[['[workspace]/targets/enterprise/images/enterprise3-rh73.craff', 'rw', 0, 20496236544L, 0]]
# set the 'files' attribute to its new value
simics> @conf.disk0_image.files = files
```

Note that by indexing *files* with the index -1, the last element of the array is accessed, which is always the one that should be set read-write, in case *files* is a list of several files.

As you can see in the example above, Simics by default does not look for files in the Simics search path when the files are used in read-write mode. If you do not provide a complete path to a read-write file, a new file will be created in the current directory.

Use this feature with caution. Make sure to take a copy of the original image before running Simics with the image in read-write mode. Remember to synchronize the storage device within the target OS before exiting Simics, for example by shutting down the simulated machine.

7.1. Working with Images

Note: You can not access directly the files attributes by doing

```
@conf.disk0_image.files[-1][0] = "rw"
```

because the *files* attribute is not an indexed attribute. In this case, Python will read the *files* attribute and change the element `[-1][0]` in the resulting list, but this list is then discarded and the change produces no effect. The *Simics Programming Guide* contains a longer discussion explaining the ambiguities introduced by indexed attributes and the `[]` location.

7.1.4 Editing Images Using Mtools

If you have an image that contains a FAT filesystem, you can use Mtools (<http://mtools.linux.lu>) to get read-write access to the image. You must have a raw binary dump of the image for Mtools to work. This can be obtained using the `craff` utility (see section 7.1.7).

If your image is partitioned (a complete disk for example), you may need to give Mtools special parameters like an *offset* or a *partition*. Refer to the Mtools documentation for more information.

7.1.5 Editing Images Using Loopback Mounting

If the host OS supports loopback devices, like, e.g., Linux and Solaris, you can mount an image on your host machine and get direct read/write access to the files within the image. If you have root permissions this allows you to easily and quickly copy files.

Note: Remember that the image must be a raw binary dump. Disk dumps supplied by Virtutech are normally in `craff` format but you can use the `craff` utility to unpack the disk image to a raw image. The resulting images have the same size as the simulated disk, so you need to have sufficient free space on your host disk to contain the entire simulated disk image.

Note: Do not try to loopback mount an image over NFS. This does not work reliably on all operating systems (Linux, for example). Instead, move the image to a local disk and mount it from there.

On Solaris 8 or later:

```
lofiadm -a disk_dump /dev/lofi/1
mount /dev/lofi/1 mnt-point
...
umount mnt-point
lofiadm -d /dev/lofi/1
```

On Linux:

```
mount <disk_dump> mnt_pnt -o loop=/dev/loopn,offset=m
```

Example:

```
# mount /disk1/rh6.2-kde-ws /mnt/loop -o loop=/dev/loop0,offset=17063424
# cd /mnt/loop
# ls
bin    dev    home   lost+found  opt    root   tmp    var
boot   etc    lib    mnt         proc   sbin   usr
#
```

As shown in the example, the disk dump containing a Red Hat 6.2 KDE WS is mounted on the `/mnt/loop` directory. The file system mounted on `/` starts on the offset 17063424 on the disk. Linux autodetects the file system type when mounting (ext2 in this example). If you want to access another kind of file system, use the `-t fs` option to the mount command. Once the file system is mounted, you can copy files in and out of the disk image.

The offset can be calculated by examining the partition table with `fdisk` (from within Simics). Use `mount` to find the partition you want to edit or examine (e.g., `/dev/hda2` is mounted on `/usr` which you want to modify). Next, run `fdisk` on the device handling this partition (such as `fdisk /dev/hda`). From within `fdisk`, change the display unit to sectors instead of cylinders with the `u` command and print the partition table with `p`. You will now see the start and end sectors of the partitions; you can get the offset by taking the start sector multiplied with the sector size (512).

When you have finished examining or modifying the disk, unmount it and touch the disk image. For example:

```
cd
umount /mnt/loop
touch /disk1/rh6.2-kde-ws
```

The modification date of the disk image does not change when if you modify the disk via the loopback device. Thus, if you have run Simics on the disk image earlier, the OS might have cached disk pages from the now modified disk image in RAM. This would cause a new Simics session to still use the old disk pages instead of the newly modified pages. Touching the image file should ensure that the OS rereads each page.

7.1.6 Constructing a Disk from Multiple Files

In some cases, you may want to populate a simulated disk from multiple files covering different parts of the disk. For example, the partition table and boot sectors could be stored in a different disk image file than the main contents of the disk. If that is the case, you cannot use the `<image>.add-diff-file` command: you must set manually the disk image *files* attribute to put each image file at its appropriate location.

Assume you are simulating a PC and want to build a disk from a main file called `hda1_partition.img` and a master boot record image file called `MBR.img`. The main partition

7.1. Working with Images

will start at offset 32256 of the disk, and the MBR (Master Boot Record) covers the first 512 bytes of the disk (typically, you would get the contents of these image files from the real disk as detailed in section 7.4). The following command in Simics's start-up script will build the disk from these two files.

```
create-std-ide-disk disk2 size = 2559836160
@image = get_component_object(conf.disk2, 'hd_image')
@image.files = [{"hda1_partition.img", "ro", 32256, 1032151040, 0},
                ["MBR.img", "ro", 0, 512, 0]]
```

Note that the two image files cover non-overlapping sections of the disk.

7.1.7 The Craff Utility

The images distributed by Virtutech, and in general most of the images created by Simics are in the `craff` file format. The `craff` utility can convert files to and from the `craff` format, and also merge several `craff` files into a single file.

In your Simics distribution you will find `craff` in `[simics]/bin`. The examples below assume that `craff` is present in your shell path.

Convert a raw dump to **craff** format

```
shell$ craff -o mydisk.craff mydisk.img
```

Convert a single **craff** file to a raw file

```
shell$ craff --decompress -o mydisk.img mydisk.craff
```

Merge multiple checkpoint files into a single **craff** file

If more than one input file is specified, they will be merged so that later files override earlier files on the command line.

```
shell$ craff -o merged.craff chkpt1.craff chkpt2.craff chkpt3.craff
```

Add a **craff** file to a raw dump, producing a new dump

```
shell$ craff --decompress -o new.img mydisk.img diff.craff
```

The input files can be any combination of raw and `craff` files.

Make a file of the differences of two dumps

```
shell$ craff --diff -o diff.craff dump1.img dump2.img
```

The resulting file, `diff.craff`, will contain only what is needed to add to `dump1.img` in order to get `dump2.img`. This is useful to save space if little has been changed.

See also the *Simics Reference Manual* for a full description of the `craft` utility and its parameters.

7.2 CD-ROMs and Floppies

7.2.1 Accessing a Host CD-ROM Drive

Accessing the CD-ROM of the host machine from inside the simulation is supported on Linux and Solaris hosts. This is done by creating a **file-cdrom** object using the **new-file-cdrom** command. First, you should insert the CD in the host machine and figure out which device name it uses.

On a Linux host, this is typically `/dev/cdrom`, which is a symbolic link to the actual CD-ROM device, e.g., `/dev/hdc`. Note that you need read/write access to the CD-ROM device for this to work.

On a Solaris host, you need to specify the raw disk device, which can be found by using the `mount` command. The line that shows where the CD is mounted will look something like this:

```
/cdrom/mydisk on /vol/dev/dsk/c0t2d0/mydisk read only/nosuid
on Fri Jul 26 11:52:52 2002
```

This means that the corresponding *raw* disk device will be called `/vol/dev/rdisk/c0t2d0/mydisk`. Note the `rdsd` instead of `dsk`.

When you have the correct device file name, you create a **file-cdrom** object and insert it into the simulated CD-ROM drive:

```
simics> new-file-cdrom /dev/cdrom file-cd0
cdrom 'file-cd0' created
simics> cd0.insert file-cd0
Inserting media 'file-cd0' into CDROM drive
```

Note that you must replace `/dev/cdrom` with the correct host device name as mentioned above, and `cd0` with the correct Simics object name. Use the **list-objects** command to find the correct object of class **scsi-cdrom** or **ide-cdrom**.

The `cd0.insert` command simulates inserting a new disk into the CD-ROM drive, and there is also a corresponding `cd0.eject` command that simulates ejecting the disk.

7.2.2 Accessing a CD-ROM Image File

A file containing an ISO-9660 image can be used as medium in the simulated CD-ROM. This image file can be created from real CD-ROM disks, or from collections of files on any disk.

An image can be created from a set of files with the `mkisofs` program, which is available on both Linux and Solaris. For example:

7.2. CD-ROMs and Floppies

```
mkisofs -l -L -o image -r dir
```

Once you have an image file, a **file-cdrom** object can be created, and then inserted into a simulated CD-ROM device in the same way as above:

```
simics> new-file-cdrom myimage.iso  
cdrom 'myimage' created  
simics> cd0.insert myimage  
Inserting media 'myimage' into CDROM drive
```

Note that **cd0** above refers to the Simics object name of the CD-ROM drive. This may, or may not be called **cd0**. To see which object name to use, try the **list-objects** command and look for an object of class **scsi-cdrom** or **ide-cdrom**.

7.2.3 Accessing a Host Floppy Drive

It is possible to access a floppy on the host machine from within Simics if the host is running Linux or Solaris. For example (assuming the floppy device object is called **flp0**):

```
simics> flp0.insert-floppy A /dev/fd0
```

Note: To boot directly from the floppy on a simulated x86 architecture you need to select the “A” drive to be the boot device (in, for example, `enterprise-common.simics`):

```
simics> system_cmp0.cmos-boot-dev A
```

7.2.4 Accessing a Floppy Image File

Sometimes it can be convenient to have copies of boot floppies as image files. To create an image of a floppy you can use the Unix command `dd`:

```
dd if=/dev/fd0 of=floppy.img
```

It is then possible to use this image file in Simics:

```
simics> flp0.insert-floppy A floppy.img
```

Note: To boot directly from the floppy on a simulated x86 architecture you need to select the “A” drive to be the boot device (in, for example, `enterprise-common.simics`):

```
simics> system_cmp0.cmos-boot-dev A
```

Floppies are also a convenient way to move small amounts of data out of the simulated machine. Write the data to the simulated floppy inside the simulated machine, and then extract it from the image.

If it is formatted as FAT filesystem, a floppy image can be manipulated with Mtools (see section 7.1.4 for more information).

7.3 Using SimicsFS

SimicsFS gives you access to the file system of your real computer inside the simulated machine. This greatly simplifies the process of importing files into the simulated machine.

SimicsFS is supported for targets running Solaris 7, 8, 9 and 10, and Linux kernel versions 2.0, 2.2 and 2.4.

SimicsFS is installed on disk dumps distributed by Virtutech. For users booting from other disks, there are a number of steps needed to configure the target system. This process is target OS specific, and is described in the following sections.

SimicsFS is not fully functional on all simulated operating systems. The following limitations apply:

Simulated OS	Limitations
Linux	Access is read-only. (Write support experimental.)
Solaris	Truncating files does not work.
Other	SimicsFS is not currently available.

7.3.1 Installing SimicsFS on a Simulated Linux System

For Linux kernel versions prior to 2.4, the SimicsFS kernel module is called `hostfs` rather than `simicsfs`, so for those kernels just replace the `simicsfs` part in file names with `hostfs` in the following description. When the instructions ask you to copy files into the simulated machine, one of the methods described elsewhere must be used (e.g., network, loopback disk access, or CD-ROM).

- Since there are lots of different Linux kernels, and a module has to match the kernel version, Simics does not provide any pre-compiled `simicsfs` modules. To build your own `simicsfs` module, download the SimicsFS source code from <https://www.simics.net/pub/simicsfs.tar.gz>. The `README_2.4` and `README_2.6` files in the archive will explain how to add the SimicsFS sources to the Linux kernel tree and compile it with the kernel you want to use.

7.3. Using SimicsFS

- Create a new directory `/lib/modules/version/kernel/fs/simicsfs/` on the simulated machine, where *version* is the simulated machine's kernel version.
- Copy the newly compiled SimicsFS module file to the directory `/lib/modules/version/kernel/fs/simicsfs/` on the simulated machine, and make sure it is called `simicsfs.o` for Linux 2.4 (and older), and `simicsfs.ko` for Linux 2.6.
- Create the mount point on the simulated machine with `mkdir /host`.
- Add the following line in the simulated machine's `/etc/fstab` (replace `/host` with your mount point):

```
special      /host      simicsfs     noauto,ro    0 0
```

Note: To use the experimental write support, change `ro` into `rw`.

- Mount SimicsFS with the command `mount /host` on the simulated machine.

SimicsFS should now be working, and by issuing `ls /host` on the simulated machine, you should get a listing of the host machine's files.

7.3.2 Installing SimicsFS on a Simulated Solaris System

These are the steps needed to install SimicsFS on a simulated Solaris, version 7, 8, 9 and 10. When the instructions ask you to copy files into the simulated machine, one of the methods described above must be used (network, loopback disk access, CD-ROM...). Note that the driver included with earlier Simics distributions was called `hostfs` and not `simicsfs`.

- Create a new directory `/usr/lib/fs/simicsfs/` on the simulated machine.
- Copy the file `[simics]/import/sun4u/mount_simicsfs` to `/usr/lib/fs/simicsfs/` on the simulated disk, and rename it to `mount`.
- Copy the file `[simics]/import/sun4u/simicsfs-solversion` (where *version* matches the version of Solaris running on your simulated machine) to `/usr/kernel/fs/sparcv9/` on the simulated machine, and rename it to `simicsfs`.
- Add the following line to `/etc/vfstab` on the simulated disk:

```
simicsfs    -      /host    simicsfs    -      no      -
```

- Create the mount point on the simulated machine with `mkdir /host`.

- When the simulated system is running, issue the command:

```
mount /host
```

You should now be able to do `ls /host` on the simulated system to get a list of the files on the host.

7.3.3 Using SimicsFS

By default, the simulated machine can access the entire file tree of the host computer from the mount point (typically `/host`).

This can sometimes be inconvenient (or dangerous, if the simulator runs untrusted or unreliable code), so it is recommended to set the directory that is visible to the simulated machine using the `<hostfs>.root` command, e.g.:

```
simics> hfs0.root /home/alice/sandbox
```

The command will take effect next time SimicsFS is mounted.

Because of implementation limitations, it is recommended that SimicsFS be chiefly used to copy files into and out from the target machine. In particular, executing binaries residing on the host machine may be unreliable.

Note: When saving a checkpoint while a SimicsFS is mounted, take care that the host files that were used at that time are kept unchanged when the checkpoint is loaded.

7.4 Importing a Real Disk into Simics

It is possible to create an image by copying data from a real disk. If the disk to be copied contains an operating system, you must have at least two operating systems on the machine, since the partition that should be copied should **not** be in use or mounted.

Before making a copy of a disk, some information about the disk should be gathered:

- The number of disk cylinders
- The number of sectors per track
- The number of disk heads
- The offset where the specific partition starts (optional)
- The size of a specific partition (optional)

These numbers can be obtained using the `fdisk` utility.

You can choose to make a copy of the whole disk or just a partition from the disk using the `dd` utility. Example:

7.4. Importing a Real Disk into Simics

```
dd if=/dev/hdb of=hdb_disk.img
```

Note: To save space, you may want to compress the disk image using the `craff` utility. See section [7.1.7](#).

The next step is to prepare the target configuration so it can use the new disk. For x86 targets, the *dredd* machine has a `$disk_files` parameter that can be set to a list of files to use in the image object of the boot disk, and also `$disk_size` that specifies the size of that disk.

```
$disk_size = 1056964608
$disk_files = ["hdb_disk.img", "ro", 0, 1056964608, 0]
```

For other machines, that do not have these parameters, attributes in the the disk object and its corresponding image objects have to be set instead.

Make sure to set the `$disk_size` correctly to reflect the size of the disk that has been copied. If only a partition has been copied, the offset where the partition starts, and the size of the partition, should be set in the file list. If the whole disk has been copied, the offset is zero and the size should be the size of the whole disk. Several partitions can be combined to form the complete disk, as described in section [7.1.6](#).

For an x86 machine, the system component will automatically set the BIOS geometry for the C: disk. It can also be set manually:

```
simics> system_cmp0.cmos-hd C 1023 16 63
```

7.4. *Importing a Real Disk into Simics*

Chapter 8

Simics Scripting Environment

The Command-Line Interface in Simics has simple scripting capabilities that can be used when writing parameterized configurations and scripts with conditional commands. For more advanced scripting, the Python language can be used.

This chapter describes how to write simple scripts in the Simics command line interface (CLI), using control flow commands, and variables. It also explains how the configuration system can be accessed from scripts, and how Python can be used for more advanced scripting.

All commands can be executed either by typing them at the prompt in the Simics console, or by writing them to a file, e.g. `example.simics`, and executing the command `run-command-file example.simics`, or for Python examples: `run-python-file example.py`.

8.1 Script Support in CLI

8.1.1 Variables

The Simics command line has support for string and integer variables. Variables are always prefixed with the `$` character. Variables that are not set have a value of 0.

```
simics> $foo = "some text"
simics> $foo
some text
simics> echo $not_used_before
0
```

There is also support for indexed variables (arrays). This is useful in loops for example.

```
simics> $foo[0] = 10
simics> $foo[1] = 20
simics> echo $foo[0] + $foo[1]
30
```

CLI also has support for local variables, described later in this chapter.

8.1.2 Command Return Values

The return value of a command is printed on the console, unless it is used as argument to some other command. Parenthesis () are used to group a command with arguments together, allowing the return value to be used as argument. The return value can also be used as name-space in another command. Variables can be used in the same way.

```
simics> $address = 0
simics> set $address 20
simics> echo "The Value at address " + $address + " is " + (get $address)
The Value at address 0 is 20
```

```
simics> $id = 0
simics> ("cpu" + $id).print-time
processor          steps          cycles      time [s]
cpu0               0              0           0.0
```

```
simics> $cpu = cpu0
simics> $cpu.print-time
processor          steps          cycles      time [s]
cpu0               0              0           0.0
```

Parenthesis can also be used to enter a multi-line command, making it easier to read scripts with nested command invocations. In the text console, the prompt will change to for code spanning more than one line.

```
simics> (echo 10
.....    + (20 - 5)
.....    + (max 4 7))
.....
32
```

8.1.3 Control Flow Commands

The script support in CLI has support for standard **if**, **else** and **while** statements.

```
simics> $value = 10
simics> if $value > 5 { echo "Larger than five!" }
Larger than five!
```

The **if** statement has a return value:

```
simics> $num_cpus = 2
simics> (if $num_cpus > 1 { "multi" } else { "single" }) + "-pro"
```

8.1. Script Support in CLI

multi-pro

Note: Multi-line **if-else** statements must have `} else {` on the same line.

It is also possible to have `else` followed by another `if` statement.

```
simics> if $b == 1 {
.....     echo 10
..... } else if $b == 0 {
.....     echo 20
..... } else {
.....     echo 30
..... }
.....
20
```

Loops can be written with the **while** command.

```
simics> $loop = 3
simics> while $loop {
.....     echo $loop
.....     $loop -= 1
..... }
.....
3
2
1
```

In **if** and **while** statements, it can be useful to have variables that are local to the scope and thus do not collide with the names of global variables. By adding **local** before the first assignment of a variable, the variable is made local.

```
simics> $global = 10
simics> if 1 > 0 {
.....     local $global = 20
.....     echo $global
..... }
..... echo $global
20
10
```

8.1.4 Integer Conversion

In some cases it is useful to interpret an integer as a signed value of a specific bit size, for example when reading four bytes from memory that should be interpreted as a signed 32 bit integer. The **signed8**, **signed16**, ..., **signed64** commands can be used in to perform the conversion. argument.

```
simics> phys_mem.set 0 0xffffffff 4
simics> phys_mem.get 0 4
4294967295
simics> signed32 (phys_mem.get 0 4)
-1
```

8.1.5 Accessing Configuration Attributes

Simics configuration attributes that are of string and integer type can be accessed directly from CLI using the `->` operator.

```
simics> echo "Will switch cpu every " + (sim->cpu_switch_time) + " cycles"
Will switch cpu every 1000000 cycles
```

8.1.6 Script Branches

Introduction to Script Branches

Script branches allow the user to write sequences of CLI commands that can wait for Simics haps at anytime without breaking the sequential flow of commands. This is typically used to avoid breaking a script into many small sections, each installed as a hap callback written in Python.

A simple example from a Simics script:

```
script-branch {
    echo "This is a script branch test - going to sleep."
    cpu0.wait-for-step 10
    echo "Processor registers after 10 steps:"
    cpu0.pregs
}
```

The example above will execute the first **echo** command at once, and then go to sleep waiting until the first 10 instructions (steps) have run. When the step counter for **cpu0** has reached 10, the branch will wake up and run the next two commands, **echo** and **pregs**.

Some commands can not be run while Simics is executing. One example is the **write-configuration** command. To issue such commands from a script branch, it is possible to stop the execution, issue the command and then resume the simulation. The following is

8.1. Script Support in CLI

an example that writes a checkpoint when the simulation reaches a login prompt, and then continues running. It assumes that a text-console called `con0` is used.

```
script-branch {
    con0.wait-for-string login
    stop
    write-configuration login.conf
    run
}
```

Waiting for Haps in Script Branches

A common use of script branches is to wait for a hap to occur before continuing the script execution. Most haps have some data associated with them, such as the exception number for the `Core_Exception` hap used in the example below. This data can be accessed from CLI by telling the **wait-for-hap** command to save it into a named indexed variable with local scope. See the hap documentation for information on what data is associated with each hap type.

```
script-branch {
    wait-for-hap Core_Exception info
    echo "Processor " + $info[0] + " got exception " + $info[1]
}
```

How Script Branches Work

When a script branch is started (using **script-branch**), it begins executing immediately, and runs until a **wait-for-** command is issued. Execution is then resumed in the main script; i.e., there is never any concurrent activity. When a hap, or some other activity, occurs that a script branch is waiting for, the branch continues executing once the currently simulated instruction is ready.

Note: Since only one branch can be active at once, any callback to Python from Simics will execute in the currently active branch, i.e., if a branch installs a callback, it is most likely that it will be called when the main branch is active.

Script Branch Commands

The following is a list of the commands related to script branches.

script-branch

Create a new script branch and start it.

list-script-branches

List all existing, but suspended, branches.

interrupt-script-branch

Interrupt a script-branch, causing it to exit.

wait-for-hap *hap* [*object*] [*index|range-start*] [*range-end*]

Suspend branch waiting for a hap to occur.

wait-for-variable *variable*

Suspend branch until a specified CLI variable is modified. This can be used for synchronization between script branches.

<**processor**>.wait-for-cycle *cycle*

Suspend branch until the specified *cycle* on the processor has been reached.

<**processor**>.wait-for-step *step*

Suspend branch until the specified *step* on the processor has been reached.

<**text-console**>.wait-for-string *string*

Suspend branch until *string* is printed on the text console.

Variables in Script Branches

Variable references in CLI are evaluated when accessed. This is important to remember when writing script branches, since some commands are executed when the branch has been suspended, and variables may have been changed. To make sure that CLI variables in script branches are untouched by other scripts, they should be made local.

The following example

```
script-branch {
    $foo = 20
    cpu0.wait-for-step 10
    echo "foo is " + $foo
}
$foo = 15
run
```

will produce the output `foo is 15` while the following script will print `foo is 20`.

```
script-branch {
    local $foo = 20
    cpu0.wait-for-step 10
    echo "foo is " + $foo
}
$foo = 15
run
```


8.2. Scripting Using Python

Canceling Script Branches

It is possible to cancel a suspended script branch by interrupting it using the **interrupt-script-branch** command. Each branch has an ID associated that can be found using **list-script-branches**, and that is returned by the **script-branch** command.

```
$id = (script-branch {  
    wait-for-variable trigger  
})  
  
...  
  
simics> interrupt-script-branch $id  
Command 'wait-for-variable' interrupted.  
Script branch 1 interrupted.
```

Script Branch Limitations

There are some limitations to script branches. The first two in the list are enforced by Simics:

- Script branches may not start new branches.
- The main branch may not issue the **wait-for-** commands.
- Breaking the simulation with multiple control-C, which forces Simics back to prompt, may confuse the Python interpreter about what thread is active. (Interrupting Simics this way typically destroy the simulation state anyway.)

8.2 Scripting Using Python

Simics provides support for the script language Python (<http://www.python.org>). By using Python the user can extend Simics, and control it in greater detail. Python can interface with Simics using functions in the *Simics API*.

8.2.1 Python in Simics

Python is normally hidden from the user by the command line interface (CLI). But since CLI is implemented in Python, it also provides simple access to the Python environment, making it easy to write your own functions and scripts.

Note: All commands in Simics are implemented as Python functions; the source code of the commands is available in the distribution.

To execute some Python code from the command line, the @ character is used to prefix the line. Example:

```
simics> @print "This is a Python line"
This is a Python line
simics>
```

For code spanning more than one line, the prompt will change to `.....` and more code can be inserted until an empty line is entered. The full code block will then be executed. (Note that whitespace is significant in Python.) Example:

```
simics> @if SIM_number_processors() > 1:
.....     print "Wow, an MP system!"
..... else:
.....     print "Only single pro :-(("
.....
Wow, an MP system!
simics>
```

Entering more than one line is useful for defining Python functions. It is also possible to execute Python code from a file, which is done with the **run-python-file** command.

If the Python code is an expression that should return a value to the CLI, the `python` command can be used, or the expression can be back-quoted. The following example selects a file with Python commands to execute depending on the number of processors in the system:

```
simics> run-python-file `"abc-%d.py" % SIM_number_processors() `
```

If the system has 2 processors, the file `abc-2.py` will be executed.

8.2.2 Accessing CLI Variables from Python

CLI variables can be accessed from Python via the `simenv` name space, for example:

```
simics> $cpu = "processor"
simics> @simenv.cpu = simenv.cpu.capitalize()
simics> $cpu
Processor
```

8.2.3 Accessing the Configuration from Python

Configuration Objects

All configuration objects are visible as objects in Python. The global Python module `conf` holds all such objects. Attribute values can be both read and written using attributes in Python. Example: (print the `pci_devices` attribute in a **pci-bus** object that is called **pcibus25B** in the machine where the example was taken from)

8.2. Scripting Using Python

```
simics> @print conf.pcibus25B.pci_devices
[[2, 0, 'glm0']]
```

To try this example in an arbitrary configuration, run **list-objects pci-bus** to find possible **pci-bus** objects to use instead of **pcibus25B**.

Any '-' (dash) character in the object name, or in an attribute name, is replaced by '_' (underscore).

Indexed attributes can be accessed using [] indexing in Python. It is also possible to index other list attributes this way, but it is inefficient since the full list is converted to a Python list before the element is extracted. Here are some examples of indexed attributes access (**sb0** is a **scsi-bus** object, and **phys_mem0** a **memory-space** object):

```
simics> @print conf.sb0.scsi_phases[1]
Arbitration
```

```
simics> @print conf.phys_mem0.memory[0x100000:0x10000f]
(157, 227, 191, 80, 3, 0, 0, 0, 130, 16, 96, 0, 131, 40, 112)
```

```
simics> @conf.phys_mem0.memory[0x100000:0x100003] = (100,101,102)
```

Warning: Python only supports 32 bit integers in keys when doing sliced indexing (no long integers). However, the Simics API treats [m:n] synonymous to [[m, n-1]], so instead of `conf.phys_mem0.memory[0x1fff80082f0:0x1fff80082f8]` (which won't work), write `conf.phys_mem0.memory[[0x1fff80082f0,0x1fff80082f7]]`

Creating Configurations in Python

In addition to using `.conf` files, it is also possible to create and load configurations from Python. The main advantage is that the configuration can be parameterized without the need of multiple `.conf` files. A part of a configuration in Python, typically written as part of a low-level machine setup, may look like:

```
scsi_bus = pre_conf_object('sb5', 'scsi-bus')

sd_image = pre_conf_object('sd5_image', 'image')
sd_image.size = 2128486400

sd = pre_conf_object('sd5', 'scsi-disk')
sd.image = sd_image
sd.scsi_bus = scsi_bus
sd.scsi_target = 1
sd.geometry = [4157200, 1, 1]
```

```
scsi_bus.targets = [[1, sd, 0]]
```

This will create one **scsi-disk**, one **image** and one **scsi-bus** pre-configuration object with the names `sd5`, `sd5_image` and `sb5`. The pre-configuration objects are Python objects that are used to build a configuration before it is actually loaded in Simics.

When all relevant pre-configuration objects have been added, they can be loaded into Simics using the *SIM_add_configuration()* function.

```
SIM_add_configuration((sd_image, sd, scsi_bus), None)
```

This example can be run from a Python (`.py`) file.

Most configurations supplied with Simics are component based and written in Python. However, the way configurations are created differs between targets. Refer to the corresponding *Simics Target Guide* for more information. Python files that are used to create configurations can be found in the file:

`[simics]/<host>/lib/<architecture>_components.py` and, to some extent, in each target architecture directory.

8.2.4 Accessing Command-Line Commands from Python

At times, it can be useful to access command-line commands from a Python script file. This is done using the *run_command(cli_string)* function, which takes a string which is then evaluated by the command-line front-end. For example, write `run_command("pregs")` to execute the **pregs** command. Any return value from the command is returned to Python.

8.2.5 The Simics API

The Simics API is a set of functions that provide access to Simics functionality from loadable modules (i.e., devices and extensions), and Python scripts. All functions in the Simics API have a name that starts with “*SIM_*”. They are described in details in the *Simics Reference Manual*.

By using the **api-help** and **api-props** commands you can get the declarations for API functions and data types. **api-help identifier** will print the declaration of *identifier*. **api-props identifier** lists all declarations where *identifier* appears.

The Simics API functions are available in the `sim_core` Python module. This module is imported into the Python environment in the frontend when Simics starts; for user-written `.py` files however, the module must be imported explicitly, i.e.,

```
from sim_core import *
```

Errors in API functions are reported back to the caller using *frontend exceptions*. The exception is thrown together with a string that describes the problem more in detail. Examples of exceptions are `General`, `Memory`, `Index`, `IOError`...

For the Python environment, Simics defines an exception subclass for each of its defined exceptions in the `sim_core` module. These are raised to indicate exceptions inside the

8.2. Scripting Using Python

API functions. When errors occur in the interface between Python and the underlying C API function, the standard Python exceptions are used; e.g., if the C API function requires an `int` argument, and the Python function is called with a `tuple`, a Python `TypeError` exception is raised.

8.2.6 Haps

A *hap* is an event or occurrence in Simics with some specific semantic meaning, either related to the target or to the internals of the simulator.

Examples of simulation haps are:

- Exception or interrupt
- Control register read or write
- Breakpoint on read, write, or execute
- Execution of a magic instruction (see section [12.1.7](#))
- Device access

There are also haps which are related to the simulator, e.g., (re)starting the simulation or stopping it and returning to prompt.

Note: In Simics documentation, the word *event* is used exclusively for events that occur at a specific point in simulated time, and *hap* for those that happen in response to other specific conditions (like a state change in the simulator or in the simulated machine).

Callback functions from any supported language can be tied to a certain hap. The callback can be invoked for all occurrences of the hap, or for a specified range. This range can be a register number, an address, or an exception number, depending on the hap.

A complete reference of the haps available in Simics can be found in the *Simics Reference Manual*.

Example of Python Callback on a Hap

This example uses functions from the Simics API to install a callback on the hap that occurs when a control register is written. It is intended to be part of a `.simics` script, that extends an UltraSPARC machine setup. The `SIM_hap_add_callback_index()` function sets the index of the control register to listen to, in this case the `%pil` register in an UltraSPARC processor.

```
@pil_reg_no = SIM_get_register_number(conf.cpu0, "pil")

# print the new value when %pil is changed
@def ctrl_write_pil(user_arg, cpu, reg, val):
    print "[%s] Write to %%pil: 0x%x" % (cpu.name, val)

# install the callback
```

```
@SIM_hap_add_callback_index("Core_Control_Register_Write",  
                             ctrl_write_pil, None, pil_reg_no)
```

In CLI, the same example would look like:

```
script-branch {  
    local $reg = ((current-processor).register-number pil)  
    while 1 {  
        wait-for-hap Core_Control_Register_Write $reg info  
        echo "[" + $info[0] + "]" Write to %pil: " + $info[2]  
    }  
}
```

Part III

Simics Networking

Chapter 9

Network Simulation

This chapter describes how Simics models networking. In the simulated environment, two or more target systems may be connected together, forming a simulated network that is fully virtualized. There are also ways to connect the simulated environment to outside systems, by creating networks that connect to simulated systems at one end, and real systems at the other.

This chapter focuses on networking based on simulated Ethernet links, but the principles for network virtualization is not limited to Ethernet. Simple communication links, such as serial connections can be simulated similarly.

9.1 Ethernet Links

Connecting simulated machines over a simulated Ethernet connection is done by creating an **ethernet-link** object that connects to the Ethernet devices in the machines. The link object can be thought of as modeling an Ethernet cable that is plugged in to the connectors on the devices.

The link object models Ethernet at the frame level, in that it performs delivery of complete frames sent from one device to any other device connected to the link, and doesn't model collisions or lower-level signaling details. This means that it can be used either to model point-to-point twisted-pair or fiber cables running at any speed, or even a 10BASE2 or 10BASE5 traditional coaxial bus, without any changes to the model. Another way of viewing **ethernet-link** is to see it as an Ethernet hub or switch to which can be connected two or more devices.

Traffic sent over the link can be anything, including TCP/IP or any other protocol stack that works on top of Ethernet. The link object doesn't need to be configured with any information about how it will be used.

New **ethernet-link** objects can be added with the **new-ethernet-link** command:

```
simics> new-ethernet-link  
Created ethernet-link ethlink0
```

9.2 Link Object Timing

All frames that are sent over a link are delivered to the receiving device/devices at a point in simulated time after the time when it was sent. The delay is the same for every frame, and is called the *latency* of the link. The latency is a configuration parameter that can be set individually for each link object.

Link objects are most often used to communicate between network devices using separate clocks, and because of how Simics handles simulated time, the clocks are not always completely synchronized. To avoid strange causality effects and indeterministic simulation, the latency of any link must not be set so low that data sent over it might reach the recipient at a point in time that it has already passed. This imposes a lower boundary on the latency, called the min-latency of the link. The value of the min-latency depends on the simulation setup, in particular whether the simulation is distributed between several Simics processes or not. See Chapter 11 for information about distributing the simulation.

The latency of a link can be displayed by the `<ethernet-link>.info` command:

```
simics> ethlink0.info
Information about ethlink0 [class ethernet-link]
=====

                Latency : 1 us
        Distribution : local
           Filtering : enabled

Devices:
        Local devices : none
        Remote devices : none

Real network connection:
                Connected : No
```

As can be seen from the example above, the default latency for an Ethernet link object is one microsecond.

The latency of a link can be specified in the `new-ethernet-link` command as a number of nanoseconds. If the latency of a link is set too low, it will be automatically adjusted to the lowest value allowed by the setup. For example, if you launch a new multi machine setup, such as *ebony-linux-multi*, and try to create a link with too low latency, Simics will adjust the latency automatically:

```
simics> new-ethernet-link latency = 100
[ethlink0 info] Adjusting latency to 1e-005 s because the min-latency changed
Created ethernet-link ethlink0
simics> ethlink0.info
Information about ethlink0 [class ethernet-link]
=====
```

9.3. IP Services

Latency : 10 us

...

9.3 IP Services

It is often useful to be able to let the simulated machines use services available on the simulated network, especially for boot-time configuration. Instead of adding a full system simulation of servers running these services, a more efficient implementation is provided as a *Simics service node*.

The **service-node** class provides a virtual network node that acts as a server for a number of protocols, and it can act as a IP router between networks. The services supported are:

- IP Based Routing
- RARP
- DHCP/BOOTP
- DNS
- Bootparam/Whoami RPC (simplified)
- TFTP
- Real network connections (see section [10](#))

There can be any number of **service-node** objects, and each object can be connected to any number of Ethernet links. In most configurations, however, there will usually be a single service node object. The links themselves do not provide any services, so a service node is needed to make the services available to target machines connected to the link.

A service node can be created using the **new-service-node** command:

```
simics> new-service-node
Created service-node sn0
```

This service node can then be connected to an Ethernet link object. For example, if you already have an **ethernet-link** called **ethlink0** and want the service-node to use the IP address 10.10.0.1 on it:

```
simics> sn0.connect ethlink0 10.10.0.1
Connecting sn0 to ethlink0
Setting IP address of sn0 on network ethlink0 to 10.10.0.1
```

Note: The service node needs to have an IP address on each link it is connected to.

The link to attach the service node to can also be specified when creating the service node. This will make the default name given to the service reflect which link object it is connected to. For example, if you already have an **ethernet-link** called **ethlink0** and want the service-node to use the IP address 10.10.0.1 on it:

```
simics> new-service-node link = ethlink0 ip = 10.10.0.1
Created service-node ethlink0_sn0
Connecting ethlink0_sn0 to ethlink0
Setting IP address of ethlink0_sn0 on network ethlink0 to 10.10.0.1
```

For each link that the service node connects to, a **service-node-device** will be automatically created to act as an Ethernet interface. The name of these device objects are a combination of the service node name and the link object name.

9.3.1 IP Based Routing

A **service-node** object can provide IP based routing between Ethernet links, allowing machines attached to different networks to communicate with each other.

Note: To use the routing mechanisms, simulated machines must use the IP address of the service node as a *gateway* for IP based traffic. Configuring a gateway requires system administration skills, and the exact procedure depends on the target operating system.

The service node contains an internal IP routing table that is used for packet routing between connected links. The routing table can be viewed using the **<service-node>.route** command:

```
simics> sn0.route
Destination  Netmask          Gateway  Port
-----
10.10.0.0    255.255.255.0    sn0_ethlink0_dev
```

The output is quite similar to `route` command available on many systems. The *destination* and *netmask* fields specify a target that can be either a network (i.e., a range of addresses) or a single host (with netmask 255.255.255.255). For packets with this target as their destination, the *port* field specifies the service node network device on which link the packet should be sent.

New entries can be added to the routing table with the **<service-node>.route-add** command. If you have a **service-node** called **sn0** connected to two links called **ethlink0** and **ethlink1**, you could, for example, set up routes like this:

```
simics> sn0.route-add 192.168.0.0 255.255.0.0 link = ethlink0
simics> sn0.route-add 192.168.1.0/26 link = ethlink1
simics> sn0.route-add 10.10.0.0 255.255.0.0 192.168.0.1 ethlink0
simics> sn0.route-add 0.0.0.0 255.255.255.255 192.168.1.1 ethlink1
simics> sn0.route
```

9.3. IP Services

Destination	Netmask	Gateway	Port
192.168.0.0	255.255.255.0		sn0_ethlink0_dev
192.168.1.0	255.255.255.192		sn0_ethlink1_dev
10.10.0.0	255.255.0.0	192.168.0.1	sn0_ethlink0_dev
default		192.168.1.1	sn0_ethlink1_dev

The destination address and the netmask identify the target, and should be given as strings in dotted decimal form. If the target is a single host, the netmask should be given as "255.255.255.255".

9.3.2 DHCP and BOOTP

A service node can act as a *Dynamic Host Configuration Protocol* (DHCP) or *Bootstrap Protocol* (BOOTP) server, responding to requests from clients that can read their network configuration from such a server. The DHCP protocol is an extension of the BOOTP protocol, and for many uses the feature set used are more or less the same. The Simics implementation uses the same configuration for both services.

The service node has a table that maps MAC addresses to IP addresses and domain name. This is used to answer DHCP or BOOTP request. An entry to this table can be added with the `<service-node>.add-host` command:

```
simics> sn0.add-host 10.10.0.1 node1 mac="10:10:10:10:10:01"
Adding host info for IP 10.10.0.1: node1.network.sim  MAC: 10:10:10:10:10:01
```

The current contents of the table can be viewed with the `<service-node>.list-host-info` command:

```
simics> sn0.list-host-info
IP          name.domain          MAC
-----
10.10.0.1   node1.network.sim     10:10:10:10:10:01
```

A pool of IP addresses for dynamic DHCP leases can be added with the `<service-node>.dhcp-add-pool` command. This is used for dynamically assigning IP addresses to new clients. When an entry from the pool is used, the new mapping is stored in the internal host info table, including an automatically generated name that can be found through DNS queries.

If a DHCP client's MAC address matches an entry in the table, it is assigned the corresponding IP address. If there is no matching MAC address, the dynamic address pools will be searched for an available IP address.

The DHCP implementation in **service-node** is simple, and might not work with all DHCP clients.

9.3.3 DNS

The service node includes the functionality of a simple *Domain Name Server* (DNS), that a simulated client can use to translate a host/domain name into an IP address and vice versa. The DNS service is based on the same host table as the DHCP service, and only answers requests for A and PTR records.

For entries in the table that will only be used for DNS requests, and not for assigning IP address by means of DHCP, the MAC address can be left out. The `<service-node>.add-host` command can be used to add table entries, and the `<service-node>.list-host-info` command prints the current table. By default, all host entries will use the `network.sim` domain.

```
simics> sn0.add-host 10.10.0.1 donut
Adding host info for IP 10.10.0.1: donut.network.sim
simics> sn0.add-host 10.11.0.1 foo other.domain
Adding host info for IP 10.11.0.1: foo.other.domain
simics> sn0.list-host-info
```

IP	name.domain	MAC	IP	name.domain	MAC
10.10.0.1	donut.network.sim		10.11.0.1	foo.other.domain	

For dynamically added DHCP addresses, a DNS name will be added for the new IP number, so that any address handed out by a DHCP server can be found by the DNS service.

When connected to a real network, the DNS service can do external lookups for names it does not recognize.

9.3.4 TFTP

The service node also supports the *Trivial File Transfer Protocol* (TFTP, see RFC 1350) and can be used to transfer files between the host system (running the simulation) and a simulated (target) client. The TFTP functionality is often used in network booting, together with the BOOTP facilities, to load OS kernels and images, and can also be used interactively from the `tftp` command found on many systems.

Files that should be transferred from the host system to the simulated client should be located in a directory on the Simics path.

Note: This is the same path as used by image objects and can be viewed with **list-directory** command and modified with the **add-directory** command.

Files transferred from the simulated client to the host, will also end up in the same directory.

Note: TFTP is based on UDP, and each package is acknowledged individually before the transfer is allowed to continue. Depending on the latency of the link, the transfer of large files can be slow. In that case, ensuring that the link can use a lower latency will increase performance.

9.4 Distributed Network Simulation

The simulation of systems connected over an Ethernet network can be distributed over several Simics instances in a distributed simulation, as described in Chapter 11. A link object can be configured to have global distribution, which means that several link object in separate Simics instances can form a shared link, over which all connected systems can communicate even if they are simulated by different Simics instances.

A link object with global distribution is automatically joined with any other link object in another Simics instance if it also has global distribution and has the same name.

The way to make a link object have global distribution is to set its *central* attribute to point to the **central-client** object. The `<ethernet-link>.info` shows the current setting. (Note that you must set up a central server and client as described in Chapter 11 for these examples to work.)

```
simics> @conf.ethlink0.central = conf.central_client
simics> ethlink0.info
Information about ethlink0 [class ethernet-link]
=====

                Latency : 1 us
        Distribution : global
                Filtering : enabled
```

9.5 Serial Links

Connecting simulated machines over a simulated serial connection is done by creating a **serial-link** object that connects to the serial devices in the machines. The link object can be thought of as modeling a serial cable that is plugged in to the connectors on the devices.

The link object models serial communication at the character level in a simplified way. The bandwidth for the connection is configured in the link object, which means that the serial devices do not need to be explicitly configured by software.

New **serial-link** objects can be added with the **new-serial-link** command:

```
simics> new-serial-link
Created serial-link serlink0
```

The serial devices can then connect to that link instead of connecting to another device:

```
simics> com0->link = serlink0
```


Chapter 10

Connecting to a Real Network

Connecting a simulator to a real network opens many new possibilities. You can, for example, easily download files to simulated machines using FTP, access the simulated machines remotely through telnet, or test software on simulated machines against real machines.

Before you can connect to a real network you may need to make some preparations on the simulation host to allow Simics to access the host's Ethernet interfaces. These are described in section 10.1. Section 10.2 describes how to select the correct host interface when running on a host with multiple Ethernet interfaces.

The examples in section 10.4 all start from a checkpoint, so if you want to try the examples you will need to prepare a checkpoint according to the instructions in section 10.3.

Simics is very flexible when it comes to connecting to real networks. There are four different types of connections, to cover almost any use. Section 10.4 describes the connection types, and how to choose which one to use.

Finally, section 10.5 describes how you can tune real network connections to improve their performance, and section 10.6 contains a troubleshooting guide in case you run into problems.

Note: Connecting a simulated network to a real network requires some knowledge of network administration issues.

10.1 Accessing Host Ethernet Interfaces

When connecting to a real network using other connection types than port forwarding, Simics needs low-level access to the simulation host's Ethernet interfaces to send and receive packets. However, operating systems do not usually allow user programs low-level access to Ethernet interfaces. You therefore have to configure the simulation host to allow Simics low-level access.

This can be done in two ways. You can give administrative privileges to a small helper program called `openif`, which will access the simulation host's Ethernet interface for Simics. This is called *raw* access in Simics. Or, you can create a virtual Ethernet (TAP) interface that Simics can access with user privileges. This is called *TAP* access in Simics.

Different connection types can use different access types. Ethernet bridging connections can use either raw or TAP access, IP routing connections can only use raw access, and host connections can only use TAP access.

To tell Simics to use raw or TAP access when creating an Ethernet bridging connection, specify *raw* or *tap* as the *host-access* argument to **connect-real-network-bridge**.

10.1.1 Raw Access

With raw access, Simics uses a small helper program, *openif*, to access the simulation host's Ethernet interfaces. *openif* is launched from Simics and executes with administrative privileges. This way you do not need administrative privileges to connect Simics to the real network once *openif* has been installed. *Simics Installation Guide* contains instructions for installing *openif*.

If *openif* is not installed in its default location, you can use the command **network-helper** to tell Simics where to find it. This command needs to be issued before connecting to the real network, for example:

```
simics> network-helper /usr/local/sbin/openif
```

Note: If *openif* is not installed so that it will run with administrative privileges, the connection to the real network will fail.

Usually, an Ethernet interface only receives packets addressed to the host itself, and ignores other packets. However, when you use Ethernet bridging without MAC address translation and raw access, the interface also needs to receive packets addressed to the simulated machines. This can be achieved by enabling promiscuous mode on the interface. On Linux, this is done by specifying the *promisc* argument to *ifconfig*, which requires administrative privileges. For example, if you want to use the *eth0* interface:

```
computer# ifconfig eth0 promisc
```

Enabling promiscuous mode on an interface means that the operating system will have to process all incoming packets. On a network with a lot of traffic, this may therefore cause a performance degradation on the simulation host. To disable promiscuous mode when it is no longer needed, use the *-promisc* argument to *ifconfig*:

```
computer# ifconfig eth0 -promisc
```

10.1.2 TAP Access

Note: TAP access is not supported on Solaris host.

10.2. Selecting Host Ethernet Interface

With TAP access Simics will connect the simulated network to a virtual Ethernet (TAP) interface provided by the operating system. Accessing the TAP interface does not require administrative privileges, so once the TAP interface has been configured you can connect Simics to the real network without administrative privileges.

The TAP interface can either be bridged to a real Ethernet interface to create an Ethernet bridging connection, or configured with an IP address to create a host connection. Either way, creating a TAP interface that Simics can use is done in two simple steps. These commands require administrative privileges:

1. Give the user running the simulation access to the `/dev/net/tun` device.

```
computer# chmod 666 /dev/net/tun
```

2. Create the TAP interface. Here the name of the user that will be using the TAP interface is assumed to be `joe` and the name of the TAP interface will be `sim_tap0`, but you should of course replace them with the correct user name and the name of the TAP interface that you want.

```
computer# tunctl -u joe -t sim_tap0  
Set 'sim_tap0' persistent and owned by uid 4711
```

To remove a TAP interface when you are done with it, use the `-d` argument to `tunctl`. Again, this requires administrative privileges:

```
computer# tunctl -d sim_tap0  
Set 'sim_tap0' nonpersistent
```

Note: The `tunctl` utility is usually not present in default Linux installations, so you may therefore have to install it yourself. It is usually included in the User Mode Linux package. For convenience, a pre-built version is included in `[simics]/hosttype/sys/bin/` (replace *hosttype* with `x86-linux` or `amd64-linux` depending on your host type).

10.2 Selecting Host Ethernet Interface

When you connect to a real network on a host with multiple network interfaces installed, Simics will select one of them for the real network connection. If the default selection is incorrect, you can use the *interface* argument of the command you use to connect to select the desired interface.

The interface name expected by the *interface* argument is the ordinary interface name used by the host operating system. You can list all network interfaces by running `/sbin/ifconfig -a` on the simulation host:

```
joe@computer$ /sbin/ifconfig -a
eth0      Link encap:Ethernet  HWaddr 00:10:18:0A:DE:EF
          BROADCAST MULTICAST  MTU:1500  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:0 (0.0 b)  TX bytes:0 (0.0 b)
          Interrupt:21

eth1      Link encap:Ethernet  HWaddr 00:0C:F1:D1:FF:09
          inet addr:10.0.0.140  Bcast:10.0.0.255  Mask:255.255.255.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:671467287 errors:0 dropped:0 overruns:0 frame:0
          TX packets:647635204 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:3725791210 (3553.1 Mb)  TX bytes:217046405 (206.9 Mb)
          Interrupt:20 Base address:0xdf40 Memory:fceef000-fceef038

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:16436  Metric:1
          RX packets:24929 errors:0 dropped:0 overruns:0 frame:0
          TX packets:24929 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:4164218 (3.9 Mb)  TX bytes:4164218 (3.9 Mb)
```

For example, to use the first interface listed above you would specify `eth0` as the *interface* argument.

10.3 Preparing for the Examples

The examples in section 10.4 use the simulated machine *enterprise*. It is a simulated x86 machine (PC) with a Red Hat Linux 7.3 installation. Unfortunately *enterprise* does not have any services running that can be used to test access from the real network to the simulated machine by default. To get around this we can turn on the *echo* service on TCP port 7.

To be able to try more than one example without repeating the configuration steps, you can prepare a checkpoint that you can then start from when trying the examples. Simply follow these steps:

1. Start Simics with `enterprise-common.simics` script in `[simics]/targets/x86-440bx` directory.
2. The simulation of the *enterprise* machine can be very fast, so you may want to run the **enable-real-time-mode** command at the Simics prompt. This will limit the simulation speed to real-time speed, so that the screen saver will not kick in annoyingly fast.

10.3. Preparing for the Examples

3. Start the simulation and let the enterprise machine boot. Log in as `root` when you get to the login prompt. No password is required.
4. Start `pico`, a simple text editor, to edit the file `/etc/xinetd.d/echo`:

```
[root@enterprise root]# pico /etc/xinetd.d/echo
```

If you prefer *Emacs* or *vi*, those editors are also available.

5. Change the line that reads

```
disable          = yes
```

to

```
disable          = no
```

You can use the arrow keys to navigate in the file.

6. Press `Ctrl` and `o` to save the file. Press `return` when asked for the file name.
7. Press `Ctrl` and `x` to exit `pico`.
8. Restart the server handling the echo port:

```
[root@enterprise root]# /etc/init.d/xinetd restart
```

You should see a couple of messages about the `xinetd` service stopping and starting again.

9. To verify that the echo service started successfully, `telnet` to the echo port.

```
[root@enterprise root]# telnet localhost 7
```

Type a line of text and press `enter`, and it should be echoed on the simulated console.

To quit the `telnet` session, first press `Ctrl` and `5`, and then type `q` and press `enter` at the `telnet>` prompt that appears.

10. Save a checkpoint using the **write-configuration** command now that the simulated machine has been configured. You can then restart from this point if you want to try several of the examples, or if you run into problems.

10.3. Preparing for the Examples

In the examples, the simulated machine sometimes needs to be reconfigured. Since it is running Linux, the Linux configuration commands are used. If you are using a simulated machine with a different operating system you should configure the simulated machine similarly, but the commands may of course be different.

In the examples, the simulation host has the IP address 10.0.0.129 and the real host that communicates with the simulated network has the IP address 10.0.0.240. You should generally replace these addresses with the address of your simulation host and another host on the your real network.

The enterprise machine uses its default IP address 10.10.0.15.

The examples assume that you have a host on the real network that you can telnet to. You should check that you can telnet from the simulation host to the other real host. Just run `telnet <ip>` in a terminal window, where `<ip>` is the IP address of the other real host. If that does not work, you should not expect to be able to telnet from the simulated machine either.

If you do not have any hosts that accept telnet connections on your network, you can test the connection by entering `GET / HTTP/1.0` and a blank line to port 80 of a web server on your network instead. This should return the HTML content of the start page of the server. Here 64.233.161.104 (www.google.com) is used, replace it with the IP address of your web server:

```
[root@enterprise root]# telnet 64.233.161.104 80
Trying 64.233.161.104...
Connected to 64.233.161.104.
Escape character is '^]'.
GET / HTTP/1.0

HTTP/1.0 302 Found
Location: http://www.google.se/cxfer?c=PREF%3D:TM%3D1118841789:S%3DumC
Vbug84n5uBWAo&prev=/
Set-Cookie: PREF=ID=a5e237e2402bdcac:CR=1:TM=1118841789:LM=1118841789:
S=HQ3jOc8_lpeVGj98; expires=Sun, 17-Jan-2038 19:14:07 GMT; path=/;
domain=.google.com
Content-Type: text/html
Server: GWS/2.1
Content-Length: 214
Date: Wed, 15 Jun 2005 13:23:09 GMT
Connection: Keep-Alive

<HTML><HEAD><TITLE>302 Moved</TITLE></HEAD><BODY>
<H1>302 Moved</H1>
The document has moved
<A HREF="http://www.google.se/cxfer?c=PREF%3D:TM%3D1118841789:S%3DumC
Vbug84n5uBWAo&prev="/>here</A>.
</BODY></HTML>
Connection closed by foreign host.
```

10.4. Connection Types

```
[root@enterprise root]#
```

Make sure that the telnet or web server you use is on the same IP subnet as the simulation host, since you may not be able to access other subnets, depending on what connection type you are using.

10.4 Connection Types

There are a four different types of connections between simulated networks and real networks in Simics. Each connection type has its advantages and drawbacks, and which one you should use depends on what you want to do.

All connection types except port forwarding require low-level access to the simulation host's Ethernet interfaces, and therefore require administrative privileges to set up. In most cases, however, you no longer need administrative privileges to access the real network from Simics once low-level access has been set up. See section [10.1](#) for details.

To help you select what type of connection to use, here is a short description of each connection type:

Port forwarding

Port forwarding is the easiest connection type to set up for simple uses. It does not require administrative privileges and does not require you to configure the simulation host or other hosts in any way.

However, port forwarding is limited to TCP and UDP traffic. Other traffic, for example, ping packets that use the ICMP protocol, will not pass through the connection. Since the port forwarding uses ports on the simulation host you will also not be able to use incoming ports that are already used by the simulation host, or ports below 1024 unless you have administrative privileges.

You need to set up a separate forwarding rule for each incoming TCP port and each incoming or outgoing UDP port. This means that if you are using an application that uses many ports, or random ports, the configuration will become cumbersome or impossible.

Outgoing TCP connections on many or random ports can be handled by NAT, so that is not a problem.

Port forwarding allows communication between the simulated machines and both the simulations host and other hosts on the real network.

Ethernet bridging connection

With an Ethernet bridging connection it appears like the simulated machines are directly connected to the real network, both to the simulated machines and the real hosts. The connection allows any kind of Ethernet traffic between the simulated and real networks.

Usually IP addresses from the IP subnet of the real network are used by the simulated machines. In that case you do not have to configure the hosts on the real network in any way.

You can not access the simulation host from the simulated machines using an Ethernet bridging connection.

IP routing connection

An IP routing connection acts just like an IPv4 router between the simulated and real networks. The connection will therefore allow any kind of IPv4 traffic between the simulated network and the real network. Other protocols, for example, IPv6 or IPX, are not supported.

Since the simulated machines and real hosts will be on different subnets, you need to configure routes to the simulated network on real hosts that need to access it.

You can not access the simulation host from the simulated machines using an IP routing connection.

Host connection

With a host connection you connect the simulation host to a simulated network, allowing any kind of Ethernet traffic between the simulation host and the simulated machines.

Host connections do not enable the simulated machines to access other hosts on the real network, unless you enable IP routing on the simulation host. In that case you need to configure routes to the simulated network on real hosts that need to access it.

Basically, if you are only going to use simple TCP services like FTP, HTTP or telnet, you should probably use port forwarding. If you can not use port forwarding but have available IP addresses on the IP subnet of the real network, or if you use other network protocols than IPv4, you should probably use an Ethernet bridging connection. If you do not have available IP addresses on the IP subnet of the real network, but can configure routes on the other hosts on the real network, you should probably use an IP routing connection. Finally, if you want to access the simulated machines from the simulation host but can not use port forwarding, you have to use a host connection.

The commands to create a connection to the real network start with **connect-real-network**, with different suffixes depending on the connection type. They come in two variants.

For each connection type there is a global command that assumes that there is at most one **ethernet-link** object. If there is no **ethernet-link** object one is created. All Ethernet interfaces of all simulated machines in the Simics process are then automatically connected to the **ethernet-link**, and the **ethernet-link** is connected to the real network. This is an easy way to connect all simulated machines in the Simics process to the real network with a single command. For example, to connect all simulated machines to the real network using an Ethernet bridging connection, you would use the global command **connect-real-network-bridge**.

If you have a more complex simulated network setup, you probably do not want to connect all simulated Ethernet interfaces to the same network. In that case, you first create your simulated network setup, and then connect specific **ethernet-links** to the real network. For each connection type there is a command with the same name as the global command that you can run on an **ethernet-link** object to connect it to the real network. For example, if you have an **ethernet-link** object named **ethlink0**, you would use the command

10.4. Connection Types

ethlink0.connect-real-network-bridge to create an Ethernet bridging connection between that particular link and the real network.

The commands related to port forwarding are an exception to this rule. They do not come in variants that can be run on **ethernet-link** objects, but instead have an *ethernet-link* argument that can be used to specify a link.

10.4.1 Port Forwarding

Port forwarding forwards traffic on TCP and UDP ports between the simulated network and the real network. There is also support for forwarding DNS queries from the simulated network to the real network. Port forwarding can be used with any kind of IP network on the host, it is not limited to Ethernet networks.

Port forwarding is probably the easiest way to access the real network if you only need simple TCP or UDP connectivity, for example, for telnet or FTP. Port forwarding is easy to set up. Simics does not need administrative privileges to run port forwarding, and you do not need to configure the simulation host or other hosts on the real network in any way.

The port forwarding is managed by a **service-node** on the **ethernet-link**. It is the **service-node** that listens for traffic both on the real and simulated networks and forwards it to the other side. All port forwarding commands except **connect-real-network** therefore require an **ethernet-link** with a **service-node** as argument.

There are really four distinct parts to Simics's port forwarding solution. There is forwarding of specific ports from the real network to the simulated network, there is forwarding of specific ports from the simulated network to the real network, there is NAT from the simulated network to the real network, and there is forwarding of DNS queries to the real network.

There is also a convenience command named **connect-real-network** that automatically sets up NAT for outgoing traffic, forwarding of DNS queries to the real network, and incoming port forwarding for some common services.

If you want to view the current port forwarding setup, you can use the **list-port-forwarding-setup** command. It will list all incoming and outgoing ports, NAT and DNS forwarding.

Note: Pinging between the simulated network and the real network will not work when using port forwarding, so ping should not be used to test if the connection is working. Ping uses the ICMP protocol, but only TCP and UDP traffic is supported with port forwarding.

The connect-real-network Command

The **connect-real-network** command is a convenience command that sets up NAT for outgoing traffic, enables forwarding of DNS queries to the real network, and opens incoming ports for FTP, HTTP and telnet to a simulated machine. This is an easy way to get inbound and outbound access for common services on a simulated machine.

The command requires a *target-ip* argument that specifies the IP address of the simulated machine that should be targeted by the incoming traffic. If you have multiple simulated machines you can run **connect-real-network** once for each machine. Simics will select different

ports on the simulation host for the incoming services for each simulated machine, and the selected ports are printed in the Simics console.

The **connect-real-network** command does not require you to specify an **ethernet-link**, unless there is more than one. If there is no **ethernet-link** or **service-node**, they will be created automatically.

Example The **connect-real-network** lets us set up all connections that are needed for most simple real network uses with one simple command. We can start from the checkpoint prepared in section 10.3, and then run the **connect-real-network** command with the IP address 10.10.0.15, which is the default address of the enterprise machine:

```
simics> connect-real-network 10.10.0.15
No ethernet-link found, creating 'ethlink0'.
No service-node found, creating 'ethlink0_sn0' with IP '10.10.0.1'.
Connecting device 'lance0' to 'ethlink0'
NAPT enabled with gateway 10.10.0.1 on link ethlink0.
Host TCP port 4021 -> 10.10.0.15:21 on link ethlink0
Host TCP port 4023 -> 10.10.0.15:23 on link ethlink0
Host TCP port 4080 -> 10.10.0.15:80 on link ethlink0
Real DNS enabled at 10.10.0.1 on link ethlink0.
```

From the output you can see that an **ethernet-link** and a **service-node** have been automatically created and connected to the simulated machine. NAPT, DNS forwarding, and incoming port forwarding for FTP, HTTP and telnet have been also enabled.

You should now be able configure the service node as default gateway on the enterprise machine and to telnet to a host on the real network as described in the example of section 10.4.1, and to configure the service node as DNS server on the enterprise machine and use it to look up real DNS names as described in the example of section 10.4.1.

If you have FTP, HTTP or telnet servers on your simulated machine, which the enterprise machine does not have by default, you should also be able to access these servers from the real network through the ports they have been assigned on the simulation host.

Incoming Port Forwarding

The **connect-real-network-port-in** command sets up port forwarding from a port on the host machine to a specific port on a simulated machines. It takes required three arguments *ethernet-link*, *target-ip* and *target-port*, that specify the **ethernet-link**, IP address and port the traffic should be forwarded to.

You can also specify what port on the simulation host should be used to receive the incoming traffic by specifying the *host-port* argument. If you do not, Simics will automatically select a port and print it on the Simics console.

The command can also take the flags *-tcp* and *-udp*, that let you specify whether you want to set up forwarding for a TCP or UDP port. If neither is specified forwarding will be set up for both the TCP and UDP port.

10.4. Connection Types

The **service-node** acts as a proxy for incoming traffic, so when you want to connect to a port on the simulated machine from the real network, you should connect to the corresponding port on the simulation host. You should not use the simulation host as gateway for the simulated network.

Any UDP packets sent to port on the simulation host are forwarded to the specified port and IP address on the simulated network. For the simulated machine to be able to return UDP packets to the real network, a separate forwarding rule must be set up using the **connect-real-network-port-out** command.

Any TCP connections to the port on the simulation host are forwarded to the specified port and IP address on the simulated network. Since TCP connections are two-way, once a connection has been established, data can be sent in both directions.

The FTP protocol needs to open additional ports when transferring files. Simics handles this by automatically opening outgoing ports for FTP when needed, so FTP will work as long as you use active mode FTP.

Example In the checkpoint we prepared in section 10.3 we started the echo service on port 7 of the simulated machine. We can now set up a port forwarding rule that lets us access the echo service from the real network. Start from the checkpoint, create an **ethernet-link** and **service-node**, connect the simulated machine to the **ethernet-link** and run the **connect-real-network-port-in** command like this:

```
simics> new-ethernet-link
Created ethernet-link ethlink0
simics> new-service-node link = ethlink0 ip = 10.10.0.1
netmask = 255.255.255.0
Created service-node ethlink0_sn0
Connecting ethlink0_sn0 to ethlink0
Setting IP address of ethlink0_sn0 on network ethlink0 to 10.10.0.1
simics> lance0.connect ethlink0
simics> connect-real-network-port-in ethernet-link = ethlink0
target-ip = 10.10.0.15 target-port = 7 host-port = 2007 -tcp
Host TCP port 2007 -> 10.10.0.15:7 on link ethlink0
```

The enterprise machine uses the IP address 10.10.0.15 and the echo service runs on TCP port 7. We use port 2007 on the simulation host, but you can use any free port.

Start the simulation. You should now be able to telnet from a real host to the echo port of the simulated machine by telnetting to port 2007 of the simulation host. In our case the simulation host has the IP address 10.0.0.129, replace it with the IP address of your simulation host:

```
bash-2.05# telnet 10.0.0.129 2007
Trying 10.0.0.129...
Connected to 10.0.0.129.
Escape character is '^]'.
Echo this!
```

Echo this!

Press Ctrl and 5.

```
^]
telnet> q
Connection to 10.0.0.129 closed.
bash-2.05#
```

Outgoing Port Forwarding

The **connect-real-network-port-out** command sets up port forwarding from a port on a **service-node** to a specific port on a host on the real network. It takes three required arguments *ethernet-link*, *target-ip* and *target-port*, that specify the **ethernet-link** the **service-node** is connected to, and the IP address and port on the real network the traffic should be forwarded.

You can also specify what port on the **service-node** should be used to receive the outgoing traffic by specifying the *service-node-port* argument. If you do not, Simics will automatically select a port and print it on the Simics console.

The command can also take the flags *-tcp* and *-udp*, that let you specify whether you want to set up forwarding for a TCP or UDP port. If neither is specified forwarding will be set up for both the TCP and UDP port.

The **service-node** acts as a proxy for outgoing traffic, so when you want to connect to a port on the host on the real network from a simulated machine, you should connect to the corresponding port on the **service-node**. You should not use the **service-node** as gateway for the real network.

Any UDP packets sent to port on the **service-node** are forwarded to the specified port and IP address on the real network. For the real host to be able to return UDP packets to the simulated network, a separate forwarding rule must be set up using the **connect-real-network-port-in** command.

Any TCP connections to the port on the **service-node** are forwarded to the specified port and IP address on the real network. Since TCP connections are two-way, once a connection has been established, data can be sent in both directions.

Example By setting up forwarding from a port on a **service-node** to port 23 of a host on the real network, we should be able to telnet to the real host by telnetting to the port on the **service-node** from the enterprise machine. We can start from the checkpoint we prepared in section 10.3, and create an **ethernet-link** and **service-node**, connect the simulated machine to the **ethernet-link** and run the **connect-real-network-port-out** command. Here we use a host on the real network with IP address 10.0.0.240, replace it with the IP address of a real host on your network:

```
simics> new-ethernet-link
Created ethernet-link ethlink0
```

10.4. Connection Types

```
simics> new-service-node link = ethlink0 ip = 10.10.0.1 ↵
netmask = 255.255.255.0
Created service-node ethlink0_sn0
Connecting ethlink0_sn0 to ethlink0
Setting IP address of ethlink0_sn0 on network ethlink0 to 10.10.0.1
simics> lance0.connect ethlink0
simics> connect-real-network-port-out ethernet-link = ethlink0 ↵
service-node-port = 2323 target-ip = 10.0.0.240 target-port = 23 -tcp
Got service node ethlink0_sn0 for ethlink0
10.10.0.1 TCP port 2323 on link ethlink0 -> host 10.0.0.240:23
```

Now start the simulation. We used the IP address 10.10.0.1 and the port 2323 for the **service-node**, so we should be able to telnet to the real host by telnetting to port 2323 of 10.10.0.1 from the enterprise machine:

```
[root@enterprise root]# telnet 10.10.0.1 2323
Trying 10.10.0.1...
Connected to 10.10.0.1.
Escape character is '^]'.
```

SunOS 5.9

```
login: joe
Password:
No directory! Logging in with home=/
Last login: Sun Jun  2 07:45:58 from 10.0.0.211
```

```
Sun Microsystems Inc.   SunOS 5.9           Generic May 2002
$ exit
[root@enterprise root]#
```

NAPT

The **connect-real-network-napt** command sets up *NAPT* (*network address port translation*, also known as just *NAT* or *network address translation*) from the simulated network to the real network. With NAPT enabled, the **service-node** will act as a gateway on the simulated network and automatically forward TCP connections to the real network.

The **connect-real-network-napt** only has one required argument, *ethernet-link*, that specifies the Ethernet link that should be connected to the real network.

You must configure the simulated machines to use the **service-node** as gateway for the real network, so that it is able to capture the outgoing traffic. The simulated machines will then be able to access hosts on the real network using their real IP addresses. If you combine NAPT with DNS forwarding, described in section 10.4.1, you will be able to use the real DNS names of hosts on the real network too.

The NAPT setup is not specific to a simulated machine, so you only need to run **connect-real-network-napt** once for each **ethernet-link**, and all simulated machines on the link get outbound access.

Since NAPT only allows new TCP connections to be opened from the simulated network to the real network, and the FTP protocol need to open new ports when transferring files, you should use passive mode FTP if you connect to an FTP server on a host on the real network from a simulated machine.

Example To try NAPT, we can start from the checkpoint we prepared in section 10.3, create an **ethernet-link** and **service-node**, connect the simulated machine to the **ethernet-link** and run the **connect-real-network-napt** command like this:

```
simics> new-ethernet-link
Created ethernet-link ethlink0
simics> new-service-node link = ethlink0 ip = 10.10.0.1
netmask = 255.255.255.0
Created service-node ethlink0_sn0
Connecting ethlink0_sn0 to ethlink0
Setting IP address of ethlink0_sn0 on network ethlink0 to 10.10.0.1
simics> lance0.connect ethlink0
simics> connect-real-network-napt ethernet-link = ethlink0
NAPT enabled with gateway 10.10.0.1 on link ethlink0.
```

Now start the simulation. Since we gave the **service-node** the IP address 10.10.0.1, the enterprise machine should be configured with 10.10.0.1 as default gateway. This is already the default, which you can see by running **route** in the the simulated console:

```
[root@enterprise root]# route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use Iface
10.10.0.0        *                255.255.255.0    U        0      0        0 eth0
127.0.0.0        *                255.0.0.0        U        0      0        0 lo
default          10.10.0.1        0.0.0.0          UG       0      0        0 eth0
```

If this had not been the case, we could have added 10.10.0.1 as default gateway with the command **route add default gw 10.10.0.1**.

You should now be able to telnet from the simulated machine to hosts on the real network. In this case we telnet to a Solaris machine with IP address 10.0.0.240, replace this with the address of another host on your real network:

```
[root@enterprise root]# telnet 10.0.0.240
Trying 10.0.0.240...
Connected to 10.0.0.240.
Escape character is '^]'.
```

10.4. Connection Types

SunOS 5.9

login: joe

Password:

Sun Microsystems Inc. SunOS 5.9 Generic May 2002

\$ exit

Connection closed by foreign host.

[root@enterprise root]#

DNS Forwarding

The **enable-real-dns** and **disable-real-dns** commands of the **service-node** class enable and disable forwarding of DNS requests to the real network by a **service-node**. This allows simulated machines to look up names and IP addresses of hosts on the real network, using the **service-node** as DNS server.

Example To try DNS forwarding, we can start from the checkpoint we prepared in section 10.3, and create an **ethernet-link** and **service-node**, connect the simulated machine to the **ethernet-link** and run the **enable-real-dns** command like this:

```
simics> new-ethernet-link
Created ethernet-link ethlink0
simics> new-service-node link = ethlink0 ip = 10.10.0.1
netmask = 255.255.255.0
Created service-node ethlink0_sn0
Connecting ethlink0_sn0 to ethlink0
Setting IP address of ethlink0_sn0 on network ethlink0 to 10.10.0.1
simics> lance0.connect ethlink0
simics> ethlink0_sn0.enable-real-dns
```

Now start the simulation. To tell the simulated machine to use the **service-node** as DNS server, the line `nameserver 10.10.0.1` is needed in the file `/etc/resolv.conf` on the simulated machine. However, this is already the default, so nothing needs to be done.

You should now be able to look up the addresses of real hosts on the simulated machine, for example, `www.google.com`:

```
[root@enterprise root]# nslookup www.google.com
Note: nslookup is deprecated and may be removed from future releases.
Consider using the 'dig' or 'host' programs instead. Run nslookup with
the '-sil[ent]' option to prevent this message from appearing.
Server:          10.10.0.1
Address:         10.10.0.1#53
```

Name: www.google.com
 Address: 216.239.59.104

10.4.2 Ethernet Bridging

Simics can act as a bridge between simulated Ethernet networks and the real Ethernet networks of the host. With this type of connection it will appear that the simulated machines are directly connected to the real network, both to the simulated machines and to the hosts on the real network. Because of this you should configure your simulated machines with IP addresses from the same subnet as the real hosts.

Since the simulated machines appear to be located on the real network, there is no need to configure routes on real hosts that communicate with it. They can find the simulated machines by sending ARP requests, just like they would find other real hosts. You will not be able to access the simulated network from the simulation host.

To create a bridged connection to the real network, use the **connect-real-network-bridge** command. It takes an argument *interface* that specifies what Ethernet interface on the host should be used. It also takes an argument *host-access* that specifies how Simics should access the host's Ethernet interface, see section 10.1. If you omit the arguments, Simics will use raw access to the first interface it finds.

By default Simics will translate MAC addresses between the simulated network and the real network, so that the host's MAC address is used for all packets sent on the real network. This has two advantages. First of all, the simulation host does not have to listen on the real network in promiscuous mode, reducing the load on it. It also avoids problems with MAC address collisions between multiple simulations connected to the same real network. Each simulated machine communicating with the real network must still be configured with a unique IP address, though.

There are a couple of drawbacks with MAC address translation. One is that it is limited to ARP and IPv4 packets; other kinds of packets will not be bridged. There may also be problems with setting up routing on simulated machines. Since the destination MAC addresses are translated, the packets may not go where you expect them.

To be able to bridge other kinds of traffic than ARP and IPv4, for example, DHCP, IPv6, or IPX, you can turn off MAC address translation by specifying the *-no-mac-xlate* flag to **connect-real-network-bridge**. This enables all kinds Ethernet traffic to be bridged between the simulated network and the real network. You must then make sure that each simulated machine connected to the real network is configured with a unique MAC address, to avoid MAC address collisions on the real network.

A big drawback with disabling MAC address translation is that the host must listen on the real network in promiscuous mode. This increases the risk of dropping packets from the real network that were intended for the simulated network, because unrelated traffic will be bridged to the simulated network. It is therefore recommended that bridging without MAC address translation only be used on dedicated networks, or on network interfaces with very little unrelated traffic.

10.4. Connection Types

Note: To use bridging without MAC address translation with raw access you must set the host's Ethernet interface to promiscuous mode, as described in section 10.1.1. If you do not, traffic will be bridged from the simulated network to the real network, but not the other way around.

When you run the **connect-real-network-bridge** command with TAP access, the operating system must be configured to act as a bridge between the virtual interface and the real interface. That is done by the following steps:

Note: When you set up bridging between a TAP interface and a real Ethernet interface, the host will no longer accept traffic on the real interface. All connections the host has open on the interface will be lost. We therefore strongly recommend that you only set up bridging on dedicated host interfaces.

1. Create a TAP interface, as described in 10.1.2
2. Create a bridge interface and connect the TAP interface and the real interface. You may also want to turn off STP (Spanning Tree Protocol) in the bridge, otherwise you will get STP traffic from the bridge into both the simulated and the real network. Here the name of the created bridge interface is `sim_br0` and the interface used is `eth1`, but you can of course use other names and interfaces.

```
computer# brctl addbr sim_br0
computer# brctl addif sim_br0 sim_tap0
computer# brctl addif sim_br0 eth1
computer# brctl stp sim_br0 off
```

3. Bring up the TAP interface and the bridge interface.

```
computer# ifconfig sim_tap0 promisc up
computer# ifconfig sim_br0 promisc up
```

To remove the bridging when you are done, do the following:

1. Bring down the TAP interface and the bridge interface.

```
computer# ifconfig sim_tap0 down
computer# ifconfig sim_br0 down
```

2. Delete the bridge interface.

```
computer# brctl delbr sim_br0
```

Note: The `brctl` utility is usually not present in default Linux installations, so you may therefore have to install it yourself. It is usually included in the `bridge-utils` package. For convenience, a pre-built version is included in `[simics]/hosttype/sys/bin/` (replace *hosttype* with `x86-linux` or `amd64-linux` depending on your host type).

Example

This example assumes that you are starting from the checkpoint prepared in section 10.3, and that you have installed the `openif` helper program according to the instructions in *Simics Installation Guide*.

To set up an Ethernet bridging connection between the real network and the simulated network, run the **connect-real-network-bridge** command. This will automatically create an **ethernet-link**, connect it to the simulated machine and set up bridging to the real network:

```
simics> connect-real-network-bridge host-access = raw
Created ethernet-link ethlink0
Connecting lance0 to ethlink0
[real_net0 info] Receive buffer size 262142 bytes.
[real_net0 info] Using mmap packet capture.
Ethernet-link 'ethlink0' connected to real network.
```

When using Ethernet bridging, the simulated machine should be configured with an unused IP address and netmask from the real network. In this case we use 10.0.0.241 and 255.255.255.0, replace it with an unused IP address and netmask from your real network:

```
[root@enterprise root]# ifconfig eth0 10.0.0.241 netmask 255.255.255.0
```

The simulated machine is now connected to the real network. Any kind of IP traffic is bridged between the simulated network and the real network. You should be able to ping any real host from the simulated machine. Replace *10.0.0.240* with the address of a host on your real network.

```
[root@enterprise root]# ping 10.0.0.240 -c 5
PING 10.0.0.240 (10.0.0.240) from 10.0.0.241 : 56(84) bytes of data.
64 bytes from 10.0.0.240: icmp_seq=1 ttl=255 time=28.0 ms
64 bytes from 10.0.0.240: icmp_seq=2 ttl=255 time=60.9 ms
64 bytes from 10.0.0.240: icmp_seq=3 ttl=255 time=40.9 ms
64 bytes from 10.0.0.240: icmp_seq=4 ttl=255 time=50.9 ms
64 bytes from 10.0.0.240: icmp_seq=5 ttl=255 time=1.10 ms

--- 10.0.0.240 ping statistics ---
5 packets transmitted, 5 received, 0% loss, time 4047ms
rtt min/avg/max/mdev = 1.106/36.409/60.950/20.734 ms
```

10.4. Connection Types

Of course, it should also be possible to ping from the real host to the simulated machine. Running `traceroute` shows that the simulated machine is connected directly to the real network; there are no routers between it and the real host. Again, replace `10.0.0.240` with another host on your real network.

```
[root@enterprise root]# traceroute 10.0.0.240
traceroute to 10.0.0.240 (10.0.0.240), 30 hops max, 38 byte packets
 1  10.0.0.240 (10.0.0.240)  2.568 ms  1.291 ms  1.292 ms
```

If the IP address of the simulated machine itself is printed and `!H` is printed after the response times, it means the simulated machine can not reach the real host, and you need to check your configuration.

You should also, for example, be able to telnet from the simulated machine to a real host. Again, replace `10.0.0.240` with another host on your real network.

```
[root@enterprise root]# telnet 10.0.0.240
Trying 10.0.0.240...
Connected to 10.0.0.240.
Escape character is '^]'.
```

```
SunOS 5.9
```

```
login: joe
Password:
Last login: Sun Jun  2 07:21:44 from 10.0.0.129
Sun Microsystems Inc.    SunOS 5.9          Generic May 2002
$ exit
Connection closed by foreign host.
```

10.4.3 IP Routing

Simics can act as an IP router between simulated Ethernet networks and the real Ethernet networks of the host. This allows any kind of IPv4 traffic between the simulated machines and hosts on the real network. It will appear to both the simulated machine and the real hosts that there is an ordinary IP router between them.

To create a routed connection to the real network, use the **connect-real-network-router** command. The arguments *ip* and *netmask* specify what IP address and netmask the real network router should use on the simulated network. It will always use the simulation host's IP address on the real network. The *gateway* argument tells the real network router what machine on the simulated network should be used as gateway to other simulated networks, in case there are multiple simulated networks with routers between them. The *interface* argument specifies what Ethernet interface on the host should be used, if there are more than one.

Simics creates a new real network router object, typically named **real_net0**, that performs the routing. It is independent of any **service-node** on the **ethernet-link**, and should use a unique IP address.

When acting as an IP router, Simics will always use *raw* access to the host's Ethernet interface; see section 10.1 for more information.

For the routing to work, you have to configure the simulated machine with a route to the real network with the real network router's IP address as gateway. You will also have to set up a route to the simulated network with the simulation host as gateway on all real hosts that are going to communicate with the simulated machines. Also, you will probably not be able to access the simulated network from the simulation host unless you set up fairly complicated routing rules.

Note: The real network router routes ICMP packets between the real network and the simulated network, but it does not implement the ICMP protocol itself. This means that it will not respond to ping. It also means that it will only show up as an unknown router if you run `tracert` between the real and simulated network.

Example

This example assumes that you are starting from the checkpoint prepared in section 10.3, and that you have installed the `openif` helper program according to the instructions in *Simics Installation Guide*.

To set up an IP routing connection between simulated network to the real network, run the **connect-real-network-router** command. This will automatically create an **ethernet-link** and real network router, and connect the simulated machine and real network router to the **ethernet-link**. Here the IP address 10.10.0.1 is used for the real network router.

```
simics> connect-real-network-router ip = 10.10.0.1
Created ethernet-link ethlink0
Connecting lance0 to ethlink0
[real_net0 info] Receive buffer size 262142 bytes.
[real_net0 info] Using mmap packet capture.
Ethernet-link 'ethlink0' connected to real network.
```

The simulated machine must have a route to the real network with the real network router as gateway, so add one using the `route` command:

```
[root@enterprise root]# route add -net 10.0.0.0 netmask 255.255.255.0
gw 10.10.0.1
```

We also need to create a route to the simulated network with the simulation host as gateway on the real hosts that are going to communicate with the simulated network. The command to do this is different depending on the operating system of the real host. In this case 10.10.0.0 and 255.255.255.0 are the IP address and the netmask of the simulated network, and 10.0.0.129 is the IP address of the simulation host. Replace these with the

10.4. Connection Types

correct addresses and netmask for your setup. Note that you must have administrative privileges to run these commands.

Linux

```
/sbin/route add -net 10.10.0.0 netmask 255.255.255.0 gw 10.0.0.129
```

Solaris

```
/sbin/route add -net 10.10.0 10.0.0.129
```

Windows

```
route add 10.10.0.0 mask 255.255.255.0 10.0.0.129
```

Once this is done, you should be able to ping the real host from the simulated machine. Replace *10.0.0.240* with the address of another host on your real network:

```
[root@enterprise root]# ping 10.0.0.240 -c 5
PING 10.0.0.240 (10.0.0.240) from 10.10.0.15 : 56(84) bytes of data.
64 bytes from 10.0.0.240: icmp_seq=1 ttl=255 time=158 ms
64 bytes from 10.0.0.240: icmp_seq=2 ttl=255 time=10.9 ms
64 bytes from 10.0.0.240: icmp_seq=3 ttl=255 time=1.10 ms
64 bytes from 10.0.0.240: icmp_seq=4 ttl=255 time=1.10 ms
64 bytes from 10.0.0.240: icmp_seq=5 ttl=255 time=1.10 ms

--- 10.0.0.240 ping statistics ---
5 packets transmitted, 5 received, 0% loss, time 4047ms
rtt min/avg/max/mdev = 1.106/34.582/158.650/62.150 ms
```

Of course it should also be possible to ping from the real host to the simulated machine too. By running `traceroute` from the simulated machine to the real host, we can see the real network router. It shows up as an unknown router (1 * * *), since it does not implement the ICMP protocol. Again, replace *10.0.0.240* with the address of another host on your real network:

```
[root@enterprise root]# traceroute 10.0.0.240
traceroute to 10.0.0.240 (10.0.0.240), 30 hops max, 38 byte packets
 1  * * *
 2  10.0.0.240 (10.0.0.240)  100.883 ms  939.796 ms  1.288 ms
```

If this works, the IP routing between the real and simulated network is working, so any kind of IP traffic can pass between them. You should, for example, be able to telnet from the real host to the echo port on the simulated machine, using the simulated machine's IP address:

```

bash-2.05# telnet 10.10.0.15 7
Trying 10.10.0.15...
Connected to 10.10.0.15.
Escape character is '^]'.
Echo echo echo echo!
Echo echo echo echo!

```

Press Ctrl and 5.

```

^]
telnet> q
Connection to 10.10.0.15 closed.

```

10.4.4 Host Connection

Simics can connect a simulated network to a virtual Ethernet (TAP) interface on the simulation host. The simulation host and the simulated machines will then be able to communicate as if they were connected to the same Ethernet network. For example, by configuring the simulation host with an IP address on the TAP interface, the simulation host and the simulated machines will be able to send IP traffic to each other.

Host connection is not supported on Solaris host, since it can only use TAP access, and TAP access is not supported on Solaris.

To connect the simulated network to the TAP interface, you first have to configure the TAP interface on the simulation host, as described in section 10.1.2. You then use the **connect-real-network-host** command, which simply takes the name of the TAP interface as the *interface* argument. The simulation host will now appear on the simulated network. It should be configured with an IP addresses from the same subnet as the simulated machines. The simulated machines will then be able to communicate with it without any further configuration.

If you enable IP routing on the simulation host you will also be able to access other hosts on the real network, very similar to using an IP routing connection as described above. You must then configure the simulated machines to use the simulation host's IP address on the TAP interface as gateway for the real network, and real hosts to use the simulation host's IP address on the real network as gateway for the simulated network.

To enable IP traffic between the simulation host and the simulated machines, you have to set up an IP address on the TAP interface. The simulation host will use this address on the simulated network. Here the IP address *10.10.0.1* and the netmask *255.255.255.0* are used, replace them with an IP address from the subnet used by the simulated machines and the corresponding netmask. This requires administrative privileges:

```

computer# ifconfig sim_tap0 10.10.0.1 netmask 255.255.255.0 up

```

10.4. Connection Types

Simulated machine configurations provided with Simics usually use IP addresses from the 10.10.0.x subnet, so you should typically select an IP address on the form 10.10.0.x and the netmask 255.255.255.0.

If you enable IP forwarding on the simulation host, you will also be able to access the simulated network from other hosts on the real network by routing the traffic through the simulation host:

```
computer# sysctl -w net.ipv4.ip_forward=1
net.ipv4.ip_forward = 1
```

To disable IP forwarding again:

```
computer# sysctl -w net.ipv4.ip_forward=0
net.ipv4.ip_forward = 0
```

Example

This example assumes that you are starting from the checkpoint prepared in section 10.3, that you have set up a TAP interface on the simulation host for host connection according to section 10.1.2, and that you have configured it with the IP address 10.10.0.1 as described above. Here the name if the TAP interface is assumed to be `sim_tap0`, replace it with the name of your TAP interface.

Since host connection is not supported on Solaris host, this example will not work if your simulation host is running Solaris.

To connect the TAP interface to the simulated network, use the **connect-real-network-host** command:

```
simics> connect-real-network-host interface = sim_tap0
Created ethernet-link ethlink0
Connecting lance0 to ethlink0
[real_net0 info] Connecting to existing TUN/TAP device 'sim_tap0'
Ethernet-link 'ethlink0' connected to real network.
```

Any kind of Ethernet traffic can now pass between the simulated network and the simulation host. You should, for example, be able to ping the simulation host from the simulated machine, at the IP address you configured on the TAP interface:

```
[root@enterprise root]# ping 10.10.0.1 -c 5
PING 10.10.0.1 (10.10.0.1) from 10.10.0.15 : 56(84) bytes of data.
64 bytes from 10.10.0.1: icmp_seq=1 ttl=64 time=1.15 ms
64 bytes from 10.10.0.1: icmp_seq=2 ttl=64 time=1.11 ms
64 bytes from 10.10.0.1: icmp_seq=3 ttl=64 time=10.9 ms
64 bytes from 10.10.0.1: icmp_seq=4 ttl=64 time=1.11 ms
64 bytes from 10.10.0.1: icmp_seq=5 ttl=64 time=1.11 ms
```

```

--- 10.10.0.1 ping statistics ---
5 packets transmitted, 5 received, 0% loss, time 4037ms
rtt min/avg/max/mdev = 1.113/3.085/10.932/3.923 ms

```

10.5 Performance

When using other connection types than port forwarding, Simics has to prevent the simulated network from being flooded with packets from the real network. If Simics buffered all incoming traffic while the simulated machine was handling it slower than it was arriving, or while the simulation was stopped, Simics would require arbitrarily large buffers for the incoming traffic.

This is prevented by limiting the amount of traffic that is allowed to enter the simulated network per simulated second. The amount of traffic allowed to enter the simulated network is determined by the *tx_bandwidth* and *tx_packet_rate* attributes of the real network object created for the connection, typically named **real_net0**. The unit of the *tx_bandwidth* attribute is bits per simulated second, and the unit of the *tx_packet_rate* attribute is packets per simulated second. You can set either to unlimited by setting them to 0, but this is not recommended unless you know that only a very limited amount of data will be received. The default is to allow 10 megabits per simulated second and an unlimited number of packets.

In addition to allowing the selected rate of traffic into the simulated network, Simics buffers traffic for an additional 0.1 seconds simulated time. This avoids dropping packets if there is a short peak in the traffic. If more packets arrives once this buffer is full, Simics will drop them.

If you want to get better performance out of the connection to the real network, you may want to alter the *tx_bandwidth* and *tx_packet_rate* attributes. A good strategy is to set the *tx_bandwidth* attribute to the amount of traffic you would expect the simulated machine to be able to handle per simulated second, and then try to increase the *tx_packet_rate* from about 5000 and see at what packet rate you get the best performance.

For example, this will set the limit to 100 megabits and 10000 packets per simulated second:

```

simics> @conf.real_net0.tx_bandwidth = 100000000
simics> @conf.real_net0.tx_packet_rate = 10000

```

10.6 Troubleshooting

A network monitoring tool such as *Ethereal* is invaluable when debugging problems with the real network connections. It is a graphical traffic analyzer that can analyze most common network protocols. Ethereal is available from <http://www.ethereal.com>.

There are some pitfalls one might encounter when trying to connect a simulated network to a real one:

10.6. Troubleshooting

Trying to access the simulation host

Accessing the simulation host from the simulated network, or the other way around, is only supported with port forwarding and host connection. You will not be able to access the simulation host from the simulated network if you set up an IP routing connection or an Ethernet bridging connection.

Real host has no route

If you are using an IP routing connection and trying to communicate with the simulated network from a real host that does not have a route to the simulated network, the real host will drop the packets intended for the simulated network or send them to the wrong router. Running `Ethereal` on the Ethernet interface of the real host will show you if, and in that case where, the packets are sent. They should be sent to the simulation host, but if a route is missing they will usually be sent to the real host's default gateway, which will probably ignore them.

Simulated OS has no route

If you are using a NAPT connection or an IP routing connection and the operating system of the simulated machine does not have a correct route to the real network, the simulated machine will drop the packets or send them to the wrong address. To view the routing setup on the simulated machine, use the command `netstat -r` on Linux and Solaris, or `route print` on Windows. Note that these commands should be executed on the *simulated* machines. The simulated OS should have a default route to the service node in the case NAPT connections, or the real network router in the case of IP routing connections.

Real host and simulation host not on the same subnet

If you are using an IP routing connection, all real hosts that should be communicating with the simulated network need to be on the same IP subnet as the simulation host. Otherwise, the real routers between the real host and the simulation host will not be aware of the simulated network and will drop any packets to it, or confuse it with a real subnet with the same address and route the packets there.

Of course, you can configure the real routers to be aware of the simulated network if you really have to communicate with the simulated network from other subnets.

This restriction does not apply to port forwarding connections and Ethernet bridging connections where the simulated machines use IP addresses from the real subnet.

Simics uses the wrong host network interface

On a host with multiple network interfaces installed, Simics will only use one of them for a real network connection. If the default selection is incorrect, use the *interface* argument of the `connect` command to select the desired network interface. See the *Selecting Host Interface* part of section [10.2](#).

Chapter 11

Distributed Simulation

When simulating more than one target system, Simics will simulate one system at a time, and periodically switch between them so that the simulation will advance similarly for all of them.

However, to better use the computing resources of a multi-CPU simulation host, or a cluster of simulation hosts, it is possible to connect a number of simulation processes so that they form a single distributed simulation session. This can be done on both a single host machine to take advantage of multiple CPUs, or over a network where the processes can connect over a TCP/IP network.

11.1 Synchronization

Ideally, the parallel simulation processes should be completely synchronized, so that anything that happens simultaneously in simulated time also happens simultaneously in real time. This would allow events triggered in one process to almost immediately have a deterministic effect in another process.

Unfortunately, this is not realistic. The amount of overhead to keep the processes synchronized would be overwhelming, leaving little time to do real work. Furthermore, the speed of the simulation varies over time, so that the second 1000 cycles of simulated time may take longer or shorter time to simulate than the first 1000 cycles etc. If all simulation processes were executed in lock-step synchronization, it would mean that most processes would constantly be waiting for the one that was currently slowest.

The remedy for this is to allow the processes some slack. The distributed simulation is not kept completely synchronized, but there is an upper bound placed on the allowed difference in time. If one process is simulating at a certain point in simulated time, all other processes must be kept at a point in simulated time that is close enough to that process.

To put it more mathematically, the absolute difference $|T_n - T_m|$ is always less than or equal to the maximum slack D for any two processes n and m .

The effect of this is that any event in the simulation can not have any deterministic effect in a simulation process other than the one triggering it, unless the effect happens long enough in simulated time after the event that triggered it. The delay that must pass is the same as the maximum slack allowed by the synchronization protocol, and the reason is that

the other process may already be that much ahead in simulated time, so it is impossible to make the event have any effect any sooner.

The maximum slack thus defines the *minimum latency* for inter-process communication.

11.2 Architecture

Synchronization in Simics is implemented by the **central** module. It defines two classes, the **central-server** class and the **central-client** class.

One Simics process is assigned the role of *server* for the distributed simulation. In this process, a **central-server** object is added, and in every Simics process that takes part in the distributed simulation, a **central-client** object is added.

The client objects are connected to the server object, and then makes sure that the Simics processes they belong to are kept synchronized with each other.

If the Simics process that contains the server will run any simulation, it too needs a **central-client** object. An alternative is to place the server in a completely dedicated Simics process without any target machine configuration. In that case it does not need a client object.

All inter-process communication that is part of the simulation is handled by *link objects*, typically Ethernet links, as described in chapter 10.

11.3 Running distributed

To create a central server object in a Simics instance, the **new-central-server** can be used:

```
simics> new-central-server
[central_server info] Listening to port 1909
[central_server info] Listening to unix socket /tmp/simics-central.user
Created central_server
```

Without arguments, this command will listen for TCP connections on the standard port (1909), and on the standard file socket name. This can be changed by specifying on the command line where it should listen for connections.

The minimum latency is by default set to 10 milliseconds, but this can also be changed when creating the central server object. It is not possible to change the minimum latency after the server is created.

Creating a central server like this does not automatically make the Simics instance in which the server runs part of the distributed simulation. If any simulation is to be done in the Simics instance acting as a central server, a central client is required in the same Simics, and can be added with the **connect-central** command:

```
simics> connect-central obj=central_server
Created central_client
[central_server info] New connection with id 0
```

11.4. Example of Distributed Simulation and Network

Connected to Simics Central

Other Simics instances can be connected using TCP connections. This works both for instances on the same host or over the host network:

```
simics> connect-central hostname
Created central_client
[central_client info] Connected to server
Connected to Simics Central
```

When running a client Simics on the same host as the server, it is better to use file sockets. This can be done by specifying the socket file name to the **connect-central**, but if the server uses the default socket file name, and runs as the same user as the client, it is simpler to specify `localhost` as the server. This will automatically be changed to a file socket if the socket is found, otherwise a TCP connection to local host is tried.

```
simics> connect-central localhost
Created central_client
[central_client info] Connected to server
Connected to Simics Central
simics> central_client.info
Information about central_client [class central-client]
=====
```

```
Identification : a string
Server : /tmp/simics-central.user
```

Another way to specify that a Simics instance should connect to a Central server is to use the `-central` flag to start Simics. This makes it very simple to use the same Simics script without modification even when running with different servers.

11.4 Example of Distributed Simulation and Network

This section is an extension of the First Steps guide in Chapter 4.

Simics allows you to distribute a simulation on several computers. The typical setup is to simulate one or more machines on each node. Simics will synchronize the nodes so that they will have the same simulated time elapsed, and also provides a way to share Ethernet links.

In this guide we will start two Simicses on the same computer and connect them together.

Create two terminals, and launch the `ebony-linux-firststeps.simics` in both shells. To be able distinguish the machines, we change the Simics prompt.

```
joe@computer: ebony$ ./simics ebony-linux-firststeps.simics
[...]
```

11.4. Example of Distributed Simulation and Network

```
simics> @SIM_set_prompt("simics1> ")
simics1>
```

Likewise, in the second terminal:

```
joe@computer: ebony$ ./simics ebony-linux-firststeps.simics
[...]
simics> @SIM_set_prompt("simics2> ")
simics2>
```

Select the first simulation again, and enter the following commands in the Simics console:

```
simics1> new-central-server
[central_server info] Listening to port 1909
[central_server info] Listening to unix socket [...]
Created central_server
simics1>
```

This will create the **central_server** object, which acts as the hub in our distributed simulation. All simulations will connect to the central server. The next step is to actually connect the first simulation to the server:

```
simics1> connect-central obj = central_server
Created central_client
[central_server info] New connection with id 0
Connected to Simics Central
simics1>
```

connect-central will create a **central-client** object named **central_client** and connect it to our **central_server** object. Note the usage of the *obj* argument. Normally, we would supply the host name of the computer running the central server. However, since the client and server are running in the same Simics process, we connect directly to the **central_server** object.

Note: You can at any time run **central_server.info** or **central_client.info** to view the current status.

Now we are going to connect the other simulation to the central as well.

```
simics2> connect-central localhost
Created central_client
[central_client info] Connected to server
Connected to Simics Central
simics2> central_client.info
```

11.4. Example of Distributed Simulation and Network

```
Information about central_client [class central-client]
=====
                Identification : computer (process 2013)
                Server : /tmp/simics-central.joe
simics2>
```

Since the **central-server** is running in another Simics process, on the same computer, we supply `localhost` as the host name.

Now both simulations are synchronized. However, the simulated machines are not able to talk to each other yet.

Both machines need a simulated Ethernet link. Create a **ethernet-link** to which we connect the simulated network card.

```
simics2> new-ethernet-link
Created ethernet-link ethlink0
simics2> emac0.connect ethlink0
simics2> ethlink0.info
Information about ethlink0 [class ethernet-link]
=====
[...]
                Distribution : local
[...]
simics2>
```

The first command will create the **ethlink0** object, and the second command connects it to the simulated network card. The **(ethernet-link).info** shows two important things: the link is connected to the simulated network card, but **not** shared via the central yet. We will do that now:

```
simics2> ethlink0->central = central_client
[ethlink0 info] Adjusting latency to 0.01 s [...]
simics2> ethlink0.info
[...]
                Distribution : global
[...]
simics2>
```

Distribution has now changed from *local* to *global*.

Finally, we must also connect the first simulation to the simulated network. Switch to the first Simics instance and enter the following commands:

```
simics1> new-ethernet-link
Created ethernet-link ethlink0
simics1> emac0.connect ethlink0
simics1> ethlink0->central = central_client
```

11.4. Example of Distributed Simulation and Network

```
[ethlink0 info] Adjusting latency to 0.01 s [...]  
simics1> ethlink0.info  
[...]  
                Local devices : <0:0> emac0  
                Remote devices : <1:0> emac0  
[...]  
simics1>
```

Here we see that the network card on the other simulated machine is added to the device list. Hopefully, the two machines are now able to talk to each other. Before we try anything, there is one more thing we must do. Since the two simulated machines are identical, they are assigned the same IP address. Thus, we must change the IP address on one of them.

Boot both machines by giving the **continue** command.

Note: You must resume *both* simulations. If you resume only one simulation it will appear to hang. This is because the simulations are synchronized through the central server.

Then, in one of the consoles, enter the following command:

```
root@firststeps: ~# ifconfig eth0 10.10.0.51  
root@firststeps: ~#
```

This will change the IP on that machine to 10.10.0.51. The other machine should have the IP address 10.10.0.50. Try pinging (from 10.10.0.51):

```
root@firststeps: ~# ping -c 2 10.10.0.50  
PING 10.10.0.50 (10.10.0.50): 56 data bytes  
64 bytes from 10.10.0.50: icmp_seq=0 ttl=64 time=10.0 ms  
64 bytes from 10.10.0.50: icmp_seq=1 ttl=64 time=10.0 ms  
  
--- 10.10.0.50 ping statistics ---  
1 packets transmitted, 1 packets received, 0% packet loss  
round-trip min/avg/max = 10.0/10.0/10.0 ms  
root@firststeps: ~#
```

You can read more about distributed simulation in [chapter 11](#).

Part IV

Developing with Simics

Chapter 12

Debugging Tools

12.1 Breakpoints

Like an ordinary debugger, Simics can run user binaries, allowing the user to set breakpoints, inspect state, single step, etc. Some difficult bugs are easier to find using various esoteric breakpoint types. In Simics you can set breakpoints on:

- memory accesses: any range and combination of read/write/execute
- time (number of cycles or instructions executed)
- instruction types, such as control register accesses
- device accesses
- output in the console

Simics is fully deterministic, allowing you to narrow down the location of difficult bugs. If your session has interactive input, you can record it using the **recorder** and replay when you need to reproduce the same execution. If Hindsight is available, you can of course freely go forward and backward in time until you found the location of the problem.

12.1.1 Memory Breakpoints

A memory breakpoint stops the simulation whenever a memory location in a specified address interval is accessed. The address interval can be of arbitrary length and the type of the memory access can be specified as any combination of *read*, *write*, and *execute*.

Physical memory breakpoints refer to addresses within a memory space, so the breakpoint itself is always connected to a specific memory space object in Simics. If this object is known by name (as **phys_mem0** in the following example), the breakpoint can be set with the **break** command:

```
simics> phys_mem0.break address = 0x10000 length = 16 -w
Breakpoint 1 set on address 0x10000, length 16 with access mode 'w'
```

Virtual memory breakpoints are handled by context objects:

```
simics> primary_context.break 0x1ff00
Breakpoint 1 set on address 0x1ff00 with access mode 'x'
```

Note that by default, all simulated processors in one Simics process share one context. If you want a virtual breakpoint to apply only to a subset of the processors, create a new context just for them:

```
simics> new-context foo
simics> cpu1.set-context foo
simics> cpu7.set-context foo
simics> foo.break 0xffffffffbfc008b8
```

The **break** command can also be used without explicitly specifying an address space or context object. Instead you can prefix the address with **p:** for a physical address, or **v:** for a virtual address. The breakpoint will refer respectively to the memory space (physical address) or context (virtual address) connected to the current front-end processor (as specified with the **pselect** command). Note that, unless you have created a new context for the current processor, the breakpoint will apply to all processors.

As you can see in the following example, Simics interprets a breakpoint address as virtual unless **p:** is explicitly specified:

```
simics> break v:0x4711
Breakpoint 2 set on address 0x4711 with access mode 'x'
simics> break p:0x4711
Breakpoint 3 set on address 0x4711 with access mode 'x'
simics> break 0x4711
Breakpoint 4 set on address 0x4711 with access mode 'x'
Note: overlaps with breakpoint 2
```

Execution breakpoints can be modified with filter rules to only trigger when instructions match certain syntactical criteria. This feature is mainly useful with breakpoints covering large areas of memory. The commands available are **set-prefix** (to match the start of an instruction), **set-substr** (to match a particular substring), and **set-pattern** (to match the bit pattern of the instruction). The commands work by modifying an existing breakpoint, so you first have to set an execution breakpoint and then modify it to match only particular expressions.

For example, to stop when an instruction with the name **add** is executed in a memory range from 0x10000 to 0x12000, use the following commands:

```
simics> break 0x10000 0x2000 -x
Breakpoint 1 set on address 0x10000, length 8192 with access mode 'x'
simics> set-prefix 1 "add"
```

12.1. Breakpoints

Simics will stop when the first add instruction is encountered. For more information, see the *Simics Reference Manual* or use the **help break** command.

12.1.2 Temporal Breakpoints

Unlike an ordinary debugger, Simics can handle temporal breakpoints, i.e., breakpoints in time. As the concept of time is based on steps and cycles, a temporal breakpoint refers to a specific step or a cycle count for a given processor object:

```
simics> cpu0.cycle-break 100
simics> cpu0.step-break 100
```

In the example above, the breakpoints are specified relative to the current time. It is also possible to set temporal breakpoints in absolute time (where 0 refers to the time when the original configuration was set up in Simics). When Hindsight is available, you can freely set time breakpoints in the past as well as in the future.

```
simics> cpu0.cycle-break-absolute 100
simics> cpu0.step-break-absolute 100
```

All the commands `cycle-break`, `step-break`, `cycle-break-absolute`, and `step-break-absolute`, can be given without prefixing them with the CPU. Note that in this case the commands will plant a breakpoint for current front-end processor (and not all processors).

12.1.3 Control Register Breakpoints

A control register breakpoint is triggered when a selected control register is accessed. The control register is specified either by name or number, and the access type can be any combination of *read* or *write*. For example:

```
simics> break-cr reg-name = asi
```

Note that the exact arguments to this command depend on the target architecture. A list of available control registers can be obtained by tab-completing the *reg-name* argument. See the documentation for **break-cr** in the *Simics Reference Manual* for more information..

12.1.4 I/O Breakpoints

An I/O breakpoint is always connected to a specific device object. The breakpoint is triggered when that device is accessed. The breakpoint is set using the **break-io** command, which takes the device name as a parameter. For example, to break on accesses to the **hme0** device, we would use the following syntax:

```
simics> break-io object-name = hme0
```

A list of devices can be obtained by tab-completing the *object-name* argument.

12.1.5 Graphics Breakpoints

The graphics-console can be used to save and set graphical breakpoints. A graphical breakpoint is a rectangular area on the simulated display that triggers a hap (`Gfx_Break_String`) whenever the pixels inside the saved breakpoint rectangle exactly match those on the display.

The following commands can be used to save and set breakpoints for a graphics console:

<gfx-console>.save-break *filename* [*comment*]

Let the user select a rectangular area inside the graphics console using the mouse pointer. To cancel the select operation press the right mouse button. The selected area will be saved as a graphical breakpoint file. You can add an optional comment that will be put in the beginning of the (binary) breakpoint file.

<gfx-console>.save-break-xy *filename left top right bottom* [*comment*]

Let the user specify a rectangular area inside the graphics console using the top left and bottom right corners coordinates. The selected area will be saved as a binary graphical breakpoint file. You can add an optional comment that will be put at the beginning of the breakpoint file.

<gfx-console>.break *filename*

Activate a previously saved breakpoint and return a breakpoint id. When a graphical breakpoint is reached, it is immediately deleted and a `Gfx_Break_String` hap is triggered. If no callbacks for this hap were registered, Simics halts execution and returns to the command prompt.

<gfx-console>.delete *id*

Delete the breakpoint associated with *id*.

12.1.6 Text Output Breakpoints

The text console can set breakpoints on the occurrence of certain character sequences in the output sent to the screen.

- To set a breakpoint, use the command ***console.break string***. Simics will stop when *string* appears in the output.
- Use ***console.unbreak string*** to remove a particular breakpoint.
- All breakpoints can be listed using the ***console.list-break-strings*** command.

Note: To find out if a specific simulated machine uses a text console, look for an object of class **text-console** in the list provided by **list-objects** once the configuration is loaded.

12.1.7 Magic Instructions and Magic Breakpoints

For each simulated processor architecture, a special no-operation instruction has been chosen to be a **magic instruction** for the simulator. When the simulator executes such an instruction, it triggers a `Core_Magic_Instruction` hap and calls all the callbacks functions registered on this hap (see chapter 8 to get more information about haps).

If the architecture makes it possible, an immediate value is encoded in the magic instruction. When the hap is triggered, this value is passed as an argument to the hap handlers. This provides the user with a rudimentary way of passing information from the simulated system to the hap handler.

Magic instructions have to be compiled in the binary files that are executed on the target. The file `magic-instruction.h` in `[simics]/src/include/simics/` defines a `MAGIC(n)` macro that can be used to place magic instructions in your program, where *n* is the immediate value to use.

Note: The declaration of the macros are heavily dependent on the compiler used, so you may get an error message telling you that your compiler is not supported. You will need to write by yourself the inline assembly corresponding to the magic instruction you want to use. The GCC compiler should always be supported.

Note: Using magic instructions in other languages than C requires the ability to insert inline assembler in a program, or at least the ability to call arbitrary functions written in assembly. For example, in Java it would be necessary to use the JNI interface. As always, check your compiler and language documentation for details on how to enter inline assembly.

A complete definition of magic instructions and the values the parameter *n* can take is provided in figure 12.1.

Target	Magic instruction	Conditions on <i>n</i>
Alpha	<i>binary</i> : 0x70000000	$n = 0$
ARM	<code>orreq rn, rn, rn</code>	$0 \leq n < 15$
IA-64	<code>nop (0x100000 + n)</code>	$0 \leq n < 0x100000$
MIPS	<code>li %zero, n</code>	$0 \leq n < 0x10000$
MSP430	<code>bis r0, r0</code>	$n = 0$
PowerPC 32-bit	<code>mr n, n</code>	$0 \leq n < 32$
PowerPC 64-bit	<code>fmr n, n</code>	$0 \leq n < 32$
SPARC	<code>sethi n, %g0</code>	$1 \leq n < 0x400000$
x86	<code>xchg %bx, %bx</code>	$n = 0$

Figure 12.1: Magic instructions for different Simics Targets

Here is a simple pseudo-code example:

```

#include "magic-instruction.h"

int main(int argc, char **argv)
{
    initialize();
    MAGIC(1);                                tell the simulator to start
                                              the cache simulation

    do_something_important();
    MAGIC(2);                                tell the simulator to stop
                                              the cache simulation

    clean_up();
}

```

This code needs to be coupled with a callback registered on the magic instruction hap to handle what happens when the simulator encounters a magic instruction with the arguments 1 or 2 (in this example, to start and stop the cache simulation).

Simics implements a special handling of magic instructions called **magic breakpoints**. A magic breakpoint occurs if magic breakpoints are enabled and if the parameter n of a magic instruction matches a special condition. When a magic breakpoint is triggered, the simulation stops and returns to prompt.

Magic breakpoints can be enabled and disabled with the commands **magic-break-enable** and **magic-break-disable**. The condition on n for a magic instruction to be recognized as a magic breakpoint is the following:

$$n == 0 \mid\mid (n \ \& \ 0x3f0000) == 0x40000$$

Note that the value 0 is included for architectures where no immediate can be specified. The file `magic-instruction.h` defines a macro called `MAGIC_BREAKPOINT` that places a magic instruction with a correct parameter value in your program.

On architectures that only offer a single magic instruction (x86 and Alpha), more information can be passed from the simulation to the magic instruction hap handler by putting data values in machine registers prior to triggering the magic instruction.

As a concrete example, on the x86, the hap argument can only be 0. This can be worked around by putting extra information into register `eax` before executing the `MAGIC(0)` magic instruction. The hap handler for magic instructions then needs to read the value from `eax` and do different things depending on its contents.

The following is an example program implementing this technique, using the `gcc` compiler's syntax for inserting inline assembler:

```

#include "stdio.h"
#include "magic-instruction.h"

#define MY_MAGIC(n) do {
    asm volatile ("movl %0, %%eax" : : "g" (n) : "eax");
}

```


12.1. Breakpoints

```
MAGIC(0);
} while (0)

int main(void)
{
    printf("Hello,\n");
    MY_MAGIC(1);
    printf("World!\n");
    MY_MAGIC(2);
    return 0;
}
```

The hap handler for this would look something like the following:

```
@def call_back_1(cpu):
    pr("call back one triggered\n")

@def call_back_2(cpu):
    pr("another one here\n")

@def hap_callback(user_arg, cpu, arg):
    eax = cpu.eax          # read value passed from program
    if eax == 1:           # take appropriate action
        call_back_1(cpu)
    elif eax == 2:
        call_back_2(cpu)
    else:
        print "Unknown callback, eax is", eax
        SIM_break_simulation("snore")

@SIM_hap_add_callback("Core_Magic_Instruction", hap_callback, None)
```

The same technique can be applied to other architectures, but you need to adapt the names of the registers involved.

Note: This method is slightly intrusive since a register will change its value. It is thus important to make sure that the compiler is aware of the change to the register so that no broken code is emitted; in the gcc compiler, this is specified in the last argument to the `asm()` statement.

12.2 Using GDB with Simics

This chapter describes how to use **`gdb-remote`**, a Simics module that lets you connect a GDB session running on your host machine to the simulated machine using GDB's remote debugging protocol, and use GDB to debug software running on the target machine.

If you load the **`gdb-remote`** module in Simics, you can use the remote debugging feature of GDB, the GNU debugger, to connect one or more GDB processes to Simics over TCP/IP. In order to do this, you need a GDB compiled to support the simulation's target architecture on whichever host you're running. The **`gdb-remote`** module only supports version 5.0 or later of GDB, but other versions may work as well. Unfortunately GDB's remote protocol does not include any version checking, so the behavior is undefined if you use other versions. For information on how to obtain and compile GDB, see section 12.2.3.

To connect a GDB session to Simics, start your Simics session and run the **`new-gdb-remote`** command, optionally followed by a TCP/IP port number, which defaults to 9123 otherwise. This will automatically load the **`gdb-remote`** module.

When a configuration is loaded, Simics will start listening to incoming TCP/IP connections on the specified port. Run the simulated machine up to the point where you want to connect GDB. To inspect a user process or dynamically loaded parts of the kernel, the easiest solution might be to insert magic instructions at carefully chosen points. For static kernel debugging, a simple breakpoint on a suitable address will solve the problem.

Note: When debugging the start-up phase of an operating system, it might happen that `gdb` gets confused by the machine state and disconnects when you try to connect. In this case, execute a few instructions and try again.

Once Simics is in the desired state, start your GDB session, load any debugging information into it, and then connect it to Simics using the **`target remote host:port`** command, where *host* is the host Simics is running on, and *port* is the TCP/IP port number as described above. Here is a short sample session using *bagle*, a Sun UltraSPARC machine running Linux:

```
(gdb) set architecture sparc:v9a
(gdb) symbol-file vmlinux
Reading symbols from vmlinux...done.
(gdb) target remote localhost:9123
Remote debugging using localhost:9123
time_init () at /usr/src/linux/include/asm/time.h:52
(gdb)
```

Note: For some architectures, you need to give a command to GDB before connecting (the **`set architecture`** command in the session above). These are tabulated in the reference manual's section on **`gdb-remote`**, and will also be printed on the Simics console when you run **`new-gdb-remote`**.

From this point, you can use GDB to control the target machine by entering normal GDB commands such as **`continue`**, **`step`**, **`stepi`**, **`info regs`**, **`breakpoint`**, etc.

12.2. Using GDB with Simics

Note that while a remote GDB session is connected to Simics, the Simics prompt behaves a little differently when it comes to stopping and resuming the simulation. While the GDB session is at prompt, it is impossible to continue the simulation from within Simics (e.g., by using the **continue** command). However, once you continue the execution from GDB, you can stop it from GDB (by pressing control-C), which causes the simulation to stop and makes both GDB and Simics return to their prompts, or you can stop the simulation from the Simics prompt (also by pressing control-C). This only makes Simics return to prompt, while GDB will still think the target program is running. In this state, you should continue the simulation from the Simics prompt before attempting to use GDB.

You can also force GDB back to prompt using the **gdb0.signal 2** command in Simics, which tells the GDB session that the simulated machine got a `SIGINT` signal. **gdb0** here refers to the configuration object created on the fly when the GDB session connected to Simics. You can connect several GDB sessions to one Simics; each connection will be associated to one **gdbnn** object.

Since GDB isn't the most stable software, especially when using remote debugging, it unfortunately hangs now and then. To force Simics to disconnect a dead connection, you can use the **gdb0.disconnect** command.

Note that the **gdb-remote** module does not have any high-level information about the OS being run inside Simics. This means that in order to examine memory or disassemble code, the data or code you want to look at has to be in the active TLB.

Note: When using **gdb-remote** with targets supporting multiple address sizes (such as x86-64 and SPARC), you must have a GDB compiled for the larger address size. For SPARC, run GDB's configure script with the `--target=sparc64-sun-solaris2.8` option.

12.2.1 Remote GDB and Shared Libraries

It takes some work to figure out how to load symbol tables at the correct offsets for relocatable object modules in GDB. This is done automatically for normal (non-remote) targets, but for the remote target, you have to do it yourself. You need to find out the actual address at which the shared module is mapped in the current context on the simulated machine, and then calculate the offset to use for GDB's **add-symbol-file** command.

To find the addresses of the shared libraries mapped into a process' memory space under Solaris, use the **/usr/proc/bin/pmap pid** command. The start address of the text segment can be obtained from the `Addr` field in the `.text` line of the output from **dump -h file**.

Under Linux, the list of memory mappings can be found in the file `/proc/pid/maps` (plain text format). The `VMA` column of the `.text` line of the output from **objdump -h file** contains the start address of the text segment.

Using these two values, *map address* and *text address*, you should use *map address + text address* as the offset to **add-symbol-file** (it has to be done this way to compensate for how GDB handles symbol loading).

The following example uses a SPARC running Linux (`sim-sh#` denotes the shell in the simulated computer):

```
sim-sh# ps
```

```

PID TTY          TIME CMD
:
461 ttyS0      00:00:00 bash
sim-sh# cat /proc/461/maps
0000000000010000-00000000000060000 r-xp 0000000000000000 08:11 90115  /bin/bash
0000000000006e000-00000000000076000 rwxp 0000000000004e000 08:11 90115  /bin/bash
:
00000000070040000-00000000070138000 r-xp 0000000000000000 08:11 106505 /lib/libc-2.1.3.so
00000000070138000-00000000070140000 ---p 000000000000f8000 08:11 106505 /lib/libc-2.1.3.so
00000000070140000-0000000007014e000 rwxp 000000000000f0000 08:11 106505 /lib/libc-2.1.3.so
:
sim-sh# objdump -h /lib/libc-2.1.3.so

/lib/libc-2.1.3.so:      file format elf32-sparc

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
:
14 .text          000ce338  000000000001e400  000000000001e400  0001e400  2**9

```

From this output, we derive that the bash process with PID 461 has `/lib/libc-2.1.3.so` located at starting address `0x70040000`. The `.text` symbols starts at address `0x1e400`, so if we connect GDB to Simics we have to add the symbols with an offset of $0x70040000 + 0x1e400 = 0x7005e400$. Before running the following commands, we stopped Simics using control-C while it was executing code in the bash process:

```

(gdb) dir ~/glibc-2.1.2/malloc
Source directories searched: /home/joe/glibc-2.1.2/malloc:$cdir:$cwd
(gdb) add-symbol-file libc.so.6 0x7005e400
add symbol table from file "libc.so.6" at
      .text_addr = 0x7005e400
(y or n) y
Reading symbols from libc.so.6...done.
(gdb) target remote localhost:9123
Remote debugging using localhost:9123
__libc_malloc (bytes=0x14) at malloc.c:2691
2691         if (victim == q)
(gdb) next
2693         q = next_bin(q);
(gdb)

```

12.2.2 Using GDB with Hindsight

`gdb-remote` supports an extension to the GDB remote protocol that allows the debugger to control the Hindsight functions in Simics. An unmodified GDB does not know about this ex-

12.2. Using GDB with Simics

tension, however; if you compile GDB yourself, you will need to start from patched sources, available at <https://www.simics.net/pub/> (see section 12.2.3). Some Simics packages come with a pre-compiled GDB with Hindsight support; you can find it in *host/sys/bin/*.

The patch adds a number of Hindsight commands to GDB; they all have the **reverse-**prefix, and are more or less just the reverse version of the standard commands:

reverse-continue

Run in reverse until a breakpoint or watchpoint is hit, or to the point where Hindsight runs out of history.

reverse-next

Run in reverse and stop at the previous source code line. Will skip subfunction calls.

reverse-nexti

Run in reverse and stop at the previous instruction. Will skip subfunction calls.

reverse-step

Run in reverse and stop at the previous source code line. Will enter subfunction calls.

reverse-stepi

Run in reverse and stop at the previous instruction. Will enter subfunction calls.

reverse-finish

Run in reverse till the point where the current function is called.

Normal break- and watchpoints set with **break** and **watch** will also be triggered when running in reverse using **reverse-continue**

A small example of how to use **reverse-next**:

```
22      for (i = 0; i < 10; i++)
(gdb) p i
$2 = 0
(gdb) n
24      c = foo (i) + c;
(gdb) p i
$3 = 1
(gdb) reverse-next
22      for (i = 0; i < 10; i++)
(gdb) p i
$4 = 0
```

The amount of history that Hindsight keeps is limited; it is only possible to reverse back to the point where Hindsight was started. If GDB recognizes that Hindsight has run out of history, it will report an error. Note that this is not a fatal error, and the debugging session can continue, without the possibility to reverse further than to the point where GDB reported the error.

```
(gdb) reverse-continue
```

Continuing.

No more history to reverse further.

```
_start () at start.c:17
```

```
17      {
```

12.2.3 Compiling GDB

If you do not want to (or cannot) use one of the GDB executables in *host/sys/bin/*, you will most likely have to compile GDB from source, even if your system already has GDB installed. The reason for this is that a given GDB executable is specialized both for the architecture of the computer you run it on (host), and the architecture of the computer that runs the programs you want to debug (target). Any GDB already installed on your computer will have target identical to host, but this is often not what you want when your target is a simulated machine.

The first step is to get the GDB source code. You can either get the unmodified GDB from <ftp://ftp.gnu.org/>, or a Hindsight-aware GDB from <https://www.simics.net/pub/>. In either case, the source will be packaged in a `.tar.gz` file.

The second step is to make sure you have all the tools necessary to compile GDB, such as GNU Make and a C compiler. On a Linux or Solaris system, you probably have them already. On Windows, you will have to install Cygwin; get it at <http://www.cygwin.com/>.

That done, unpack and configure GDB like this:

```
~> tar xzfv gdb-6.3.tar.gz
~> cd gdb-6.3
~/gdb-6.3> ./configure --target=powerpc64-elf-linux
```

(On Windows, be sure to enter these commands in the **bash** shell installed as part of Cygwin.)

The **--target** flag to **configure** specifies which target architecture your new GDB binary will be specialized for (in this example, a 64-bit PowerPC). These flags are tabulated in the reference manual's section on **gdb-remote**, and will also be printed on the Simics console when you run **new-gdb-remote**.

```
~/gdb-6.3> make
```

The build process takes a while; when done, it will have left a `gdb` executable in the "gdb" subdirectory. You can execute it directly from that location:

```
~/gdb-6.3> ./gdb/gdb
```

12.3 Symbolic Debugging Using Symtable

As an alternative to `gdb-remote`, Simics comes with some symbolic debugging facilities of its own in the `symtable` module. It is less full-featured than GDB but is easy to use, and it can be scripted in Python.

Note: Not all Simics targets have full `symtable` support implemented. Check the reference manual for more information.

12.3.1 Symtables and Contexts

Each processor in the simulated system has a *current context*, which represents the virtual address space currently visible to code running on the processor. This context is embodied by a **context** object. A context object has various attributes, such as virtual-address breakpoints and symbolic information for the address space (contained in a **symtable** object).

The correctness of the simulation does not depend on contexts in any way; the concept of multiple virtual address spaces is useful for *understanding* the simulated software, but not necessary for just running it. What contexts to create and how to use them is entirely your business; Simics does not care.

By default, each processor has the object **primary-context** as its current context. You may create new contexts and switch between them at any time. This allows you, for example, to maintain separate debugging symbols and breakpoints for different processes in your target machine. When a context is used in this manner (active when and only when a certain simulated process is active), the context is said to *follow* the process.

One handy tool when trying to make a context follow a simulated process is process trackers. A process tracker knows something about the target machine and its operating system—just enough to be able to tell when a given process is active (Simics itself knows nothing about the abstractions—such as processes—implemented by the simulated software). When listening to the haps triggered by the process tracker, switching contexts at the right moment is a breeze.

Simics comes with process trackers for some targets, but far from all. Chapter 21 describes process trackers in more detail, including how to build your own.

12.3.2 Sample Session

Here we inspect a user-space program—the `zsh` shell—running on a 32-bit PowerPC target. Two things are required for this session: a `zsh` binary built with debug info (see section 12.3.5), and its source code.

Start by creating a process tracker:

```
simics> new-linux-process-tracker kernel = ppc32-linux-2.4.17
Using parameters suitable for ppc32 Linux 2.4.17.
New process tracker tracker0 created.
simics> tracker0.add-processor cpu0
simics> new-context-switcher tracker = tracker0
```

New context switcher switcher0 created.

Note that we had to tell the process tracker which kernel we use—or rather, what we have to tell it is the value of a number of numerical parameters:

```
simics> tracker0.status
Status of tracker0 [class linux-process-tracker]
=====

          Processors : cpu0
    Processor type : ppc32

Process tracking parameters:
          ts_comm : 582
          ts_next : 72
    ts_next_relative : 0
          ts_pid : 128
          ts_prev : 76
          ts_state : 0
    ts_thread_struct : 624
```

ppc32-linux-2.4.17 is simply a convenient name for the set of parameters that work with the 32-bit PowerPC Linux 2.4.17 kernel; such predefined parameter sets exist for some of the kernels in the machines shipped by Virtutech. If you want to do process tracking on a kernel for which there is no such predefined parameter set, you will want to look up the process tracker's **autodetect-parameters** command in the Reference Manual.

We also created a context switcher. It handles the rather boring task of listening to the process tracker and actually switching contexts at the right moment, so that one context will follow the process that runs `zsh`. (Context switchers are covered in more detail in section [21.3](#).)

Now create the symbol table and load the symbols. Note that for this to work, the `zsh` binary must have been built with debug info.

```
simics> new-symtable zsh_sym
Created symbol table 'zsh_sym'
zsh_sym set for context primary_context
simics> zsh_sym.load-symbols ~/zsh-4.2.3/Src/zsh
found load segment at 0x10000000
[symtable] Symbols loaded at 0x10000000
```

Tell the context switcher to use a special context for the `zsh` process. Make sure that the new context uses the symbol table:

```
simics> switcher0.track-bin zsh zsh_context
Context 'zsh_context' will be tracking the first process
```


12.3. Symbolic Debugging Using Symtable

that executes the binary 'zsh'.

```
simics> @conf.zsh_context.symtable = conf.zsh_sym
```

We would like to start debugging the program at the beginning of its main function. The symbol table can tell us where that is:

```
simics> psym main
{int (int, char **)} 0x100001f8
simics> whereis (sym main)
in main() at /home/jane/zsh-4.2.3/Src/main.c:92
```

sym is like **psym**, except that it only returns the value, and not its type—which is exactly what other commands are expecting as input.

Let us set a breakpoint at **main**, and let the simulation run:

```
simics> zsh_context.break -x (sym main)
Breakpoint 1 set on address 0x100001f8 with access mode 'x'
1
simics> c
```

Note that as long as you do not execute a binary named “zsh”, you can run whatever program you want without triggering this breakpoint. That is because it is set on the **zsh_context** context, which will not be activated until **zsh** is run:

```
$ ls /
bin  dev  home  lib          mnt  proc  sbin  usr
boot etc  host  lost+found  opt  root  tmp   var
$ sh -c 'echo foo'
foo
$ zsh
```

Now the simulation stops:

```
Code breakpoint 1 reached.
main (argc=0, argv=0x0) at /home/jane/zsh-4.2.3/Src/main.c:92
92      {
[cpu0] v:0x100001f8 p:0x079d41f8 stwu r1,-32(r1)
```

We can single-step through the code:

```
simics> zsh_context.step

93      return (zsh_main(argc, argv));
simics>
```

12.3. Symbolic Debugging Using Symtable

```
zsh_main (argc=1, argv=0x7ffffe14) at /home/jane/zsh-4.2.3/Src/init.c:1205
1205      {
simics>

1206          char **t;
simics>

1209          setlocale(LC_ALL, "");
simics>

1212          init_jobs(argv, environ);
simics>

init_jobs (argv=0x1, envp=0x7ffffe14) at /home/jane/zsh-4.2.3/Src/jobs.c:1465
1465      {
```

(Just pressing Return at the prompt repeats the last stepping command.)

Note how the function `setlocale` was skipped. It is part of the C library linked into `zsh`, which was not compiled with line number information.

We can also examine the contents of variables (note that some C expressions must be quoted to prevent the command-line parser from trying to parse them):

```
simics> psym envp
(char **) 0x7ffffe14
simics> psym "envp[0]"
(char *) 0x7ffffefa "zsh"
```

Looking at the stack, we can see that we have made two function calls that have not returned since we started single-stepping:

```
simics> stack-trace
#0 0x10043914 in init_jobs (argv=0x1, envp=0x7ffffe14)
    at /home/jane/zsh-4.2.3/Src/jobs.c:1465
#1 0x1003d394 in zsh_main (argc=1, argv=0x7ffffe14)
    at /home/jane/zsh-4.2.3/Src/init.c:1219
#2 0x10000220 in main (argc=1, argv=0x7ffffe14)
    at /home/jane/zsh-4.2.3/Src/main.c:93
#3 0x10129830 in __libc_start_main () in zsh
#4 0x0 in ?? ()
```

12.3.3 Source Code Stepping

There are other source code stepping functions besides **step**. **next**, for example, steps to the next source line without descending into function calls like **step** does. (This is exactly what happened when **step** skipped `setlocale`, but **next** will do this with every function call

12.3. Symbolic Debugging Using Symtable

whether or not we have line number information for them.) And **finish** runs the simulation until the current function returns:

```
simics> zsh_context.finish

zsh_main (argc=1, argv=0x7ffffe14) at /home/jane/zsh-4.2.3/Src/init.c:1219
1219      typtab['\0'] |= IMETA;
```

If Hindsight is enabled, all these stepping commands have reverse counterparts: **rstep**, **rnext**, and **uncall**.

Note: The stepping commands expect to step through a single single-threaded process. If the context does not properly follow a single process, or if that process is multi-threaded, they may terminate too soon, or too late, or not at all.

The reason for this is that all but the simplest stepping commands rely on the stack pointer to be well-behaved – in particular, that it keeps pointing to the same stack. The presence of multiple threads – or multiple processes not hidden by a process tracker – breaks this assumption.

12.3.4 Symbolic Breakpoints

We saw earlier how **sym** could be used to set a breakpoint on a function. **pos** can be used to set a breakpoint on a source line:

```
simics> pos jobs.c:1465
268712212
simics> hex (pos jobs.c:1465)
0x10043914
simics> zsh_context.break -x (pos jobs.c:1465)
Breakpoint 26 set on address 0x10043914 with access mode 'x'
26
```

It is also possible to set a breakpoint on data (a watchpoint). The following example sets a data breakpoint on the variable “`argc`”, causing the simulation to stop whenever this variable is read from or written to. The second parameter is the extent of the breakpoint, in bytes.

```
simics> zsh_context.break -r -w (sym "&argc") (sym "sizeof argc")
Breakpoint 27 set on address 0x7ffffd88, length 4 with access mode 'rw'
27
```

See section [12.1](#) for more information about how to use breakpoints.

12.3.5 Reading Debug Information from Binaries

Symbolic information is normally read from file using the `<symtable>.load-symbols` command as in the example above. Currently only ELF binaries can be used, and the debug info must be in the STABS format. Also, the files must be present on the host machine—Simics cannot read directly from the file system of the simulated machine.

Here are some things to think about when preparing a binary for debugging:

- On some platforms, the GCC compiler does not use the STABS format by default. Use the `-gstabs+` option to force STABS (with some GCC extensions) to be used.
- Some versions of the Sun WorkShop (Forte) C compiler do not put the debug information in the final executable, but expect a debugger to read it from the object files directly. This is not supported by Simics, so be sure to use the `-xs` option when compiling.
- If getting sensible stack traces is important, adhere to the target machine's calling and stack frame conventions. In other words, avoid optimizations such as GCC's `-fomit-frame-pointer`.
- Currently, only C is supported, not C++. It is possible to debug programs built from a mixture of C and C++ source, but then only symbols from the C part (and those declared `extern "C"`) will be reliably recognized, for name mangling reasons.
- It is possible to debug dynamically loaded code by specifying the base address of each module when using `load-symbols`, but it is easier to just link the code statically when possible. See section 12.2.1 for how to find the base address on some systems.

12.3.6 Loading Symbols from Alternate Sources

Sometimes it is desirable to read symbols from a source other than a binary file—perhaps all you have is a text file listing the symbols. The `<symtable>.plain-symbols` command reads symbols from a file in the output format of the BSD `nm` command. Example:

```
000000000046b7e0 T iunique
000000000062ba40 B ivector_table
00000000005a6338 D jiffies
```

The hexadecimal number is the symbol value, usually the address. The letter is a type code; for this purpose, D, B, and R are treated as data and anything else as code.

The symbols do not have any C type or line number information associated with them, but you will at least be able to print stack traces and find the location of statically allocated variables.

12.3.7 Multiple Debugging Contexts

The process tracker and context switcher can without problem handle separate contexts (and symbol tables) for two or more processes at once:

12.3. Symbolic Debugging Using Symtable

```
simics> new-symtable encode_sym
Created symbol table 'encode_sym'
encode_sym set for context primary_context
simics> encode_sym.load-symbols ~/sharutils-4.3.80/src/uuencode
[symtable] Symbols loaded at 0x10000000
simics> switcher0.track-bin uuencode context = encode_context
Context 'encode_context' will be tracking the first process
that executes the binary 'uuencode'.
simics> @conf.encode_context.symtable = conf.encode_sym
```

```
simics> new-symtable decode_sym
Created symbol table 'decode_sym'
simics> decode_sym.load-symbols ~/sharutils-4.3.80/src/uudecode
[symtable] Symbols loaded at 0x10000000
simics> switcher0.track-bin uudecode context = decode_context
Context 'decode_context' will be tracking the first process
that executes the binary 'uudecode'.
simics> @conf.decode_context.symtable = conf.decode_sym
```

Here, we have created separate contexts for the programs `uuencode` and `uudecode`, loaded their symbols into two **symtables**, and asked the context switcher to associate these contexts with the first processes that execute “`uuencode`” and “`uudecode`”, respectively.

We would like to step through `uudecode` first:

```
simics> decode_context.step
```

The simulation just runs freely now, waiting for us to reach a source line while **decode_context** is active—which will happen as soon as we start `uudecode`:

```
$ ls -laR / | uuencode - | uudecode | wc
```

```
main (argc=0, argv=0x0) at /home/jane/sharutils-4.3.80/src/uudecode.c:432
432     {
```

We started four programs at once, here. First, `ls` prints a listing of every file on the target machine. This listing is fed to `uuencode`, which encodes it, then feeds the coded result to `uudecode`, which decodes it. The decoded file listing (which is identical to the original listing produced by `ls`) is then fed to `wc`, which counts the number of words in it and prints the result on the terminal.

This description makes it sound like the four programs are run in sequence, one after the other. This is not the case. They all run simultaneously—or rather, since this is a single-processor system, they run interleaved. It works more or less like this:

1. `ls` runs, accumulating output in a buffer. When that buffer is full, it makes a system call that passes the buffered output to the next process, `uuencode`.
2. `uuencode` runs until it has consumed all available input, or until it has to flush its output buffer.
3. `uudecode` runs until it has consumed all available input, or until it has to flush its output buffer.
4. `wc` runs until it has consumed all available input.
5. Repeat until `wc` has finished.

Things may not happen in exactly that order; the only constraint is that a process that is waiting for input cannot be run until some input is available. As long as the amount of data to be passed is large enough to fill the programs' output buffers several times over (and a directory listing of the entire file system should be large enough), execution will alternate between the different programs. So, when we reach the first source line in `uudecode`, `uuencode` should still be running.

Let's first step a few lines in `uudecode`:

```
simics> decode_context.step

433      int opt;
simics>

437      program_name = argv[0];
```

This is just like when we were debugging a single program. Now, let's test our assumption by stepping to the next line in `uuencode`. The `step` command will let the simulation run until we reach a new line *in that context*, so we will either stop when `uuencode` gets to run again, or continue running forever if `uuencode` has already finished.

```
simics> encode_context.step

try_putchar (c=34) at /home/jane/sharutils-4.3.80/src/uuencode.c:130
130      if (putchar (c) == EOF)
```

As expected, `uuencode` was still running. In fact, it was busy outputting text for `uudecode` when it was interrupted so that `uudecode` could be started.

If we step to the next line in `uudecode` again, we see that it has now run to the point where it was blocking, waiting for `uuencode` to produce more output:

```
simics> decode_context.step

read_stdin (inname=0x100464ec "stdin", outname=0x7fffbcf8 "-")
at /home/jane/sharutils-4.3.80/src/uudecode.c:126
```

12.3. Symbolic Debugging Using Symtable

```
126             if (fgets ((char *) buf, sizeof(buf), stdin) == NULL)
```

12.3.8 Scripted Debugging

It is often useful to access data symbolically from Python scripts. Scripts access the debugging facilities using the `symtable` interface and attributes of the `symtable` class. These are documented in the *Simics Reference Manual*.

For instance, here is a short script to print out the contents of one of the linked lists that `zsh` uses. It uses the `eval_sym` function, which takes a C expression and returns a *(type, value)* pair. The expression parsed by `eval_sym` may contain casts, struct member selection and indexing.

```
eval_sym = SIM_get_class_interface("symtable", "symtable").eval_sym

def eval_expr(cpu, expr):
    return eval_sym(cpu, expr, [], 'v')

def ptr_str(typed_val):
    (type, val) = typed_val
    return "((%s)0x%x)" % (type, val)

def print_linklist(list):
    cpu = current_processor()
    ll = eval_expr(cpu, list)
    first = eval_expr(cpu, ptr_str(ll) + "->first")
    l = []
    def print_tail(node):
        type, val = node
        if val == 0:
            return # end of list
        type, val = eval_expr(cpu, ptr_str(node) + "->dat")
        type, val = eval_expr(cpu, ptr_str(("char *", val)))
        l.append(val)
        next = eval_expr(cpu, ptr_str(node) + "->next")
        print_tail(next)
    print_tail(first)
    print l
```

`zsh` uses these lists for lots of things, among them to store the directory stack. After having given the command `pushd` a few times on the `zsh` prompt, we can inspect the directory stack by stopping in the `bin_cd` function and printing the linked list “`dirstack`”:

```
simics> zsh_context.break -x (sym bin_cd)
Breakpoint 1 set on address 0x1000282c with access mode 'x'
1
```

12.3. Symbolic Debugging Using Symtable

```
simics> c
Code breakpoint 1 reached.
bin_cd (nam=0x300001a8 "/usr/bin", argv=0x0, ops=0x101b0000, func=2147481984)
    at /home/jane/zsh-4.2.3/Src/builtin.c:772
772     {
simics> @print_linklist("dirstack")
['/var/tmp', '/tmp', '/usr', '/home', '/sbin', '/bin', '/root']
```


Chapter 13

Profiling Tools

Unlike most profiling tools, which instrument the target source code or object code, Simics can profile a workload non-intrusively. This allows you to profile without disturbing the execution. Simics will profile arbitrary code, including device drivers, dynamically generated code, and code for which you do not have the source. Also unlike most profiling tools, Simics collects profiling data exactly, not by sampling the execution and relying on statistics.

The following sections describe how to make Simics collect profiles of instruction execution and memory reads and writes, and how to examine the collected data.

13.1 Instruction Profiling

Note: Instruction profiling is not yet implemented for targets other than SPARC, PowerPC, x86, and MIPS, and is only available when Simics is started with *-stall*.

Simics can maintain an exact execution profile: every single taken branch is counted, showing you exactly what code was executed and how often branches were taken.

To get started with instruction profiling, type **cpu0.start-instruction-profiling** at the prompt (assuming the machine you are simulating has a processor called **cpu0**). This has the same effect as if you had executed the following two commands:

```
simics> new-branch-recorder cpu0_branch_recorder physical
simics> cpu0.attach-branch-recorder cpu0_branch_recorder
```

These commands create a branch recorder and attach it to the **cpu0** processor. All branches taken by **cpu0** will now be recorded in **cpu0_branch_recorder**.

You need to use these separate commands instead of **start-instruction-profiling** if you want a branch recorder with a particular name, or if you want to record virtual instead of physical addresses.

A branch recorder remembers the source and destination address of each branch, as well as the number of times the branch has been taken. It does *not* remember in which order the branches happened, so it cannot be used to reconstruct an execution trace (if you want that, you have to use the **trace** module, which is slower and generates much more profiling data).

There is however enough information to compute a number of interesting statistics. To get a list of what the branch recorder can do, type:

```
simics> cpu0_branch_recorder.address-profile-info
cpu0_branch_recorder has 6 address profiler views:
View 0: execution count
    64-bit addresses, granularity 4 bytes
View 1: branches from
    64-bit addresses, granularity 4 bytes
View 2: branches to
    64-bit addresses, granularity 4 bytes
View 3: interrupted execution count
    64-bit addresses, granularity 4 bytes
View 4: exceptions from
    64-bit addresses, granularity 4 bytes
View 5: exceptions to
    64-bit addresses, granularity 4 bytes
```

An address profiler is an object whose data can be meaningfully displayed as a count for each address interval. The output of the last command indicates that **cpu0_branch_recorder** is an address profiler with six separate *views*; i.e., there are six separate ways of displaying its data as counts for each four-byte interval.

View 0 is the execution count per instruction, conveniently represented as one count per four-byte address interval since the instructions of the simulated machine (a SPARC in this case) are all four bytes long and aligned on four-byte boundaries. Views 1 and 2 count the number of times the processor has branched from and to each instruction, except when the branch is caused by an exception (or exception return); those branches are counted separately by views 4 and 5. View 3, finally, counts the number of times an instruction was started but not completed because an exception occurred.

When you are done recording for the moment, type:

```
simics> cpu0.detach-branch-recorder cpu0_branch_recorder
```

to detach the branch recorder; it will stop recording new branches, but the already collected branches will remain in the branch recorder (until you type **cpu0_branch_recorder.clean** at the prompt). You can reattach it again at any time:

```
simics> cpu0.attach-branch-recorder cpu0_branch_recorder
```

Section [13.3](#) explains how to access the recorded data.

13.1.1 Virtual Instruction Profiling

Branch recorders can record virtual instead of physical addresses; just say so when you create them:

13.2. Data Profiling

```
simics> new-branch-recorder ls_profile virtual
```

Once it has been created, the new branch recorder object behaves just the same as a physical profiler would, except for one thing: when you want a physical profile, you typically expect the profiler to collect statistics from the whole system, but when you want a virtual profile, you probably are interested in one process (or the kernel) only.

To be consistent with the name of the branch recorder we just created, let us collect a virtual profile of a run of the `ls` program. The problem is to have the branch recorder attached when `ls` is running, and detached when other processes are running. This is the same problem that we had in chapter 12.3, when we wanted a **context** object to be active precisely when a given process was active.

Unfortunately, we cannot use the same convenient tools for branch recorders as we did for contexts, since Simics does not yet support that. We can use the process tracker directly, though, with Python scripts very similar to those in section 21.2:

```
def exec_hap(user_arg, tracker, tid, cpu, binary):
    if binary.endswith("ls"):
        def set_profiler(user_arg, tracker, tid, cpu, active):
            if active:
                cpu.branch_recorders = [conf.ls_profile]
            else:
                cpu.branch_recorders = []
        SIM_hap_add_callback_obj_index("Core_Trackee_Active", tracker,
                                       0, set_profiler, None, tid)
    SIM_hap_add_callback_obj("Core_Trackee_Exec", conf.tracker0,
                             0, exec_hap, None)
```

This script assumes the existence of a process tracker called **tracker0**, and the branch recorder **ls_profile** that we created above.

Note: Remember to use the CLI command **<tracker>.activate** (or, equivalently, call the *activate* function of the `tracker` interface) before trying to use a tracker.

First, it listens for the `Core_Trackee_Exec` hap, which is triggered when any process in the system calls the `exec` system call. When that happens, and the binary being executed is called “`ls`”, the script starts listening for the `Core_Trackee_Active` hap for that process, attaching the branch recorder **ls_profile** to the processor when the process becomes active, and detaching it when the process becomes inactive.

13.2 Data Profiling

In addition to recording branches, Simics can record memory reads and writes by physical address—all of it simultaneously if you like. Start it like this:

```
simics> cpu0.add-memory-profiler read
[cpu0] New read memory profiler added: cpu0_read_mem_prof
```

This creates a new data profiler called **cpu0_read_mem_prof**, which will keep track of all reads performed by **cpu0** until you detach it:

```
simics> cpu0.remove-memory-profiler read
```

Like branch recorders, data profilers retain their recorded data after being detached, and may be reattached later. They may also be attached to another processor, even while they are still attached to the first one.

And, like branch recorders, data profilers are also address profilers. This means that while they may be different under the hood, they present their data the same way:

```
simics> cpu0_read_mem_prof.address-profile-info
cpu0_read_mem_prof has 1 address profiler view:
View 0: data profiler
        64-bit addresses, granularity 32 bytes
```

However, unlike branch recorders, data profilers can only record physical addresses.

Section 13.3 explains how to access the recorded data.

Cache models provide profiling on cache hits and misses. See chapter 18 for more information.

13.3 Examining the Profile

Examining the collected profile is the same for any address profiler, no matter if it is really a data profiler, a branch recorder, or something else. For the examples, we use the branch recorder from section 13.1.

Run a few million instructions or so with the branch profiler attached to the processor to get some data to look at, and then type:

```
simics> cpu0_branch_recorder.address-profile-data
View 0 of cpu0_branch_recorder: execution count
64-bit physical addresses, profiler granularity 4 bytes
Each cell covers 36 address bits (64 Gigabytes).
```

column offsets:

	0x1000000000*	0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
0x0000000000000000: 17000k	
0x0000008000000000:
0x0000010000000000:
0x0000018000000000:

13.3. Examining the Profile

```

0x0000020000000000: . . . . . . . .
0x0000028000000000: . . . . . . . .
0x0000030000000000: . . . . . . . .
0x0000038000000000: . . . . . . . .
0x0000040000000000: . . . . . . . .
0x0000048000000000: . . . . . . . .
0x0000050000000000: . . . . . . . .
0x0000058000000000: . . . . . . . .
0x0000060000000000: . . . . . . . .
0x0000068000000000: . . . . . . . .
0x0000070000000000: . . . . . . . .
0x0000078000000000: . . . . . . . 141

```

16999781 (17000k) counts shown. 0 not shown.

(View 0 is the default, so we did not have to specify it explicitly.) This gives us an overview of the address space. Since we did not specify what address interval we wanted to see, we got the smallest interval that contained all counts. By giving arguments to **address-profile-data**, we can zoom in on the interesting part where almost all the action is. A few orders of magnitude closer, it looks like this:

```

simics> cpu0_branch_recorder.address-profile-data address = 0x1f800000 table-bits = 19
View 0 of cpu0_branch_recorder: execution count
64-bit physical addresses, profiler granularity 4 bytes
Each cell covers 12 address bits (4 kilobytes).

```

column offsets:

	0x1000*	0x0	0x1	0x2	0x3	0x4	0x5	0x6	0x7
0x000000001f800000:	31	760	9472k	204633
0x000000001f808000:	14380	39728	420	1	2751	38226	119215	560668	
0x000000001f810000:	375	301	5979k	1036	
0x000000001f818000:	.	.	390	
0x000000001f820000:	
0x000000001f828000:	
0x000000001f830000:	
0x000000001f838000:	16	129960	
0x000000001f840000:	16048	194576	12930	
0x000000001f848000:	.	1264	152	.	.	18930	12338	174868	
0x000000001f850000:	5736	6
0x000000001f858000:	52	
0x000000001f860000:	
0x000000001f868000:	
0x000000001f870000:	
0x000000001f878000:	

```
16999640 (17000k) counts shown. 141 not shown.
```

Here, we have zoomed in on the 19 bits of address space containing the address 0x1f800000. It is also possible to specify exactly what address interval you are interested in; see the *Simics Reference Manual* or the online help for **address-profile-data**.

If you want to see the current numbers every time you disassemble code, use the command **aprof-views** to select the views you want:

```
simics> cpu0.aprof-views add = cpu0_branch_recorder view = 0
simics> si 5
[cpu0] <v:0x00000000f00066b4> <p:0x000000001f8066b4> 1183160 sllx %o1, 63, %o1
[cpu0] <v:0x00000000f00066b8> <p:0x000000001f8066b8> 1183160 jmp1 [%o7 + 8], %g0
[cpu0] <v:0x00000000f00066bc> <p:0x000000001f8066bc> 1183160 andn %o0, %o1, %o0
[cpu0] <v:0x00000000f0012358> <p:0x000000001f812358> 1181028 cmp %o0, %i0
[cpu0] <v:0x00000000f001235c> <p:0x000000001f81235c> 1181028 bcs,pt %xcc, 0xf0012350
```

Note that since the disassembly printed by commands such as **c** and **si** are the instruction just about to execute, the statistics reflect the way things were immediately before that instruction executed.

For most real-world profiling tasks, the above methods of looking at profile data are inadequate. Unfortunately, Simics does not currently offer a way to display large sets of profiling data in a user-friendly way; what it does offer is a few primitive operations (such as accessing individual counters and summing all counters in an interval) that users can call from their own scripts. See the Profiling API section in the *Simics Reference Manual* for details.

Part V

Advanced Simics Usage

Chapter 14

Startup Options

This chapter describes the most important command-line options accepted by Simics. A complete list can be found in the *Simics Reference Manual*.

The common way of starting Simics from the command line is to run the `./simics` script in your workspace directory. Simics can also be started by running a similar start script in the `[simics]/bin` directory.

```
joe@computer:[workspace]$ ./simics targets/x86-440bx/enterprise-common.simics
```

This will run the `enterprise-common.simics` script, which will start the *Enterprise* machine with its default configuration.

14.1 Simulation Modes

The *simulation mode* (or *execution mode*) is used to control the availability of certain features: instruction and data profiling, memory timing, and the micro-architectural interface. To allow for maximum performance when these features are not used, they are compiled as separate processor implementations. By default, Simics uses the fastest possible mode (“normal”).

To select a simulation mode, the corresponding command line flag (`-stall` or `-ma`) is passed to Simics. For example, to select “Stall mode”:

```
$ cd [workspace]
$ ./simics -stall ...
```

Simics does not support switching processor implementations at runtime. To change from one mode to another, use Simics’s checkpoint feature to save the simulation state and restart Simics in the new mode from the checkpoint:

```
simics> write-configuration at-workload-start
simics> quit
```

and then start from the new checkpoint in the new simulation mode:

```
$ ./simics -stall -c at-workload-start
```

14.1.1 Normal mode

Normal mode is the fastest execution mode, and is optimized for emulation-style usage. The following features are not available in Normal mode:

- Instruction and data profiling is not available.
- Attaching timing-models and snoop devices to memory spaces is not supported. This means that the tracing and cache modules do not work correctly, and that stalling is not possible.
- The micro-architectural interface is not available.

14.1.2 Memory Timing with -stall

Stall mode supports stalling and instruction/data profiling.

Note that not all Simics processor models implement a complete support for stalling and profiling. Refer to chapter 13 (Profiling) and chapter 16 (Stalling) for more information.

14.1.3 Micro Architectural Simulation with -ma

The micro-architecture mode selects in implementation with support for micro-architectural simulation. For performance reasons, this mode is usually only used for those subsets of the workload for which micro-architectural modeling is relevant. Typically, the system is booted in normal mode until the workload starts. A checkpoint is then taken, and Simics is restarted in micro-architecture mode from that checkpoint.

Changing to micro-architectural mode usually requires additional configuration information (like a processor timing model for example) before the simulation can be run. See chapter 17 and *Simics Micro-Architectural Interface* document for more information.

14.2 Common Options

Below is a list of the most common options used when starting Simics. All possible startup options are listed in the *Simics Reference Manual*.

-c file

Reads a specified configuration file. This is equivalent to running the command **read-configuration file** after starting Simics.

-central <addr[:port]>[:<file>]

Connects to an existing Simics instance with a Simics Central server running. This option takes an IP address and an optional port number (the default is 1909), or a Unix domain socket (a file name). If no port number is supplied, and the local host's address is supplied, Simics will try to connect to the default Unix domain socket.

14.2. *Common Options*

-h

Prints a list of the possible startup flags along with a short description.

-v

Print the Simics version number.

Chapter 15

The Command Line Interface

The Simics Command Line Interface (CLI) is an advanced text based user interface with built-in help system, context sensitive tab-completion, and scripting support (both built-in and using Python).

15.1 Invoking Commands

Commands are invoked by typing them at the command line followed by their arguments. The synopsis part of a command documentation (you can see many examples in the *Simics Reference Manual*) explains how to call a command. Here are two examples:

SYNOPSIS

command1 *-x -y -small* [cpu-name] address (size|name)

SYNOPSIS

command2 files ...

Arguments starting with a hyphen are flags and are always optional. Flags can be more than one character long so it is *not* possible to write *-xy* for *-x -y*. The order of the flags is not significant and they can appear anywhere in the argument list.

Arguments enclosed within square brackets are optional; in the example above, it is not necessary to specify *cpu-name*. *address*, on the other hand, is required. The last argument to **command1** is either a size or a name, but not both. Such arguments are called *polyvalues* and can be of different types. Size and name are called sub-arguments.

If an argument is followed by three dots as the file argument in **command2** it indicates that the argument can be repeated one or more times.

The type of the arguments, e.g., if they are integers or strings, should be evident from their names. For example *size* should be an integer and *name* a string if not documented otherwise.

Integers are written as a sequence of digits beginning with an optional minus character for negative numbers. Hexadecimal numbers can be written by prefixing them with 0x, octal numbers with 0o, and binary numbers with 0b. Integers may contain “_” characters to make them easier to read. They are ignored during parsing. For example:

```
simics> 170_000
170000
simics> 0xFFFF_C700
4294952704
```

Strings are written as is or within double quotes if they contain spaces or begin with a non-letter.

Here are some possible invocations of the commands above:

```
simics> command1 -small cpu0 0x7fff_c000 14 -y

simics> command1 0x7fffc000 foo

simics> command1 -x "Pentium 4" 0x7fff_c000 -8

simics> command2 "/tmp/txt" "../bootdisk" floppy
```

In the first example `cpu-name` is passed as the string `cpu0` and `size` as the integer 14. In the second invocation `cpu-name` has been omitted and `name` is set to the string `foo`. The third example illustrated the use of a string containing a space. In all **command1** examples the address is set to the hexadecimal value `0x7fffc000`. **command2** takes a list of at least 1 string.

A few commonly used commands have aliases. For example, it is possible to write `c` for **continue** and `si` for **step-instruction** for example. Command aliases are documented with their corresponding command in the *Simics Reference Manual*.

15.1.1 How are Arguments Resolved?

Simics tries to match the provided arguments in same order as they appear in the synopsis. If the type of the next argument is identical to what is typed at the command line the argument will match. If there is a mismatch and the argument is optional, the argument will be skipped and the next argument will be matched, and so on. If a mismatching argument is not optional, the interpreter will fail and explain what it expected. For polyvalues, the argument will match if one of its sub-arguments matches.

There are situations however when this method is not sufficient. For example, when two arguments both have the same type and are optional, there is no way to know which argument to match if only one is given. This is resolved by naming the arguments: `arg-name=value`. For example **command1** in the example above can be invoked like this:

```
simics> command1 size=32 -y address = 0xf000 -small cpu-name=cpu0
```

Thus there is no ambiguity in what is meant and in fact this is the only way to specify a polyvalue with sub-arguments of the same type. Note also that named arguments can be placed in any order.

15.1. Invoking Commands

15.1.2 Namespace Commands

Configuration objects (such as devices or CPUs) that define user commands usually place them in a separate namespace. The namespace is the name of the object. Interfaces may also define commands, in which case all objects implementing these interfaces will inherit of the commands in their own namespace.

Namespace commands are invoked by typing the name of the object, followed by a dot and the command name: *object.command*, e.g.,

```
simics> cache0.print-status
```

All namespace commands are listed in the *Simics Reference Manual* under the class or interface they belong to.

15.1.3 Expressions

The CLI allows expressions to be evaluated, for example:

```
print -x 2*(0x3e + %g7) + %pc
```

The precedence order of the operators is as follows (highest first):

\$	read Simics variable
%	get register value
[]	variable indexing
->	attribute access
pow	power of
~	bitwise not
*, /	multiplication, division
+, -	addition, subtraction
<<, >>	left, right shift
&	bitwise and
^	bitwise xor
	bitwise or
<, <=, ==, !=, >=, >	comparison
not	boolean not
and	boolean and
or	boolean or

Parentheses can be used to override the priorities. Commands which return values can also be used in expressions if they are enclosed within parentheses:

```
print -x (cpu0.read-reg g7)
```

Values can be saved in variables for later use. You set a variable by simply giving an assignment command such as **\$var = 15**.

15.1.4 Interrupting Commands

Any command which causes the simulation to advance can be interrupted by typing control-C. The simulator will gracefully stop and prompt for a new command. If Simics hangs for some reason, possibly due to some internal error, you can usually force a return to the command line by pressing control-C two or more times in a row.

Note: Pressing control-C several times may damage some internal state in the simulator so should be used only in last resort.

15.2 Tab Completion

The command line interface has a tab-completion facility (if the `readhist` helper program has been started—Simics should do that automatically when supported). It works not only on commands but on their arguments as well. The philosophy is that the user should be able to press the tab key when uncertain about what to type, and Simics should fill in the text or list alternatives.

For example `com<tab>` will expand to the command beginning with `com` or list all commands with that prefix if there are several. Similarly, `disassemble <tab>` will display all arguments available for the command. In this case Simics will write:

```
address =      count =      cpu-name =
```

to indicate that these alternatives for arguments exists. Typing `disassemble cp<tab>` will expand to `disassemble cpu-name =` and a further tab will fill in the name of the CPU that is defined (or list all of them).

15.3 Help System

The most useful Simics commands are grouped into categories. To list these categories, just type `help` at the command prompt. The list should look like this:

```
simics> help
[...]
To get you started, here is a list of command categories:
Breakpoints
Changing Simulated State
Command-Line Interface
Configuration
Debugging
Disk
Ethernet
Execution
Files and Directories
```


15.3. Help System

```
Haps
Help
Inspecting Simulated State
Light Edition
Logging
Memory
Modules
Output
Profiling
Python
Real Network
Registers
Remote Access
Simics Central
Simics Search Path
Speed
Symbolic Debugging
Test
Tracing
[...]
```

Note that since Simics's configuration can change between sessions and even dynamically through loading modules, the commands and command categories may look different.

Type **help category** for a list of commands, e.g., **help "Changing Simulated State"** will list all commands belonging to that category:

```
simics> help "Changing Simulated State"
Commands available in category Changing Simulated State
set-pc                set the current processor's program counter
set                   set physical address to specified value
<image>.set           set bytes in image to specified value
load-file             load file into memory
load-binary           load binary (executable) file into memory
<processor>.write-reg write to register
write-reg             write to register
penable              switch processor on
```

Type **help command** to print the documentation for a specific command.

The **help** command can do much more than printing command documentation: it gives you access to nearly all Simics documentation on commands, classes, modules, interfaces, API types and functions, haps and more according to the configuration loaded in the simulator. All documentation is also available in the *Simics Reference Manual*.

Here are some more examples of usage of the **help** command:

```
simics> help ptime
```

```
[... ptime command documentation ...]

simics> help cpu0.disassemble
[... <processor>.disassemble command documentation ...]

simics> help <processor>.disassemble
[... <processor>.disassemble command documentation ...]

simics> help cpu0
[... <ultrasparc-iii-i> class documentation ...]

simics> help ultrasparc-iii-i
[... <ultrasparc-iii-i> class documentation ...]

simics> help processor
[... <processor> interface documentation ...]

simics> help cpu0.reorder_buffer_size
[... <ultrasparc-iii-i>.reorder_buffer_size attribute documentation ...]

simics> help ultrasparc-iii-i.windowed_registers
[... <ultrasparc-iii-i>.windowed_registers attribute documentation ...]

simics> help Core_Exception
[... Core_Exception hap documentation ...]

simics> help SIM_get_mem_op_type
[... SIM_get_mem_op_type() function declaration ...]

simics> help sparc-u3-turbo
[... sparc-u3-turbo module documentation ...]
```

When a name matches several help topics (for example, a command and an attribute, or a module and a class), **help** will print out the first topic coming in this order: command categories, commands, classes, interfaces, haps, modules, attributes, API functions and symbols. It will also inform you at the end of the documentation output that other topics were matching your search:

```
simics> help cheerio-hme
[... cheerio-hme class documentation ...]
```

Note that your request also matched other topics:
 module:cheerio-hme

15.4. Simics's Search Path

If you type **help module:cheerio-hme**, the module documentation will be printed instead:

```
simics> help module:cheerio-hme
[... cheerio-hme module documentation ...]
```

You can use specifiers like `module:` or `class:` at any time. It will also allow the **help** command to provide you better tab-completion, since only items in the selected category of documentation will be proposed. The following specifiers are available: `object:`, `class:`, `command:`, `attribute:`, `interface:`, `module:`, `api:`, `hap:` and `category:`.

Note: By default, **help** does not propose tab-completion for modules and API symbols, because they tend not to be the most searched for and would clutter the tab-completion propositions unnecessarily. You can get tab-completion for those by specifying `module:` or `api:` in front of what you are looking for.

The **apropos** command can search for keywords in the documentation provided by **help**. Type **apropos keyword** to get a list of all documentation topics matching this keyword.

```
simics> apropos step
The text 'step' appears in the documentation
for the following items:

Command      <processor>.cycle-break-absolute
Command      <processor>.step-break
Command      <processor>.step-break-absolute
Command      cycle-break-absolute
Command      log-setup
Command      print-event-queue
[...]
Attribute    <ultrasparc-iii>.steps
Hap          Core_Step_Count

simics>
```

15.4 Simics's Search Path

Many Simics commands will look up files based on the current directory. This may be impractical when writing scripts or building new configurations, so Simics provides two features to ease directory handling:

- Simics recognizes some special path markers that are translated before being used:

%simics%

Translated to the current Simics installation directory, so that `%simics%/scripts/`

`foo.simics` will be translated to, for example, `/home/joe/simics/scripts/foo.simics`. Note that if you change the version of Simics you are using, `%simics%` will change as well, so you should use it to refer only to files that you know are present in all Simics versions..

%script%

Translated to the directory where the currently running script is located. A possible usage is to let a script call another one in the same directory, independently of what the current directory is.

For example, if the directory `/home/joe/scripts/` contains the scripts `foo.simics` and `bar.simics`, even if the user uses Simics with `/home/joe/workspace/` as current directory, it will be possible for `foo.simics` to call `bar.simics` by issuing the command:

```
run-command-file %script%/bar.simics
```

`%simics%` and `%script%` are always translated to absolute paths, so they never interact with the next feature, called *Simics's search path*.

- Simics has a list of paths called *Simics's search path* where files will be looked up when using some specific commands (among others, **load-binary**, **load-file**, **run-command-file**, and **run-python-file**). The file is first looked up in the current directory, then in all entries of Simics's search path, in order.

Let us assume for example that Simics's search path contains `/home/joe/scripts/` and that the current directory is `/home/joe/workspace`. If the command:

```
simics> run-command-file test/start-test.simics
```

is issued, Simics will look for the following files to run:

1. `/home/joe/workspace/test/start-test.simics`
2. `/home/joe/scripts/test/start-test.simics`

Simics's search path can be manipulated using the **add-directory**, **clear-directories** and **list-directories** commands. Simics's search path is also used when looking for images files belonging to checkpoints or new configuration. This is described in section [6.2.2](#).

Note: Although the Simics search path is saved in the `sim` object in checkpoints, allowing image files that were found through it to be opened again by the checkpoint, it is not available until the object creation phase. Module initialization code should not rely on the Simics path since that code is run before the `simobject` object from the checkpoint has been created.

15.5 Using the Pipe Command

The **pipe** command, which only is available on Unix hosts, lets you send the output of a Simics command to a shell command through a pipe:

```
simics> pipe "help" "grep Tracing"
```

This will run **help** (which lists all Simics commands categories) and send its output to the standard input of the **grep trace** process. **grep** will discard all lines not containing “Tracing” and forward the rest to its standard output, which will be printed on the Simics terminal.

The **pipe** command can be used to send all the output of a command to a file:

```
simics> pipe "stepi 1000" "cat > trace.txt"
```

Or you can use it to view information using the shell command **more**:

```
simics> pipe "pregs -all" more
```

Note that you have to enclose both the Simics command (the first argument) and the shell command (the second argument) in double quotes if they contain whitespace or other non-letter characters.

A related command is **!**, which runs a shell command from the Simics prompt:

```
simics> !uname -a
```

This will run the **uname -a** shell command and print its output in the Simics console.

Chapter 16

Memory Transactions

Both CPUs and devices can initiate memory transactions in Simics. For a CPU, the requested virtual address is first translated by the MMU, and an access is performed using the physical address obtained. Device transactions usually skip the translation phase, but are otherwise handled in the same way.

In order to know what a given physical address corresponds to, Simics uses a concept called *memory-space*. A memory-space maps a physical address to an object that can accept a memory transaction, like RAM, a flash-memory or a device. An access performed in a memory-space is automatically propagated to the right *target* (device or memory).

Memory-spaces can also contain other memory-spaces as targets which will in-turn map the physical address to a specific object, thus creating a hierarchical organization of memory-spaces. This is often used to separate different types of mapping, or to implement architecture specific differences.

For example, the memory mappings of a one-processor Serengeti system (in this case `sarek-common.simics`) at boot time is based on two memory-spaces: **phys_io0** (the I/O space) and **phys_mem0** (the memory space). They contain the following mappings:

```
simics> phys_mem0.map
```

base	object	fn	offs	length
0x0000000000000000	memory	0	0x0	0x10000000
0x0000007fff07ffff0	hfs	0	0x0	0x10


```
simics> phys_io0.map
```

base	object	fn	offs	length
0x0000040000400000	mmu0	0	0x0	0x48
0x000004000c000000	schizo24	0	0x0	0x800000
0x000004000c800000	schizo25	0	0x0	0x800000
0x000007fff0000000	cpuprom	0	0x0	0xbd3b0
0x000007fff0102000	fpost_code	0	0x0	0x2000
0x000007fff0104000	fpost_data	0	0x0	0x2000
0x000007fff0800060	empty0	0	0x60	0x10
0x000007fff091e000	empty1	0	0x0	0x120

Simics allows you to observe and sometimes modify the behavior of the transactions that go through the memory system. The rest of this chapter will present the basic concepts behind Simics's memory system and how to interact with it.

16.1 Observing Memory Transactions

Memory-spaces provide a *memory hierarchy interface* for observing and modifying memory transactions passing through them. This interface is in fact composed of two different interfaces acting at different phases of a memory transaction execution:

- The `timing-model` interface provides access to a transaction *before* it has been executed (i.e., it has just arrived at the memory-space).

To connect an object to the timing model interface, just set the `timing_model` attribute of the corresponding memory-space with the value of the object you want to connect:

```
simics> @conf.phys_mem0.timing_model = conf.listening_object
```

The timing model interface can also be used to change the timing and the execution of a memory transaction, as well as to modify the value of a store going to memory. This is described in more detail in section 16.2 and in the *Simics Programming Guide*.

- The `snoop-memory` interface provides access to a transaction *after* it has been executed.

Connecting an object to the snoop memory interface is done in the same way as for the timing model interface:

```
simics> @conf.phys_mem0.snoop_device = conf.listening_object
```

The `trace` module, for example, automatically connects itself (or rather one of its objects) to this interface when a trace is started. The advantage of using this interface is that the value of load operations is accessible, since the operation has already been performed.

The snoop memory interface can also be used to modify the value of a load operation. This is described in more detail in the *Simics Programming Guide*.

Note: Both interfaces can be used simultaneously, even by the same object. This property is used by the `trace` module, which is in fact connected both to the `timing-model` and the `snoop-memory` interfaces. The reason for this double connection is explained later in this chapter.

16.2 Stalling Memory Transactions

When an object attached to the *timing-model* interface receives a memory-transaction, it is allowed to modify the timing of the transaction by returning a *stall time* (as a number of processor cycles). Upon receiving a non-zero stall time, Simics will block the transaction for that number of processor cycles before executing it.

This behavior is a key to modeling caches and memory hierarchies in Simics, particularly in multi-processor simulations. A complete description of cache simulation in Simics is available in chapter 18. The *Simics Programming Guide* explains how to make your own timing-model object to observe and stall memory transactions, and provides more information about the different types of stalling available.

To be able to stall transactions, the simulation must be started with the *-stall* flag. This will tell Simics to set up the processors so that stalling is allowed. It is otherwise disabled for performance reasons.

Stalling a transaction is not always possible, depending on the processor model you are using in the simulation:

- UltraSPARC processors can stall both data accesses and instruction fetches.
- x86, PowerPC and MIPS processors can stall data accesses.

Note: As mentioned above, a transaction may go through several memory-spaces in hierarchical order before being executed. Each of these memory-spaces may have a timing-model connected to them. However, if the transaction is stalled by one timing model, the other timing models may see the transaction being reissued before it is executed.

16.3 Observing Instruction Fetches

For performance reasons, instruction fetches are not sent to the memory hierarchy interface by default. You can activate instruction fetches for each processor by using the command `<cpu>.instruction-fetch-mode`. It can take several values:

no-instruction-fetch

No instruction fetches are sent to the memory hierarchy interface.

instruction-cache-access-trace

An instruction fetch is sent every time a different cache line is accessed by the processor. The size of the cache line is set by the processor attribute *instruction-fetch-line-size*. This option is available for UltraSPARC processors.

This option is meant to be used for cache simulation where successive accesses to the same cache line do not modify the cache state.

instruction-fetch-trace

All instruction fetches are sent to the memory hierarchy interface. This option is often

implemented as *instruction-cache-access-trace* with a line size equal to the size of one instruction. This option is available for UltraSPARC and x86 processors.

This option is meant to provide a complete trace of fetch transactions.

16.4 Simulator Translation Cache (STC)

In order to improve the speed of the simulation, Simics does not perform all accesses through the memory spaces. The Simulator Translation Caches (STCs) try to serve most memory operations directly by caching relevant information. In particular, an STC is intended to contain the following:

- The current logical-to-physical translation for the address;
- A count of number of accesses to the address.

The general idea is that the STC will contain information about “harmless” memory addresses, i.e., addresses where an access would not cause any device state change or side-effect. A particular memory address is mapped by the STC only if:

- The given logical-to-physical mapping is valid.
- An access would not affect the MMU (TLB) state.
- There are no breakpoints, callbacks, etc associated with the address.

The contents of the STCs can be flushed at any time, so models using them to improve speed can not rely on a specific address being cached. They can however let the STCs cache addresses when further accesses to these addresses do not change the state of the model (this is used by cache simulation with **g-cache**; see chapter 18).

The STCs are activated by default. They can be turned on or off at the command prompt, using the **stc-enable/disable** functions. An object connected to the `timing model` interface can also mark a memory transaction so that it won’t be cached by the STCs. For example, the **trace** module uses that method to ensure that no memory transaction will be cached, so that the trace will be complete.

Note that since transactions are inserted into the STCs when they are executed, only objects connected to the *timing model* interface can influence the STCs’ behavior. The *Simics Programming Guide* provides a complete description of the changes authorized on a memory transaction when using the memory hierarchy interface.

16.5 Summary of Simics Memory System

To summarize the different concepts introduced in this chapter, here is a description of the path followed by a processor transaction through Simics memory system.

1. The CPU executes a load instruction.
2. A memory transaction is created.

16.5. Summary of Simics Memory System

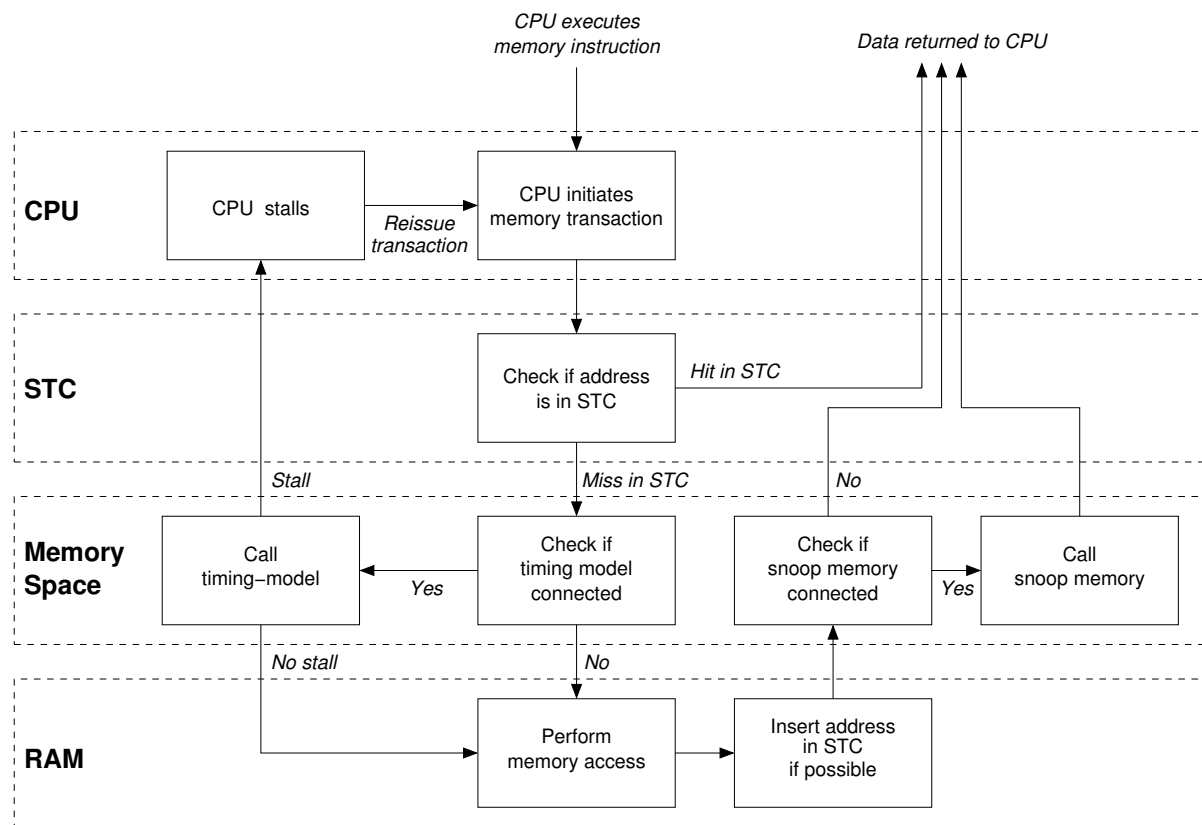


Figure 16.1: Transaction Path through Simics Memory System

3. If the address is in the STC, the data is read and returned to the CPU using the cached information.
4. If the address is not in the STC, the transaction is passed along to the CPU memory-space.
5. If a timing-model is connected to the memory-space, it receives the transaction.
 - (a) If the timing model returns a non-zero stalling time, the processor is stalled and the transaction will be reissued when the stall time is finished.
 - (b) If the timing model return a zero stall time, the memory-space is free to execute the transaction.
6. The memory-space determines the target object (in this example, a RAM object).
7. The RAM object receives the transactions and executes it.
8. If possible, the transaction is inserted in the STC.
9. If a snoop-memory is connected to the memory-space, it receives the transaction.
10. The transaction is returned to the CPU with the correct data.

Store operations works in the same way, but no data is returned to the CPU.

Note: Simics's memory system is more complex than what is presented here, but from the point of view of a user timing-model or snoop-memory, this diagram explains correctly at which point the main different events happen.

Chapter 17

Understanding Simics Timing

This chapter provides an overview of the mechanisms behind Simics simulation system, and gives more formal definitions to some terms and concepts used in previous chapters.

17.1 Events

Simics is an *event driven simulator* with a maximum time resolution of a *clock cycle*. In a single-processor system, the length (in seconds) of a clock cycle is easily defined as 1 divided by the processor frequency set by the user. As described later in this chapter, Simics can also handle multiprocessor systems with different clock frequencies, but we will focus on single-processor systems for the rest of this section.

As Simics makes the simulated time progress, cycle by cycle, events are triggered and executed. Events include device interrupts, internal state updates, as well as *step* executions. A *step* is the unit in which Simics divides the execution of a flow of instructions going through a processor; it is defined as the execution of an instruction, an instruction resulting in an exception, or an interrupt. Steps are by far the most common events.

Steps and cycles are fundamental Simics concepts. Simics exposes two types of events to users: events linked to a specific step, and events linked to a specific cycle. Step events are useful for debugging (execute 1 step at a time, stop after executing 3 steps, etc.), whereas time events are rather independent from the flow of execution (sector read operation on the hard disk will finish in 1 ms).

For each executed cycle, events are triggered in the following order:

1. All events posted for this specific *cycle*, except step execution.
2. For each step scheduled for execution on this cycle:
 - (a) All events scheduled for that specific *step*.
 - (b) A step execution.

Events belonging to the same category are executed in FIFO order: posted first is executed first.

17.2 Instruction Execution Timing

Simics in-order

In the default model, the execution of a step takes no time by itself, and steps are run in program order. This is called the Simics *in-order* model. It implements the basic instruction set abstraction that instructions execute discretely and sequentially. This minimalistic approach makes simulation fast but does not attempt to model execution timing in any way.

Normally one step is executed every cycle, so that the step and cycle counts are the same. See the section **Changing the Step Rate** for how to change this.

Stalling

The in-order model can be extended by adding *timing models* to control the timing of memory operations, typically using the memory hierarchy interface described in chapter 16. When timing models are introduced, steps are no longer atomic operations taking no time. A step performing a memory operation (whether an instruction fetch or a data transaction) can stall for a number of cycles. Cycle events are performed during the stalling period as time goes forward. Step events are performed just before the step execution, as in the default model. Note that in this mode, Simics still executes one step at a time, with varying timing for each step, so the simulation is still performing an in-order execution of the instruction flow. The basic step rate can also be changed; see the section **Changing the Step Rate** below.

Simics MAI

Simics also has an extended timing control mode called *Micro-architectural Interface* (MAI). When running in MAI mode, step execution is completely under user control. Cycle events are still executed one cycle at a time, but no step execution is performed unless a user *processor timing model* is driving it. This user timing model is usually called every cycle to make instructions advance in the processor pipeline, stall, allocate resources, etc. Whenever an instruction is committed by the processor timing model, Simics increases the step counter, executes all step events related to this step and finally commits the step execution. From the point of view of event handling, Simics MAI is similar to the standard model, but the number of steps executed (or, in this case, committed) per cycle is entirely under user control.

Simics MAI transcends the standard model by providing an infrastructure for parallelized and speculative execution, allowing a much more accurate timing simulation of complex processors. Simics MAI mode is also called Simics *out-of-order* mode. A complete description of Simics MAI is available in the *Simics Micro-Architectural Interface* document. Note that Simics MAI is only available for the UltraSPARC and x86/x86-64 processor models.

Choosing an Execution Mode

Choosing an execution mode is matter of trade-off between performance and accuracy. As the timing control becomes more refined, the simulation slows down. Simics allows you to choose the execution mode dynamically (using checkpoints), so it is possible to use a simple

17.2. Instruction Execution Timing

model to reach interesting parts of the simulation quickly, then switch to a more complex model.

Changing the Step Rate

The *step rate* is the number of steps executed each cycle, disregarding any stalling. It is expressed as the quotient q/p . By default, $p = q = 1$; this schedules one step to run in each cycle. This can be changed by setting the *step_rate* processor attribute to $[q, p, r]$ where the last r parameter is a remainder to keep track of the relative phase of the cycle and step counters; set it to zero if you are not interested in sub-cycle accuracy. For example,

```
@conf.cpu0.step_rate = [3, 4, 0]
```

will set the step rate of **cpu0** to 3/4; that is, three steps every four cycles.

If $q < p$, then some cycles will execute no step at all; if $q > p$, then some cycles will execute more than one step. The step rate parameters are currently limited to $1 \leq p \leq 128$ with $p = 2^k$ for some integer k , and $1 \leq q \leq 128$.

Setting a non-unity step rate can be used to better approximate the timing of a target machine averaging more or less than one instruction per cycle. It can also be used to compensate for Simics running instructions slower than actual hardware when it is desirable to have the simulated time match real time; specifying a lower step rate will cause simulated time go faster.

The step rate is sometimes called IPC (instructions per cycle), and its inverse, the *cycle rate*, may be called CPI (cycles per instruction). The actual rates will depend on how many extra cycles are added by stalling.

Let us look at an example using a single Ebony card. We will first run 1 million steps with the default settings:

```
+-----+      Copyright 1998-2005 by Virtutech, All Rights Reserved
| Virtutech |      Version:
| Simics   |      Build:
+-----+
www.simics.com      "Virtutech" and "Simics" are trademarks of Virtutech AB

simics> c 1000000
[cpu0] v:0xffff8a610 p:0x1ffff8a610 mftbu r5
simics> ptime
processor          steps          cycles      time [s]
cpu0               1000000        1000000      0.010
```

The processor has run 1 million steps, taking 1 million cycles to execute them. Let us set the cycle rate to the value mentioned above, 3 steps for every 4 cycles:

```
simics> @conf.cpu0.step_rate = [3, 4, 0]
simics> cb 1200000
```

```

simics> c
Caught time breakpoint
[cpu0] v:0xffff8a634 p:0x1ffff8a634 bl 0xffff8a608
simics> ptime
processor          steps          cycles    time [s]
cpu0              1900000        2200000    0.022
simics>

```

When running the next 1.2 million cycles, the processor executes only 900000 steps, which corresponds to the 3/4 rate that we configured.

Suspending Time or Execution

It is possible to set the step rate to infinity, or equivalently, to suspend simulated time while executing steps. This is done by setting the *step_per_cycle_mode* processor attribute to one of the following values:

"constant"

Steps are executed at the constant and finite rate specified in the *step_rate* attribute

"infinite"

Steps are executed with no progress in simulated time

While time is suspended, the cycle counter does not advance, nor are any time events run. To the simulated machine this appears as if all instructions are run infinitely fast.

Conversely, it is possible to set the step rate to zero, thus suspending execution while letting simulated time pass. This can be done by stalling the processor for a finite time (see **Stalling** above) or by *disabling* the processor for an indefinite time. Disabling and re-enabling processors is done with the commands **<processor>.enable** and **<processor>.disable**.

Using the same example as above, we set the step per cycle mode to "infinite" to prevent the simulated time from advancing:

```

simics> @conf.cpu0.step_per_cycle_mode = "infinite"
simics> c 1000000
[cpu0] v:0xffff8a614 p:0x1ffff8a614 cmpw r3,r5
simics> ptime
processor          steps          cycles    time [s]
cpu0              1000000          0        0.000
simics>

```

The processor has executed 1 million steps but the simulated time has not advanced. Note that setting this mode would probably prevent a machine like Ebony from booting since many hardware events (like interrupts) are time-based.

17.3 Multiprocessor Simulation

Simics can model systems with several processors, each with their own clock frequency. In this case the definition of how long a cycle is becomes processor-dependent. Ideally, Simics would make time progress and execute one cycle at a time, scheduling processors according to their frequency. However, perfect synchronization is exceedingly slow, so Simics *serializes* execution to improve performance.

Simics does this by dividing time into segments and serializing the execution of separate processors within a segment. The length of these segments is referred to as the *quantum* and is specified in seconds (this is similar to the way operating systems implement multitasking on a single-processor machine: each process is given access to the processor and runs for a certain time quantum). The processors are scheduled in a round-robin fashion, and when a particular processor *P* has finished its quantum, all other processors will finish their quanta before execution returns to *P*. The length of the time quantum can be set by using the command **cpu-switch-time**. The argument to **cpu-switch-time** is specified in cycles (referring to the first processor in the system) rather than absolute time.

As in the single-processor case, instruction execution and latency are defined with execution modes and timing interfaces. Simics does not define the order in which the processors are serialized, which means that if causality is to be preserved, processor-to-processor communications must have a minimum latency of one quantum. Another consequence of serializing the execution is that Simics will maintain strict sequential consistency. However, through careful use of the memory hierarchy interface, the user can choose to simulate other consistency models.

As an example, consider a dual-processor system where the first processor runs at 4 MHz and the second at 1 MHz. Setting **cpu-switch-time** to 10 will give a quantum of 2.5 simulated microseconds. During each quantum, the first processor will execute 10 steps, and the second 2 or 3 steps, not necessarily in that order. Breakpoints do not affect this schedule, so that interaction remains non-intrusive.

Note that if you are single-stepping (**step-instruction**) on a processor *P*, which has just executed the last cycle of a quantum, the next single-step will cause all other processors to advance an entire quantum and then *P* will stop after one step. This behavior makes it convenient to follow the execution of instructions on a particular processor. You can use the **<processor>.ptime** command to see the flow of time on each particular processor in the simulated machine.

For a multi-processor simulation to run efficiently, the quantum should not be set too low, since a CPU switch causes simulator overhead. It should not be set below 10, and should preferably be set to 50 or higher. The default value is 1000. For a perfectly synchronized simulation, set the switch time to 1 (which will give a very slow simulation but is useful for detailed cache studies, for example). Note that all of the above remains essentially the same when running a distributed simulation (see next section).

Time events in Simics are executed when the processor on which they were posted run the triggering cycle during its quantum. However, it is possible to post *synchronizing* time events that will ensure that all processors have the same local time when the event is executed, independently of the time quantum. Synchronizing events can not be posted less than one time quantum in the future unless the simulation is already synchronized.

17.3. Multiprocessor Simulation

Simics MAI has limited support for multiprocessor simulation; processors are always scheduled in a round-robin fashion, one cycle at a time.

Let us have a look at a 2-machines setup containing two SPARC SunFire machine (with one processor each) to illustrate multiprocessor simulation. The processor in the first machine runs at 168MHz; the other runs at 56MHz (equal to 168/3). The time quantum (configured via the **cpu-switch-time** command) is 1000 cycles of the first processor, or 6 microseconds.

```
+-----+      Copyright 1998-2005 by Virtutech, All Rights Reserved
| Virtutech |      Version:
| Simics    |      Build:
+-----+
www.simics.com      "Virtutech" and "Simics" are trademarks of Virtutech AB
```



```
simics> @conf.d1_cpu0.freq_mhz
168
simics> @conf.d2_cpu0.freq_mhz
56
simics> @conf.sim.cpu_switch_time
1000
simics> c 10000
[d1_cpu0] v:0xffffffff0001364 p:0x1fff0001364 bne,pt %xcc, 0xffffffff0001360
simics> ptime -all
processor          steps          cycles      time [s]
d1_cpu0            10000          10000        0.000
d2_cpu0             3333           3333        0.000
```

While the first processor executed 10000 steps, the second processor completed 3333 steps, which corresponds to the ratio between the two frequencies (168MHz compared to 56MHz). Let us now examine the effects of the time quantum:

```
simics> c 30
[d1_cpu0] v:0xffffffff0001364 p:0x1fff0001364 bne,pt %xcc, 0xffffffff0001360
simics> ptime -all
processor          steps          cycles      time [s]
d1_cpu0            10030          10030        0.000
d2_cpu0             3333           3333        0.000
```

Although the first processor ran 30 steps further, the second processor has not run the 10 steps that we would expect, and the frequency ratio is not respected anymore. This is the effect of the 1000 cycles time quantum: the first processor is scheduled for the next 1000 cycles and no other processor will be run until the quantum is finished. If we switch to the second processor and try to make it run one step further, we will observe the following:

17.3. Multiprocessor Simulation

```
simics> pselect d2_cpu0
simics> c 1
[d2_cpu0] v:0xffffffff0001364 p:0x1fff0001364 bne,pt %xcc, 0xffffffff0001360
simics> ptime -all
processor          steps          cycles      time [s]
d1_cpu0            11000          11000        0.000
d2_cpu0            3334           3334         0.000
```

The second processor has run 1 step further as requested, but the first had to finish its time quantum before the second processor could be allowed to run, which explains its step count of 11000 compared to 10030 before. Let us now set the time quantum to 1:

```
simics> cpu-switch-time 1
The switch time will change to 1 cycles (for CPU-0) once all 2
processors have synchronized.
simics> c 1
[d2_cpu0] v:0xffffffff0001368 p:0x1fff0001368 nop
simics> ptime -all
processor          steps          cycles      time [s]
d1_cpu0            11000          11000        0.000
d2_cpu0            3335           3335         0.000
```

Note that the new time quantum length will only become valid once all processors have finished their current time quantum. This is why stepping one more step forward with the second processor hasn't affected the first yet. Now let us select the first processor again, and run three steps:

```
simics> pselect d1_cpu0
simics> c 3
[d1_cpu0] v:0xffffffff0001368 p:0x1fff0001368 nop
simics> ptime -all
processor          steps          cycles      time [s]
d1_cpu0            11003          11003        0.000
d2_cpu0            3668           3668         0.000
simics> c 3
[d1_cpu0] v:0xffffffff0001368 p:0x1fff0001368 nop
simics> ptime -all
processor          steps          cycles      time [s]
d1_cpu0            11006          11006        0.000
d2_cpu0            3669           3669         0.000
simics>
```

17.3. Multiprocessor Simulation

All processors finished their 1000 cycles time quantum and started to run with the new 1 cycle value, which means that they are now advancing in lockstep. For every 3 steps performed by the first processor, the second executes 1.

Chapter 18

Cache Simulation

18.1 Introduction to Cache Simulation with Simics

By default, Simics does not model any cache system. It uses its own memory system to obtain high speed simulation, and modeling a hardware cache model would only slow it down.

Note: Although it is often said that Simics models a processor, such as an UltraSPARC III or an x86 P4, remember that Simics is an *instruction-set* simulator, not a processor simulator. Transactions coming out of a real x86 P4 processor have already gone through the L1 and L2 caches, so they consist mainly of cache misses to be fetched from memory. In Simics, all transactions performed by the processor execution core are made visible.

For simplicity and performance, Simics does not model incoherence. In Simics, the memory is *always* up to date with the latest CPU and device transactions. Memory accesses take no time to execute and are always atomic.

The possibility to observe and alter memory transactions (both in timing and in execution) makes Simics very suitable for cache simulation:

Cache Profiling

The goal is to gather information about the cache behavior of a system or an application. Unless the application runs on multi-processors, takes a lot of interrupts or runs a lot of system-level code, the timing of the memory operations is often irrelevant, thus no stalling is necessary. The `timing-model` interface is a good place to be informed of all transactions sent by the processor.

Note that this type of simulation *does not modify* the execution of the target program. It could be done by using Simics as a simple memory transaction trace generator, and then computing the cache state evolution afterward. However, doing the cache simulation at the same time as the execution enables a number of optimizations that Simics models make good use of.

Cache Timing

The goal is to study the timing behavior of the transactions, in which case a transaction

to memory should take much more time than, for example, a transaction to an L1 cache. This is useful when studying interactions between several CPUs, or to grossly estimate the CPI of an application. To be able to stall the processor, a cache timing model should be connected to the `timing-model` interface. Simics models can be used for such a simulation.

This type of simulation *modifies* the execution, since interrupts and multi-processor interaction will be influenced by the timing provided by the cache model. However, unless the target program is not written properly, the execution will always be correct, although different from the execution obtained without any cache model.

Cache Content Simulation

It is possible to change Simics coherency model by allowing a cache model to contain data that is different from the contents of the memory. Such a model needs to use both the `timing-model` and the `snoop-memory` interfaces to properly handle the memory transactions (it must be able to change the values of loads and stores or to prevent their execution to main memory).

Note that this kind of simulation is difficult to do and requires a well-written, bug-free cache model, since it can prevent the target program from executing properly.

Simics comes with cache models that allow for the two first types of cache simulation:

- **g-cache** is the standard cache model. It handles one transaction at a time in a flat way: all needed operations (copy-back, fetch, etc.) are performed in order and at once. The cache returns the sum of the stall times reported for each operation.
- **g-cache-ooo** is a more complex model adapted to Simics MAI. It handles multiple outstanding transactions and keeps track of their current status (copy-back or fetch ongoing). **g-cache-ooo** is the standard cache model for Simics out-of-order; it is described in detail in *Simics Micro-Architectural Interface*.

The source code of these caches is available in `[simics]/src/extensions/`.

18.2 Simulating a Simple Cache

Adding a **g-cache** to a system is pretty straightforward. You can append the following code to the script creating your simulated machine:

```
@cache = pre_conf_object('cache', 'g-cache')
@cache.cpus = conf.cpu0
@cache.config_line_number = 256
@cache.config_line_size = 32
@cache.config_assoc = 1
@cache.config_virtual_index = 0
@cache.config_virtual_tag = 0
@cache.config_replacement_policy = 'random'
@cache.penalty_read = 0
```

18.3. Example Machines

```
@cache.penalty_write = 0
@cache.penalty_read_next = 0
@cache.penalty_write_next = 0

@SIM_add_configuration([cache], None)
```

This command will tell Simics to create a new object called **cache**, of type **g-cache**. It also sets some parameters to describe the cache (number of lines, size, associativity, type of index and tag, replacement policy, and stall penalties).

Now start Simics with the `-stall` flag to load the configuration with the cache in a mode where cache simulation is working properly. The only thing left is to connect the cache somewhere so that it will receive memory accesses. To connect it to the main memory, execute:

```
@conf.phys_mem0.timing_model = conf.cache
```

Note: The name of the main physical memory space can be different from one simulated machine to another, but the names `phys_mem` and `phys_mem0` are the most commonly used.

From now on, accesses to main memory will go through the cache and cache hits/misses will be simulated. If you run the simulation for a while, you will get information about the cache with the **cache.status** and **cache.statistics** commands.

Note: For performance reasons, instruction fetches are not sent to the caches unless you explicitly ask for it. Refer to section [16.3](#) for more information.

18.3 Example Machines

Some machines in the `[simics]/targets/` directories are pre-configured with a simple data/instruction cache per processor. You will recognize them by their name: `name-gcache-common.simics`.

18.4 A More Complex Cache System

Using **g-cache**, we can simulate more complex cache systems. Let us consider the structure described by figure [18.1](#).

What we want to simulate is a system with separate instruction and data caches at level 1 backed by a level 2 cache. We will connect this memory hierarchy to an x86 processor. The dotted components in the diagram represent elements that are introduced in Simics to complete the simulation. Let us have a look at these components:

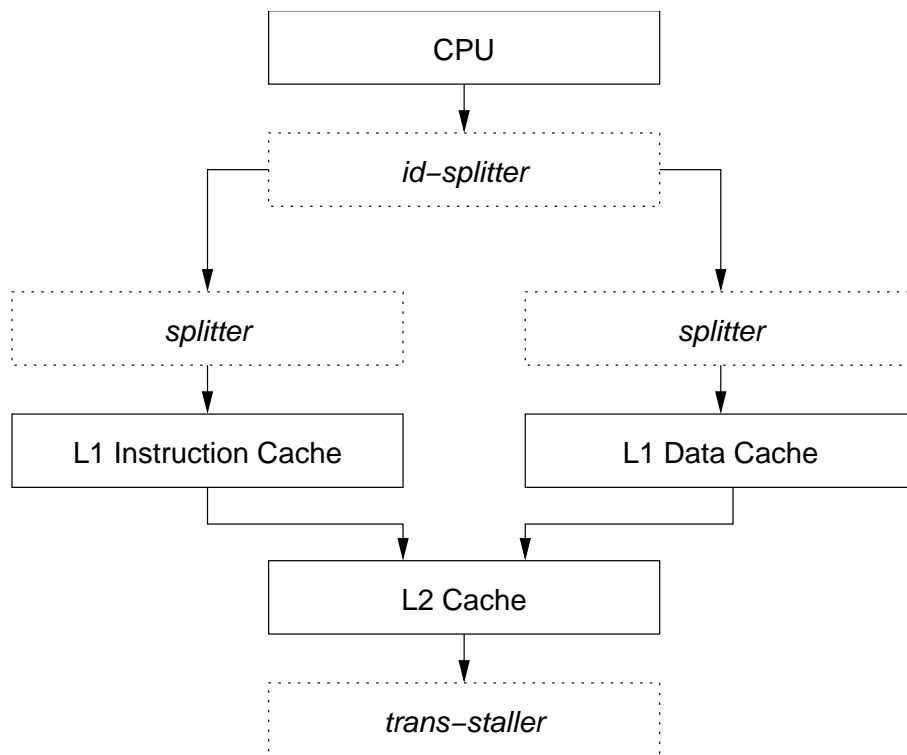


Figure 18.1: A More Complex Cache System

18.4. A More Complex Cache System

id-splitter

This module is used by Simics to separate instruction and data accesses and send them to separate L1 caches.

splitter

Since we are simulating an x86 machine, accesses can cross a cache-line boundary. To avoid that, we connect two splitters before the caches. The splitters will let un-cacheable and correctly aligned accesses go through untouched, whereas others will be split in two accesses.

trans-staller

The **trans-staller** is a very simple device that will simulate the memory latency. It will stall all accesses by a fixed amount of cycles.

In terms of script, this configuration gives us the following:

```
#
# Transaction staller for memory
#
@staller = pre_conf_object('staller', 'trans-staller')
@staller.stall_time = 200
#
# l2 cache: 512Kb Write-back
#
@l2c = pre_conf_object('l2c', 'g-cache')
@l2c.cpus = conf.cpu0
@l2c.config_line_number = 4096
@l2c.config_line_size = 128
@l2c.config_assoc = 8
@l2c.config_virtual_index = 0
@l2c.config_virtual_tag = 0
@l2c.config_write_back = 1
@l2c.config_write_allocate = 1
@l2c.config_replacement_policy = 'lru'
@l2c.penalty_read = 10
@l2c.penalty_write = 10
@l2c.penalty_read_next = 0
@l2c.penalty_write_next = 0
@l2c.timing_model = staller
#
# instruction cache: 16Kb
#
@ic = pre_conf_object('ic', 'g-cache')
@ic.cpus = conf.cpu0
@ic.config_line_number = 256
@ic.config_line_size = 64
@ic.config_assoc = 2
```

```

@ic.config_virtual_index = 0
@ic.config_virtual_tag = 0
@ic.config_replacement_policy = 'lru'
@ic.penalty_read = 3
@ic.penalty_write = 0
@ic.penalty_read_next = 0
@ic.penalty_write_next = 0
@ic.timing_model = l2c
#
# data cache: 16Kb Write-through
#
@dc = pre_conf_object('dc', 'g-cache')
@dc.cpus = conf.cpu0
@dc.config_line_number = 256
@dc.config_line_size = 64
@dc.config_assoc = 4
@dc.config_virtual_index = 0
@dc.config_virtual_tag = 0
@dc.config_replacement_policy = 'lru'
@dc.penalty_read = 3
@dc.penalty_write = 3
@dc.penalty_read_next = 0
@dc.penalty_write_next = 0
@dc.timing_model = l2c
#
# transaction splitter for instruction cache
#
@ts_i = pre_conf_object('ts_i', 'trans-splitter')
@ts_i.cache = ic
@ts_i.timing_model = ic
@ts_i.next_cache_line_size = 64
#
# transaction splitter for data cache
#
@ts_d = pre_conf_object('ts_d', 'trans-splitter')
@ts_d.cache = dc
@ts_d.timing_model = dc
@ts_d.next_cache_line_size = 64
#
# instruction-data splitter
#
@id = pre_conf_object('id', 'id-splitter')
@id.ibranch = ts_i
@id.dbranch = ts_d

```

18.5. Workload Positioning and Cache Models

```
@SIM_add_configuration([staller, l2c, ic, dc, ts_i, ts_d, id], None)
```

Once this is done, we can simply plug the **id-splitter** to the main memory:

```
@conf.phys_mem0.timing_model = conf.id
```

Note the way the penalties have been set: we don't use *_next* penalties but let the next level report penalties in case they are accessed. In this configuration, a read hit in L1 would take 3 cycles; a read miss that goes to memory would take $3 + 10 + 200 = 213$ cycles if no copy-back is performed in the L2 cache. There's no best way to set up the penalties so it's up to you to decide how your model should behave.

18.5 Workload Positioning and Cache Models

You may have noticed that connecting the caches to the memory-space can be done separately from defining the caches. It is possible to connect and disconnect the caches at any time during the simulation. You may, for example, boot an operating system and set up the workload with Simics in *-fast* mode, save a checkpoint and reload it in *-stall* mode when you want to start cache simulation.

To get decent cache statistics, you should run a few million instructions to warm up the caches before actually starting to do measurements on the cache behavior of your workload. Note that this is only a rough advice; the precise warm-up time needed will depend on the cache model and the workload.

18.6 Using g-cache

Let us have a more detailed look at **g-cache**. It has the following features:

- Configurable number of lines, line size and associativity. Note that the line size must be a power of 2, but the only restriction on the cache structure is that the number of lines divided by the associativity must be a power of two.
- Physical/virtual index and tag.
- Configurable write allocate/write back policy.
- Random, true LRU or cyclic replacement policies. It is easy to add new replacement policies.
- Sample MESI protocol.
- Support for several processors connected to one cache.
- Configurable penalties for read/write accesses to the cache, and read/write accesses initiated by the cache to the next level cache.
- Cache miss profiling.

Transactions are handled one at a time; all operations are done in order, at once, and a total stall time is returned. The transaction is not reissued afterward. Here's a short description of the way **g-cache** handles a transaction (we assume a write-back, write allocate cache):

- If the transaction is uncacheable, **g-cache** ignores it.
- If the transaction is a read hit, **g-cache** returns *penalty_read* cycles of penalty.
- If the transaction is a read miss, **g-cache** asks the replacement policy to provide a cache line to allocate.
- The new cache line is emptied. If necessary, a copy-back transaction is initiated to the next level cache. In this case, a penalty of *penalty_write_next* is counted, added to the penalty returned by the next level.
- The new data is fetched from the next level, incurring *penalty_read_next* cycles penalty added to the penalty returned by the next level.
- The total penalty returned is the sum of *penalty_read*, plus the penalties associated with the copy-back (if any), plus the penalties associated with the line fetch.

Note the usage of *penalty_read/write* and *penalty_read/write_next*: a write-through cache would always take a *penalty_write_next* penalty independently of the fact that a write is a hit or a miss, but **g-cache** always reports hits and misses according to the line lookup, not based on the fact that a transaction is propagated to the next level.

18.7 Understanding g-cache Statistics

The following statistics are available in a **g-cache**:

Total number of transactions

Count of all transactions received by the cache, including all transactions listed below.

Device data read/write

Number of transactions, e.g., DMA transfers, performed by devices against the memory-space to which the cache is connected. Device accesses are otherwise ignored by the cache and passed as-is to the next level cache.

Uncacheable data read/write, instruction fetch

Number of uncacheable transactions performed by the processor. Uncacheable transactions are otherwise ignored by the cache and passed as-is to the next level cache.

Data read transactions

Cacheable read transactions counted by the cache.

Data read misses

Cacheable read transactions that missed in the cache.

18.8. Speeding up g-cache simulation

Data read hit ratio

$1 - (\text{cacheable read misses} / \text{cacheable read transactions})$

Instruction fetch transactions

Cacheable instruction fetch transactions counted by the cache.

Instruction fetch misses

Cacheable instruction fetch transactions that missed in the cache.

Instruction fetch hit ratio

$1 - (\text{cacheable instruction fetch misses} / \text{cacheable instruction fetch transactions})$

Data write transactions

Cacheable write transactions counted by the cache.

Data write misses

Cacheable write transactions that missed in the cache. This is not directly related to the number of transactions sent to the next level cache, which also depends on the write allocation and write-back policies selected for the cache.

Data write hit ratio

$1 - (\text{cacheable write misses} / \text{cacheable write transactions})$

Copy-back transactions

Copy-back transactions performed by the cache to flush modified cache lines.

18.8 Speeding up g-cache simulation

By default, **g-cache** will try to use the STCs to minimize the number of transactions that it has to handle while still providing correct statistics and behavior.

Simulating a data cache

If you are only interested in data accesses, **g-cache** can use the Data STC and still provide correct statistics. To allow **g-cache** to use the Data STC, you should set the *penalty_read* and *penalty_write* to 0 (so that cache hits do not take any penalty). **g-cache** will then allow cache hit transactions to be saved in the Data STC and use its internal counters to report a correct number of total transactions, and thus correct ratios.

Note that when using the Data STC counters, **g-cache** can not determine to which memory space the accesses reported by the DSTC belong to, so you need to connect **g-cache** to all the memory spaces to which the processor is talking. In practice, the processor often talks to one main memory space and nothing special needs to be done. Sun's UltraSPARC machines, however, have separate physical memory and physical I/O spaces. The cache should then be connected to both of them. Another way of solving this problem is to connect a small module that will prevent accesses to other memory spaces from being cached in the Data STC.

Simulating an instruction cache

When simulating an instruction cache, **g-cache** is able to use the Instruction STC to speed up the simulation and report the number of instruction misses, but it won't report the correct number of total transactions. If you wish to have a correct total amount of instruction fetches, you need to disable ISTC usage at the command-line with the **istc-disable** command.

You can always prevent **g-cache** from using the STCs if you encounter one of the limitations mentioned later in this chapter, by setting the *config_block_STC* attribute of the cache to 1.

18.9 Cache Miss Profiling

g-cache can use Simics's data profiling support to profile cache misses. You can use the **add-profiler** and **remove-profiler** commands to add and remove profiler to the cache for a specific type of events.

g-cache can profile the following type of events:

- number of data read misses per virtual address, physical address, or instruction
- number of data write misses per virtual address, physical address, or instruction
- number of instruction fetch misses per virtual or physical address

For example, if you would like to see which parts of the code are responsible for read and write misses, you could create profilers that count read and write misses per instruction. (This example assumes that your cache is called **dc**.)

```
simics> dc.add-profiler type = data-read-miss-per-instruction
[dc] New profiler added for data-read-miss-per-instruction: 2
dc_prof_data-read-miss-per-instruction
simics> dc.add-profiler type = data-write-miss-per-instruction
[dc] New profiler added for data-write-miss-per-instruction: 2
dc_prof_data-write-miss-per-instruction
```

This creates two profiler objects and attaches them to the proper slots in the cache object. The profilers are initially empty, so we have to run for a while to give them time to collect some interesting data:

```
simics> c 10_000_000
[cpu0] v:0x00000000000002590 p:0x0000000003c7e590 sll %i0, 1, %o1
```

Now, we can ask either profiler what data it has gathered:

```
simics> dc_prof_data_read_miss_per_instruction.address-profile-data
View 0 of dc_prof_data_read_miss_per_instruction: dc prof: data-read-miss-per-instruction
```

18.9. Cache Miss Profiling

64-bit virtual addresses, profiler granularity 4 bytes

Each cell covers 2 address bits (4 bytes).

column offsets:

	0x1*	0x00	0x04	0x08	0x0c	0x10	0x14	0x18	0x1c
0x00000000000002480:	.	.	.	23331
0x000000000000024a0:	.	.	15492	545	.	.	90	.	.
0x000000000000024c0:	.	.	.	112
0x000000000000024e0:
0x00000000000002500:
0x00000000000002520:
0x00000000000002540:
0x00000000000002560:
0x00000000000002580:
0x000000000000025a0:
0x000000000000025c0:
0x000000000000025e0:	136	.	.

39706 counts shown. 0 not shown.

Since these two profilers are instruction indexed, it also makes sense to display their counts in a disassembly listing:

```
simics> cpu0.aprof-views add = dc_prof_data_read_miss_per_instruction
simics> cpu0.aprof-views add = dc_prof_data_write_miss_per_instruction
simics> cpu0.disassemble %pc 32
v:0x00000000000002590 p:0x0000000003c7e590    0 0  sll %i0, 1, %o1
v:0x00000000000002594 p:0x0000000003c7e594    0 0  ldub [%o1 + %o2], %o1
v:0x00000000000002598 p:0x0000000003c7e598    0 0  and %l0, %o1, %o1
v:0x0000000000000259c p:0x0000000003c7e59c    0 0  sll %o1, 16, %i0
v:0x000000000000025a0 p:0x0000000003c7e5a0    0 0  sra %i0, 16, %i0
v:0x000000000000025a4 p:0x0000000003c7e5a4    0 0  jmp1 [%i7 + 8], %g0
v:0x000000000000025a8 p:0x0000000003c7e5a8    0 0  restore %g0, %g0, %g0
v:0x000000000000025ac p:0x0000000003c7e5ac    0 0  jmp1 [%i7 + 8], %g0
v:0x000000000000025b0 p:0x0000000003c7e5b0    0 0  restore %g0, -1, %o0
v:0x000000000000025b4 p:0x0000000003c7e5b4    0 0  illtrap 0
v:0x000000000000025b8 p:0x0000000003c7e5b8    0 0  illtrap 0
v:0x000000000000025bc p:0x0000000003c7e5bc    0 0  illtrap 0
v:0x000000000000025c0 p:0x0000000003c7e5c0    0 0  illtrap 0
v:0x000000000000025c4 p:0x0000000003c7e5c4    0 0  illtrap 0
v:0x000000000000025c8 p:0x0000000003c7e5c8    0 0  save %sp, -96, %sp
v:0x000000000000025cc p:0x0000000003c7e5cc    0 0  sethi %hi(0x41c00), %o0
v:0x000000000000025d0 p:0x0000000003c7e5d0    0 0  add %o0, 840, %o1
v:0x000000000000025d4 p:0x0000000003c7e5d4    0 0  ldub [%o1 + 0], %o0
```

18.10. Using g-cache with Several Processors

```
v:0x000000000000025d8 p:0x0000000003c7e5d8    0 0  subcc %o0, 1, %o0
v:0x000000000000025dc p:0x0000000003c7e5dc    0 0  stw %o0, [%o1 + 0]
v:0x000000000000025e0 p:0x0000000003c7e5e0    0 0  bpos 0x25f0
v:0x000000000000025e4 p:0x0000000003c7e5e4    0 0  mov -1, %i0
v:0x000000000000025e8 p:0x0000000003c7e5e8    0 0  jmpl [%i7 + 8], %g0
v:0x000000000000025ec p:0x0000000003c7e5ec    0 0  restore %g0, %g0, %g0
v:0x000000000000025f0 p:0x0000000003c7e5f0    0 0  sethi %hi(0x4f000), %o0
v:0x000000000000025f4 p:0x0000000003c7e5f4    0 0  add %o0, 616, %o0
v:0x000000000000025f8 p:0x0000000003c7e5f8   136 0  ldub [%o0 + 0], %o2
v:0x000000000000025fc p:0x0000000003c7e5fc    0 0  add %o2, 1, %o1
v:0x00000000000002600 p:0x0000000003c7e600    0 0  stw %o1, [%o0 + 0]
v:0x00000000000002604 p:0x0000000003c7e604    0 0  ldub [%o2 + 0], %o2
v:0x00000000000002608 p:0x0000000003c7e608    0 0  and %o2, 255, %i0
v:0x0000000000000260c p:0x0000000003c7e60c    0 0  jmpl [%i7 + 8], %g0
```

In the listing above, we see that in the 32 instructions following the current program counter, one load instruction is responsible for 136 read misses, and no instructions have caused any write misses.

For more on getting information out of profilers, see section [13.3](#).

18.10 Using g-cache with Several Processors

Support for several processors talking to one cache is integrated in **g-cache**. You just need to specify the list of CPUs connected to the cache in the *cpus* attribute. Note that it is possible for the cache to use the STCs as described above even with several processors.

You can use the sample MESI protocol to connect several caches in a multiprocessor system. You need to provide the cache with a list of the other caches snooping on the bus using the *snoopers* attribute. If you have higher-level caches that are not snooping on the bus, you need to set the *higher_level_caches* attribute so that invalidation is done properly. Note that the sample MESI protocol was written to handle simple cases such as several L1 write-through caches with L2 caches synchronizing via MESI. To model more complex protocol, you will need to modify **g-cache**.

If you use LRU replacement with several processors, you may have problems with the way Simics schedules processor execution (read the chapter [17](#) for more information). You may want to lower the CPU switch time from its default value to prevent accesses “in the future” from changing the way LRU replacement is behaving.

18.11 g-cache Limitations

Virtually tagged caches are usually informed of the state of the MMU in order to immediately invalidate lines whose virtual-to-physical mapping changes. **g-cache** is not MMU-aware, so when looking for a hit with virtual tags, it tries to match both the virtual tag and the physical tag. This approach prevents cache accesses from triggering a hit on lines with invalid translations, but it makes it possible for some invalid mappings to be present in the cache (and shown by the **status** command for example).

18.11. *g-cache Limitations*

Some special instructions (atomic instructions in particular) can cause the STC counters to be off by about one memory access per million, which influences the total number of transactions reported by the cache. The UltraSPARC architecture should be free of this bug, while others may still trigger it in some circumstances. The workaround is to avoid using the STC if a very precise total transaction count is needed.

g-cache doesn't understand control transactions like cache flushes and cache line locking.

Chapter 19

Memory Spaces

19.1 Memory Space Basics

Simics memory-spaces are handled by the generic **memory-space** class. A **memory-space** object implements interface functions for memory accesses, and it has attributes specifying how the mappings are setup.

The most important attribute in a memory-space is the *map* attribute. This is a list of mapped objects that may contain devices, RAM and ROM objects, and other memory-spaces. In addition to the *map* attribute, there is also a *default_target* attribute that is used for accesses that don't match any of the targets in the map list.

Python example of a **memory-space** object where already created objects are mapped into a new memory space:

```
@mem = pre_conf_object('phys_mem', 'memory-space')
@mem.map = [[0x00000000, conf.ram0, 0, 0, 0xa0000],
            [0x000a0000, conf.vga0, 1, 0, 0x20000],
            [0x000c0000, conf.rom0, 0, 0, 0x10000],
            [0x000f0000, conf.rom0, 0, 0, 0x10000],
            [0x00100000, conf.ram0, 0, 0x100000, 0xff00000],
            [0xfe00000, conf.apic0, 0, 0, 0x4000],
            [0xffe81000, conf.hfs0, 0, 0, 16],
            [0xffff0000, conf.rom0, 0, 0, 0x10000]]
@mem.default_target = [conf.pci_mem0, 0, 0, conf.pci_mem0]
@SIM_add_configuration([mem], None)
```

The fields for an entry in the map list are as follows:

base

The start address of the mapping in the memory space.

object

Reference to the mapped object.

function

An object specific identification number for this mapping. The function number is

typically used by objects that have several mappings. When an object is accessed, it can use the function number to figure out what mapping it was accessed through.

offset

The start offset in the target object. This is often used when a memory-space is mapped in another memory-space. Example: memory-space B is mapped in space A at base 0x4000 with length 0x1000, and with offset 0x2000. If an access is made in space A at address 0x4010, it will be forwarded to space B at address 0x2010. Without any offset in the mapping, the resulting address would have been 0x10.

length

The size of the mapping in bytes.

target

(optional) If *object* isn't the final destination for the access, then *target* is a reference to the actual target object. This is described in more details in the section about different mapping types.

priority

(optional) The priority of the mapping. Default is 0 (highest priority), and the lowest is 256. If mappings overlap, then the priority field specified what mapping that has precedence. It is an error if overlapping mappings have the same priority. Usually overlapping mappings should be avoided, but for bridges that catch unclaimed accesses in specific address ranges the priority field is useful. There are also devices that have overlapping mappings that have different priorities, and in that case the priority field in the map list can be used.

align-size

(optional) The *align-size* can be used if a target does not support large accesses. Accesses that crosses an alignment boundary will be split into several transactions by the Simics memory system. By default this will be set to 4 bytes for port space devices, 8 bytes for other devices and 4096 or 8192 for memory.

reverse-endian

(optional) Some device mappings reverse the byte order of data, to support mixed-endian environments. The *reverse-endian* field can be used to model this behavior. It will reverse the byte-order of data for mappings that have an *align-size* of 2, 4 or 8 bytes.

Note: In Simics versions prior to 2.0, the priority field was a “reverse-endian” flag. Old checkpoints will automatically be updated with this field cleared, since the handling of endian swapping has changed. This old “reverse-endian” flag was obsoleted since devices in 2.0 read and write data with explicit endianness, and reversing the byte-order of all accesses in a memory-space often was incorrect.

19.2 Memory Space Commands

All mappings in a memory-space can be viewed with the `<memory-space>.map` command. Example:

```
simics> phys_io0.map
base          object          fn offs          length
0x0000040000400000 chmmu0          0 0x0          0x48
0x000004000c000000 schizo24          0 0x0          0x800000
0x000004000c800000 schizo25          0 0x0          0x800000
0x000007fff0000000 serengeti_cpuprom 0 0x0          0xbd3b0
0x000007fff0102000 serengeti_fpost_code 0 0x0          0x2000
0x000007fff0104000 fpost_data0          0 0x0          0x2000
0x000007fff0800060 empty0_0          0 0x60          0x10
0x000007fff091e000 empty1_0          0 0x0          0x120
```

Another useful command is `devs`, that lists all mapped devices in the system.

```
simics> devs
Count    Device          Space          Range          Func
      0  chmmu0          phys_io0          0x0000040000400000 - 0x0000040000400047    0
      0  e2bus24B_1      bus_pcicfg24B      0x0000000000000800 - 0x00000000000008ff  255
      0  empty0_0          phys_io0          0x000007fff0800060 - 0x000007fff080006f    0
      0  empty1_0          phys_io0          0x000007fff091e000 - 0x000007fff091e11f    0
      0  fpost_data0        phys_io0          0x000007fff0104000 - 0x000007fff0105fff    0
      0  glm0              bus_pcicfg25B      0x0000000000000100 - 0x000000000000010fff  255
...

```

19.3 Memory Mapping Types

There are a few different kinds of mappings that can be specified in the map attribute. All use the format described in the previous section.

Device Mapping

The most common kind of mapping. It is used for devices, RAM and ROM objects. The *target* field is not set.

Translator Mapping

Sometimes the address has to be modified between memory-spaces, or the destination memory-space depends on the address or some other aspect of the access such as the initiating processor. In these cases a *translator* can be used. A translator mapping is specified with the translator in the *object* field, and the default target as *target*. The translator has to implement the `TRANSLATE` interface. When an access reaches a translator mapping, the *translate* function in the `TRANSLATE` interface is called. The translator can then modify the address if necessary, and specify what destination

memory-space to use. If it doesn't specify any new memory-space, the default one from the configuration is used. The following fields can be changed by the translator: `physical_address`, `ignore`, `block_STC`, `inverse_endian` and `user_ptr`.

Translate to RAM/ROM Mapping

Used to map RAM and ROM objects with a translator first. The *object* field is set to the translator, and *target* is set to the RAM/ROM object.

Space-to-space Mapping

Map one memory-space in another. Both *object* and *target* should be set to the destination memory-space object.

Bridge Mapping

A bridge mapping is typically used for mappings that are setup by some kind of bridge device. The purpose of a bridge mapping is to handle accesses where nothing is mapped, in a way that corresponds to the bus architecture. For a bridge mapping, the *object* field is set to the bridge device, implementing the `BRIDGE` interface. The *target* field is set to the destination memory-space. If both a translator and bridge is needed, they must be implemented by the same object. If an access is made where nothing is mapped, the memory-space by default returns the `Sim_PE_IO_Not_Taken` pseudo exception. But if the access was made through a bridge mapping, the bridge device will be called to notify it about the unmapped access. It can then update any internal status registers, specify a new return exception, and set the data that should be returned in the case of a read access. Since the bridge is associated with the mapping and not the memory-space itself, several bridges can exist for one space, and devices doing accesses directly to the memory-space in question will not affect the bridge for non-mapped addresses. In the latter case, the device itself has to interpret the `Sim_PE_IO_Not_Taken` exception. The `Sim_PE_IO_Error` exception, indicating that a device returned an error is also sent to the bridge. Finally, bridges are called for accesses that generate `Sim_PE_Inquiry_Outside_Memory`, i.e. an inquiry access where nothing is mapped. In this case the bridge may have to set a default return value, such as `-1`.

19.4 Avoiding Circular Mappings

Since memory-spaces can be mapped in other memory-spaces, it is possible to create loops where accesses never reach any target device. One typical case is a PCI memory-space with bridges that has both downstream and upstream mappings. In a real system, the bridge typically doesn't snoop its own transactions and there will be no loop. But in Simics there are usually *n* bridge devices mapped between different memory-spaces. Instead the bridge will create space-to-space mappings. To avoid loops in this case, bridges should be careful when setting up memory-space mappings. Simics will make sure that an access does not loop by terminating it after 512 memory-space transitions. If this limit is reached, it is considered a configuration error and Simics will stop.

Chapter 20

PCI Support in Simics

20.1 Introduction

Simics models a PCI bus using several objects of various kinds:

- A **pci-bus** object that models the bus to which PCI devices are connected, keeping track of which device is in which slot.
- Three **memory-space** objects that model the three different access types to the PCI bus: one for configuration accesses, one for I/O accesses and one for memory accesses.
- An object acting as a bridge, linking the PCI bus to the rest of the system. This can be a host-to-PCI bridge (like a Northbridge chipset) or a PCI-to-PCI bridge (linking the PCI bus to a parent PCI bus).
- A number of PCI devices connected to the PCI bus object. Note that in Simics, each function of a multi-function PCI device must be represented by a separate object.

Let us look at a more concrete example from the simulated *Ebony* board:

```
OBJECT pcibus0 TYPE pci-bus {
    conf_space: pciconf0
    memory_space: pcimem0
    io_space: pciio0
    bridge: pci0
    interrupt: (pci0)
    pci_devices: ((0, 0, pci0, 1), (3, 0, dec0, 1))
}
OBJECT pciconf0 TYPE memory-space {
}
OBJECT pciio0 TYPE memory-space {
}
OBJECT pcimem0 TYPE memory-space {
}
OBJECT pci0 TYPE ppc440gp-pci {
```

```

        pci_bus: pcibus0
        ...
    }
OBJECT dec0 TYPE DEC21140A-dml {
    pci_bus: pcibus0
    ...
}

```

- The **pcibus0** object represents the PCI bus. The *pci-devices* attribute shows that two PCI devices are connected: **pci0** is the bridge (presented below), inserted as function 0 in slot 0; **dec0** is a DEC21140A network card inserted as function 0 in slot 3.
- The three memory-spaces are **pciconf0** for configuration, **pciio0** for I/O and **pcimem0** for memory. Note that these memory-spaces are *empty*. The configuration memory-space will be filled in automatically as devices are connected to the PCI slots. The other two spaces will be filled in when software enables different base address registers in the PCI devices. In most cases, these spaces need never be manipulated manually.
- The **pci0** object is a host-to-PCI bridge present in the 440GP chipset. It creates a bridge between the PPC 440GP processor and the PCI devices.
- The **dec0** PCI device is connected to the PCI bus in slot 3, function 0.

Once the configuration presented above is loaded into Simics, a number of parameters will be configured automatically. Let us look at the result of this automatic configuration. The PCI bus object now reports the devices that are connected:

```

simics> pcibus0.info
Information about pcibus0 [class pci-bus]
=====

        Bridge device : pci0
    Interrupt devices : pci0
        PCI Bus Number : 0x0

        Config space : pciconf0
            IO space : pciio0
        Memory space : pcimem0

Connected devices:
    Slot 0 function 0 : pci0
    Slot 3 function 0 : dec0

```

The configuration memory-space has been automatically configured to include the configuration registers defined by the devices connected on the bus. The exact addresses where

20.2. Configuration Space

the configuration registers are mapped are determined using the PCI configuration address described in the next section.

```
simics> pciconf0.map
base          object          fn offs          length
0x0000000000000000 pci0      255 0x0      0x100
0x0000000000000180 dec0      255 0x0      0x100
```

The PCI I/O and memory spaces are empty since no base address registers have been configured yet. The bridge, however, catches all unmapped accesses in the PCI memory space to redirect them towards the main memory. This enables PCI devices to perform DMA accesses to and from other PCI devices or main memory.

```
simics> pciio0.map
base          object          fn offs          length
simics> pcimem0.map
base          object          fn offs          length
- default -   pci0            5 0x0      -
               target -> plb
```

20.2 Configuration Space

The PCI configuration accesses are implemented with a generic Simics memory-space object. The space is populated by the **pci-bus** object based on the list provided by the *pci_devices* attribute. It does not need to be configured manually.




Figure 20.1: PCI Type 1 Configuration Address in Simics

The configuration space in Simics uses the Type 1 addressing layout (see figure 20.1), but with the lower bits used for byte addressing. No **IDSEL** signals are used. This is done to simplify the implementation, and does not impose any restriction of the PCI model.

In the previous example, the configuration space was configured with two devices as follow:

```
base          object          fn offs          length
0x0000000000000000 pci0      255 0x0      0x100
0x0000000000000180 dec0      255 0x0      0x100
```

The addresses are computed as explained by figure 20.1, which gives an address of 0x0 for **pci0** (bus 0, device or slot 0, function 0) and an address of 0x1800 for **dec0** (bus 0, device or slot 3, function 0).

In many systems, typically 32-bit ones, the configuration space is not visible in the global memory-space of the processor. Instead the bridge has a register pair that is used to access the configuration space: one register for the address and one for data.

PCI-to-PCI bridges will automatically add mappings to the configuration space of the primary bus for all subordinate buses. This is done at run-time, since both the secondary and subordinate bus is configured by the simulated software.

20.3 Memory and I/O Spaces

The PCI memory and I/O spaces are implemented as generic Simics memory-spaces where the spaces controlled by the base address registers of the PCI devices will be mapped.

Let us imitate the transactions that the simulated software would issue to map the space controlled by the base address register 1 of the **dec0** object above:

```
# Issue a write to BAR1 register (offset 0x14)
simics> pciconf0.set address = 0x1814 value = 0xFF000000 size = 4 -1

# Enable memory mappings in Command register (offset 0x4)
simics> pciconf0.set address = 0x1804 value = 0x0002 size = 2 -1

# Print the PCI memory mappings
simics> pcimem0.map
```

base	object	fn	offs	length
0x00000000ff000000	dec0	0	0x0	0x80
- default -	pci0	5	0x0	-

```
target -> plb
```

The address space controlled by the BAR1 register is now mapped at address in 0xFF000000 in the PCI memory.

PCI bridges in Simics typically create space-to-space mappings when adding a PCI memory space to the global memory-space in a system. This way memory transactions go directly between memory-spaces without having to pass the bridge object. There are, however, cases when the bridge is needed, for example for accesses where nothing is mapped. In these cases, the bridge is associated with the mapping (see section on memory mapping types), and the bridge will be called for non-mapped accesses.

The bridge mapping is also used to prevent memory mappings loops. A bridge typically doesn't forward upstream a transaction that it sent downstream previously. But in Simics the memory-spaces are connected directly to each other, and there is no bridge doing the forwarding. Since the bridge creates the mappings as "bridge mappings", Simics will make sure that the transaction isn't forwarded using the same bridge twice in a row. Circular mappings are not a problem in a correctly configured system, but can occur for memory accesses to a non-existing device.

20.4. PCI Interrupts

In the same way as with the configuration space, a PCI-to-PCI bridge will create mappings in the memory and I/O spaces for upstream and downstream accesses. These mappings are created at run-time, since they are dynamic and set up by the simulated software.

20.4 PCI Interrupts

There are four interrupt pins defined for PCI (A, B C and D). A PCI device uses the **pci-bus** object to raise and lower the interrupt signal for a specific interrupt pin. The **pci-bus** object then passes the interrupt to the host-to-PCI bridge that translates the interrupt to a system architecture specific interrupt.

20.5 Expansion ROM

All PCI devices that are based on the generic PCI device code in Simics can have support for expansion ROMs, depending on the device implementation. To add an expansion ROM image, a **rom** object should be created that contains the ROM memory image. This **rom** object is then referenced from the `expansion_rom` attribute in the PCI device. The Expansion ROM base address registers are handled automatically.

20.6 PCI Express

At the level of simulation provided by Simics, there is very little difference between a PCI bus and a PCI Express switch. The major change is the extension of the configuration registers to 4096 bytes per device instead of 256. This is taken into account if you define a PCI Express system using the **pcie-switch** class instead of the **pci-bus** class. The PCIe switch will check if the device claims to be PCIe compatible (whether it implements the `pci_express` interface or not) before allowing it to be connected. Other PCIe features depend on the device used as root complex for the system.

20.7 Other PCI Features

20.7.1 Master Abort

Accesses that pass a host-to-PCI bridge, or a PCI-to-PCI bridge, but have a non-mapped address as target will call the master abort handling in the bridge. This is implemented in Simics using the bridge mapping, where the most recent bridge will be called signaling that the access wasn't claimed by any device. The bridge can then perform the appropriate master abort semantics, and return a `Sim_PE_No_Exception` pseudo exception. If a PCI device issues a transaction on the local bus that is not claimed by any device, the memory-space access will return the pseudo exception `Sim_PE_IO_Not_Taken` directly to the device. This is the same exception that the bridge will get in case of a bridge mapping. It is important to use bridge mappings for PCI memory-spaces, since the processor should not get the `Sim_PE_IO_Not_Taken` exception on accesses, but instead get informed by the bridge about the error.

20.7.2 Target Abort

A PCI device can signal a target abort by returning the `Sim_PE_IO_Error` pseudo exception on accesses. It should also be prepared to receive this error when doing memory accesses itself on the PCI bus. PCI bridges implemented in Simics may receive the `Sim_PE_IO_Error` exception from the `PCI_BRIDGE` interface if the access was done through a bridge mapping, for example a CPU-initiated access to a PCI device below the bridge.

20.7.3 Message Signaling Interrupt

There is no generic support in Simics for Message Signaling Interrupts (MSI). This is something that can be implemented on a per device (or bridge) basis.

20.7.4 Special Cycles

The PCI system in Simics does support PCI Special Cycles. However, most PCI bridges and devices modeled do not have Special Cycles support implemented.

20.7.5 System Error (SERR#)

System Error is used in PCI to signal unrecoverable errors. A PCI device can assert the SERR# line on the PCI bus upon errors, and the host-to-PCI bridge informs the Operating System of the error. System Error is supported by the PCI system in Simics, but not all modeled bridges and devices implement it.

20.7.6 Parity Error (PERR#)

Simics currently does not support PCI Parity Error signaling. Please contact Virtutech if you need this modeled.

20.7.7 Interrupt Acknowledge

Interrupt Acknowledge is a rarely used feature of PCI. It is implemented in Simics's `pci-bus` object, but most bridges and devices do not support it.

20.7.8 VGA Palette Snooping

Simics currently does not support VGA Palette Snooping. Please contact Virtutech if you need this modeled.

Chapter 21

Driving Context Changes

As explained in section [12.3.1](#), each processor has a *current context*, which represents the virtual address space currently visible to code running on the processor. Things such as virtual-address breakpoints and symbol information are properties of contexts, so when the current context of a processor is changed, those things change with it.

Chapter [12.3](#) explains how this is used in practice. This chapter focuses on another issue: how to actually change the current context.

As a general guideline, you probably want a separate context for each userspace process and kernel that you are interested in (and a default context for everything you are not interested in). After all, one of the abstractions maintained by most operating systems is that each process has its own virtual address space, so if you, for example, set a virtual breakpoint at an address in one process, you probably do not want it to trigger in other processes.

The bad news is that since Simics does not need to know about the abstractions—such as processes—implemented by the operating system, it (quite correctly) does not try to. This is actually very good news most of the time, since Simics can then run any and all code without having to be modified, but in this case it is bad news: Simics knows nothing about processes on the simulated machine, and so cannot change contexts automatically. You will have to do that, and section [21.1](#) explains how.

The good news is that Simics comes with add-on modules, *process trackers*, that each understand a particular architecture/OS combination, and can help you tell which process is active when. They even come with source code, which is helpful if you want to write your own process tracker for some architecture/OS combination that does not yet have one. Process trackers are covered in section [21.2](#).

Simics also comes with a **context-switcher** module. It automates the task of listening to what the process tracker says and switching contexts accordingly. Section [21.3](#) describes how to use it.

21.1 Switching Contexts Manually

Changing the current context of a processor is very simple. Every processor starts out with **primary-context** as current context, so for this to be interesting we will need to create another context first:

```
simics> new-context my-little-context
```

Then, simply set the *current-context* attribute of the processor, either directly, or via the **set-context** command:

```
simics> cpu0.set-context my-little-context
```

And we are all done.

The tricky thing is to know when to change the context. This generally involves monitoring the state of the simulated machine, and looking for more or less subtle signs of relevant change. For example, assume that we are interested in the kernel. What we would like to do then is to change to the kernel context when the processor enters supervisor mode, and change back to **primary-context** when it enters user mode.

Conveniently, Simics triggers a `Core_Mode_Change` hap whenever the privilege level changes. The following Python script listens for that hap, and changes context accordingly.

```
def set_context(user_arg, cpu, old_mode, new_mode):
    if new_mode == Sim_CPU_Mode_Supervisor:
        cpu.current_context = conf.my_little_context
    else:
        cpu.current_context = conf.primary_context
SIM_hap_add_callback_obj("Core_Mode_Change", conf.cpu0,
                        0, set_context, None)
```

Note that, in this example, all of userspace is covered by a single context (**primary-context**). Distinguishing between different userspace processes is, as mentioned earlier, not something that Simics can do by itself. It needs a process tracker for that; this is the topic of the next section.

21.2 Process Trackers

As mentioned in the previous section, Simics triggers haps for hardware events such as processor privilege level changes, making it easy to change contexts in response to those. It does not trigger haps in response to software events, such as when the currently executing process changes, because it knows nothing about things such as processes. But if you know how the operating system works, you can inspect the information that Simics does provide, and figure out all kinds of stuff.

A *process tracker* does just that; inspects the simulated operating system, and triggers haps when interesting things occur.

Typically, a new tracker needs to be written for each architecture/OS combination. Simics comes with two trackers:

- **cpu-mode-tracker**, which only distinguishes between user and supervisor mode (but works on all targets and operating systems).

21.2. Process Trackers

- **linux-process-tracker**, which tracks processes on Linux on PowerPC, UltraSPARC and x86 targets.

Their source code can be found in `src/extensions`. The intention is that this will be helpful if you would like to write a tracker for some other architecture/OS combination.

Note that a tracker is watching over a specific set of processors. This set should typically contain all processors that run an operating system together, and no other processors. That way, processes that migrate between processors will not get lost, and processes running in different operating system instances will not be mixed up. If there is more than one operating system running in the simulation, they will need separate trackers.

A new **cpu-mode-tracker** can be created like this:

```
simics> new-cpu-mode-tracker name = tracker0
New cpu mode tracker tracker0 created.
simics> tracker0.add-processor cpu0
```

Creating a **linux-process-tracker** requires a little more work:

```
simics> new-linux-process-tracker name = tracker0
New process tracker tracker0 created.
simics> tracker0.add-processor cpu0
simics> tracker0.autodetect-parameters
```

Notice the added command **tracker0.autodetect-parameters**. This makes the process tracker examine memory to figure out what operating system the target machine is actually running. This means that the target system must be booted before this command is issued. If that is not an option, the OS version can be specified explicitly as an argument to **new-linux-process-tracker**. The process tracker just created should be suitable to try out the examples below.

If the trackers distributed with Simics do not fit your needs, there are some things to think about when creating a new tracker. (This part is beneficial to read even if you use the trackers shipped with Simics, as it explains the interface to the tracker, with some usage examples.) A tracker must do two things:

- Monitor a specific set of processors, and trigger the `Core_Trackee_Active` hap when a trackee becomes active or inactive on one of them.
- Answer questions about trackees by implementing the `tracker` interface.

(It says “tracker” instead of “process tracker”, and “trackee” instead of “process” in the text above because the things being tracked are not necessarily processes. **cpu-mode-tracker**, for example, tracks processor privilege levels.)

If the tracker tracks processes on a Unix-like operating system, it may additionally do the following:

- Monitor the processors, and trigger the `Core_Trackee_Exec` hap when a process calls the `exec` system call on one of them.

- Answer questions about processes by implementing the `tracker_unix` interface.

And if it wants to play nice with automatic configuring of process trackers, there is one more thing to do:

- Expose any parameter settings it requires through the `tracker_settings` interface.

These haps and interfaces are documented in the Reference Manual. The tracker typically implements this functionality by listening to haps such as `Core_Mode_Change` and `Core_Exception`, and knowing in which registers and memory locations the operating system stores interesting information.

Clearly, using just the basic interface (the `Core_Trackee_Active` hap and the `tracker` interface), it is easy to make a context follow the currently active trackee (this example assumes that there is a tracker called **tracker0**):

```
def set_context(user_arg, tracker, tid, cpu, active):
    if active:
        cpu.current_context = conf.my_little_context
    else:
        cpu.current_context = conf.primary_context
    current_tid = conf.tracker0.iface.tracker.active_trackee(
        conf.tracker0, conf.cpu0)
    SIM_hap_add_callback_obj_index("Core_Trackee_Active", conf.tracker0,
                                   0, set_context, None, current_tid)
```

Here, `current_tid` is set to the ID of the trackee that is active *when we run this code*. `SIM_hap_add_callback_obj_index` is then used to cause our callback function `set_context` to be called every time the trackee with this ID becomes active or inactive. For example, if **tracker0** is a process tracker, `current_tid` will be an ID representing the currently executing process (or the operating system—whichever was executing at the time). The callback function then ensures that **my_little_context** is active whenever that process is active, and that **primary_context** is active at all other times.

Using the Unix process tracker interface as well (the `Core_Trackee_Exec` hap and the `tracker_unix` interface), we can do more complicated things, such as waiting for a specific binary (`ifconfig` in this example) to be executed, then follow any process that executes it:

```
def exec_hap(user_arg, tracker, tid, cpu, binary):
    if binary.endswith("ifconfig"):
        def active_hap(user_arg, tracker, tid, cpu, active):
            if active:
                cpu.current_context = conf.my_little_context
            else:
                cpu.current_context = conf.primary_context
        SIM_hap_add_callback_obj_index("Core_Trackee_Active", tracker,
                                       0, active_hap, None, tid)
```


21.3. Switching Contexts Automatically

```
SIM_hap_add_callback_obj("Core_Trackee_Exec", conf.tracker0,  
                        0, exec_hap, None)
```

Section 21.3 reveals how to do this without having to write scripts.

Switching contexts is not all that can be done with process trackers. Here is another example that prints the number of steps a user process is running continuously in user mode. In this case we are looking at the program “ls”.

```
start_cycle = 0  
  
def exec_hap(user_arg, tracker, tid, cpu, binary):  
    global total_cycles, start_cycle  
    if binary.endswith("ls"):  
        def active_hap(user_arg, tracker, tid, cpu, active):  
            global total_cycles, start_cycle  
            if active:  
                start_cycle = SIM_cycle_count(cpu)  
            else:  
                print "ls ran for", (SIM_cycle_count(cpu) - start_cycle), "cycles"  
                start_cycle = SIM_cycle_count(cpu)  
        SIM_hap_add_callback_obj_index("Core_Trackee_Active", tracker,  
                                      0, active_hap, None, tid)  
    SIM_hap_add_callback_obj("Core_Trackee_Exec", conf.tracker0,  
                            0, exec_hap, None)
```

Section 12.3 has an example of how to use process trackers in symbolic debugging.

Note: Remember to use the CLI command `<tracker>.activate` (or, equivalently, call the *activate* function of the `tracker` interface) before trying to use a tracker.

21.3 Switching Contexts Automatically

If all you want to do is let a context follow a process with a given PID, or follow processes that execute a given binary, you do not have to program elaborate scripts like the ones at the end of section 21.2. There is a module called **context-switcher** that will do these common tasks for you so that you do not have to talk to the tracker directly:

```
simics> switcher0.track-pid pid = 4711 context = my-little-context  
Context 'my_little_context' is now tracking process 4711.  
simics> switcher0.track-bin binary = "emacs" context = emacs-context  
Context 'emacs_context' will be tracking the first process  
that executes the binary 'emacs'.
```

(These commands assume the existence of a **context-switcher** object called **switcher0**.)

21.3. *Switching Contexts Automatically*

Of course, you might want to do more complicated things; in that case, feel free to write all the scripts you want, or modify **context-switcher** to fit your needs; it, too, comes with source code.

Chapter 22

Understanding Hindsight

22.1 Command Overview

Hindsight™ is the technology utilized by Simics to achieve reverse execution.

Before any reverse operations can be performed, at least one time bookmark must be added. Depending upon how Simics is invoked (and various user preferences), a time bookmark denoted “start” is sometimes added automatically at the beginning of the simulation. Reverse operations are possible in the region following the first (i.e. oldest) bookmark. Bookmarks can be managed through the commands

```
simics> set-bookmark label
simics> list-bookmarks
simics> delete-bookmark [label | -all].
```

The main reverse execution commands are

```
simics> reverse
simics> reverse-to position
simics> skip-to position.
```

The **reverse** and **reverse-to** commands run the simulation backwards until a breakpoint occurs or till the oldest time bookmark is reached. Skipping is somewhat similar to reversing; the main difference is that intermediate breakpoints are ignored. Skipping can also be much faster than reversing.

The **skip-to** and **reverse-to** commands take either an absolute step count or a time bookmark as argument.

Many forward executing commands used for debugging has a corresponding reverse variant obtained by adding a reverse prefix. The reverse variant of **step-instruction** is for instance **reverse-step-instruction**.

Any external input (like a human typing on a virtual serial console) is replayed when the simulation is being run forward after a reversal; this guarantees that the simulation will follow the same path as originally. For the same reason, all external input is ignored until

the point is reached where the first reverse operation was initiated. It is possible to override this behavior with the

```
simics> clear-recorder
```

command. This command discards all recorded input and allows an alternate future to take place.

Note that external changes to the simulation not under the control of the recorder can make Hindsight operate somewhat unpredictably. It is recommended that any time bookmarks before a non-replayable change of the simulation state is deleted explicitly. Examples of such external changes are the modification of a CPU register by hand or loading a boot image from the command line.

22.2 Performance

Hindsight optimizes for certain reverse operations that are expected to be common. One example of this is that skipping to bookmarks can be faster than skipping to some other location.

The usage of time bookmarks has a certain impact on overall performance since it implicitly enables Hindsight support. Normal performance is always obtained if all time bookmarks are removed.

It is possible to tune certain reverse execution parameters in order to optimize operations for a particular usage pattern (although the default settings should work well in most situations). Tradeoffs exists between:

- reverse performance
- forward performance
- memory utilization
- scope of reversibility

The **rexec-limit** command is the primary tool for adjusting the balance, e.g.

```
simics> rexec-limit steps = 20000000
simics> rexec-limit size_mb = 200
```

The steps limit indicate that the scope of interest is at most the specified number of steps. By imposing a steps limit, resources can be spent more effectively with the drawback that reversal past the limit may not possible. By default, no steps limit is imposed.

The size limit imposes a limit on the amount of memory Hindsight may use. If the limit is exceeded, reverse performance will be traded for less memory consumption.

Index

Symbols

!, 181
-ma, 170
-stall, 170
@, 89
%, 49
[simics], 13
[workspace], 13

A

address space, 18
api-apropos, 92
api-help, 92
apropos, 179

B

BOOTP, 101
Bootstrap Protocol, 101
branch recorder, 161
brctl, 121, 122
breakpoint, 139
 control register access, 141
 graphics, 142
 I/O, 141
 memory, 139
 set-pattern, 140
 set-prefix, 140
 set-substr, 140
 symbolic, 155
 temporal, 141
 text output, 142

C

c, 49
caches
 simulation, 197
 workload positioning, 203
callback, 17, 93

cd, 51
CD-ROM, 76
 image files, 76
checkpoint, 17, 58
CLI, 83
 -> operator, 86
 attributes, 86
 if-else statement, 84
 local variable, 85, 88
 script branch, 86
 variable, 83
 while statement, 84
clock cycle, 189
command line interface, 17, 173
 accessing commands from Python, 92
 argument resolving, 174
 commands, 50
 expression, 175
 help system, 176
 namespace commands, 175
 operators, 175
 tab completion, 176
 variable, 175
commands
 namespace, 175
component, 17, 60, 61, 64
configuration, 17, 53
 access from Python, 90
 object, 90
 scripted, 91
configuration object, 17
connecting to a real network, 105
connector, 61
context, 17, 151, 221
 context switcher, 152, 225
 current, 17, 151, 221
 changing, 221

context switcher, [152](#), [225](#)
 context-switcher, [221](#)
 continue, [49](#)
 CPI, [191](#)
 CPU usage, [24](#)
 craff, [17](#), [75](#)
 current context, [17](#), [151](#), [221](#)
 cycle, [18](#), [189](#)
 cycle rate, [191](#)

D

date, [51](#)
 debug information, [156](#)
 debugging, [146](#), [151](#), [221](#)
 GDB, [146](#), [150](#)
 hindsight, [148](#)
 memory spaces, [156](#)
 remote, [146](#)
 scripted, [159](#)
 shared libraries, [147](#)
 symbolic, [146](#), [151](#)
 device, [18](#)
 DHCP, [101](#)
 dirs, [51](#)
 disabling processors, [192](#)
 disks, [69](#)
 building from multiple files, [74](#)
 CD-ROM images, [76](#)
 copying real, [80](#)
 floppy, [77](#)
 host CD-ROM, [76](#)
 images in craff format, [75](#)
 loopback mounting, [73](#)
 MBR, [74](#)
 display, [50](#)
 Dynamic Host Configuration Protocol, [101](#)

E

echo, [51](#)
 ELF, [156](#)
 enable-real-time-mode, [24](#)
 enabling processors, [192](#)
 Ethernet, [97](#), [132](#)
 event, [18](#), [189](#)
 execution
 suspending, [192](#)

execution timing, [190](#)
 extension, [18](#)

F

file-cdrom, [76](#)
 floppy, [77](#)
 images, [77](#)
 Forte
 C compiler, [156](#)

G

g-cache, [197](#)
 workload positioning, [203](#)
 GCC, [156](#)
 GDB, [146](#)
 compiling, [150](#)
 gdb-remote, [146](#)
 generic-cache, [197](#)
 get, [49](#)
 glossary, [17](#)
 graphics breakpoints, [142](#)

H

hap, [18](#), [93](#)
 help system, [176](#)
 Hindsight, [18](#), [227](#)
 host machine, [24](#)

I

if, [84](#)
 in-order, [190](#)
 interrupt-script-branch, [88](#)
 IP address, [99](#), [100](#)
 IP addresses, [101](#)
 IPC, [191](#)

L

laptop, [24](#)
 latency, [98](#), [132](#)
 link object, [97](#), [103](#), [132](#)
 list-failed-modules, [50](#)
 list-modules, [50](#)
 list-script-branches, [88](#)
 load-module, [50](#)
 load-persistent-state, [58](#)
 local, [85](#)

INDEX

loopback mounting, 73
ls, 51

M

MAC addresses, 101
magic instruction, 35
 passing arguments, 143
magic-break-disable, 144
magic-break-enable, 144
memory mappings, 147
micro architectural mode, 170
min-latency, 98
minimum latency, 132
module, 18
 commands, 50

N

namespace
 commands, 175
NAPT, 117
NAT, 117
network
 BOOTP, 101
 bridging, 111, 120
 connecting to real network, 105
 DHCP, 101
 distribution, 103
 DNS, 102
 Ethernet, 97
 gateway, 100
 host connection, 112
 latency, 98
 NAPT, 117
 NAT, 117
 port forwarding, 111, 113
 routing, 100, 112
 Serial, 103
 TAP, 105
 TFTP, 102
nm
 output format, 156

O

openif, 105
operators
 precedence, 175

P

pfregs, 49
pipe, 181
plain-symbols, 156
popd, 51
port forwarding, 111, 113
pregs, 49
print, 51
process
 follow, 151
process tracker, 151, 221, 222
processor
 disabling, 192
 enabling, 192
profiling, 161
pselect, 49
pstats, 49
ptime, 49
pushd, 51
pwd, 51
Python, 18, 89, 159

R

read-reg, 49
recorder, 139
remote GDB, 146
reverse execution, 18, 227
run-command-file, 50
run-python-file, 50, 90
run_command, 92

S

save-persistent-state, 58
sc, 49
script, 83
 commands, 49
script branch, 86
 commands, 87
 local variable, 88
 wait-for, 87
 wait-for-hap, 87
serial, 103
service node, 99
service-node, 99
set, 49
set-pattern, 140

set-prefix, [140](#)
 set-substr, [140](#)
 si, [49](#)
 Simics Central, [170](#)
 SimicsFS, [18](#)
 simulated time, [189](#), [192](#)
 STABS, [156](#)
 stall, [190](#)
 stall mode, [170](#)
 stalling, [192](#)
 stalling period, [190](#)
 STC, [19](#)
 step, [19](#), [189](#)
 step rate, [191](#), [192](#)
 step-cycle, [49](#)
 stepi, [49](#)
 Sun Workshop
 C compiler, [156](#)
 symbols
 loading, [156](#)
 symtable, [151](#)

T

tab completion, [176](#)
 TAP, [105](#)
 TCP/IP, [97](#)
 TFTP, [102](#)
 time
 suspending, [192](#)
 timing models, [190](#)
 trackee, [223](#)
 tracker, [223](#)
 tuncctl, [107](#)
 tutorial, [29](#)

U

undisplay, [51](#)

V

variable, [83](#), [88](#), [90](#)

W

wait-for-cycle, [88](#)
 wait-for-hap, [87](#), [88](#)
 wait-for-step, [88](#)
 wait-for-string, [87](#), [88](#)

wait-for-variable, [88](#)
 watchpoint, [155](#)
 while, [84](#)
 workspace, [19](#)
 write-reg, [49](#)

X

x, [49](#)

Z

zsh, [151](#)



Virtutech, Inc.

1740 Technology Dr., suite 460
San Jose, CA 95110
USA

Phone +1 408-392-9150
Fax +1 408-608-0430

<http://www.virtutech.com>